



Apache Flink Streaming

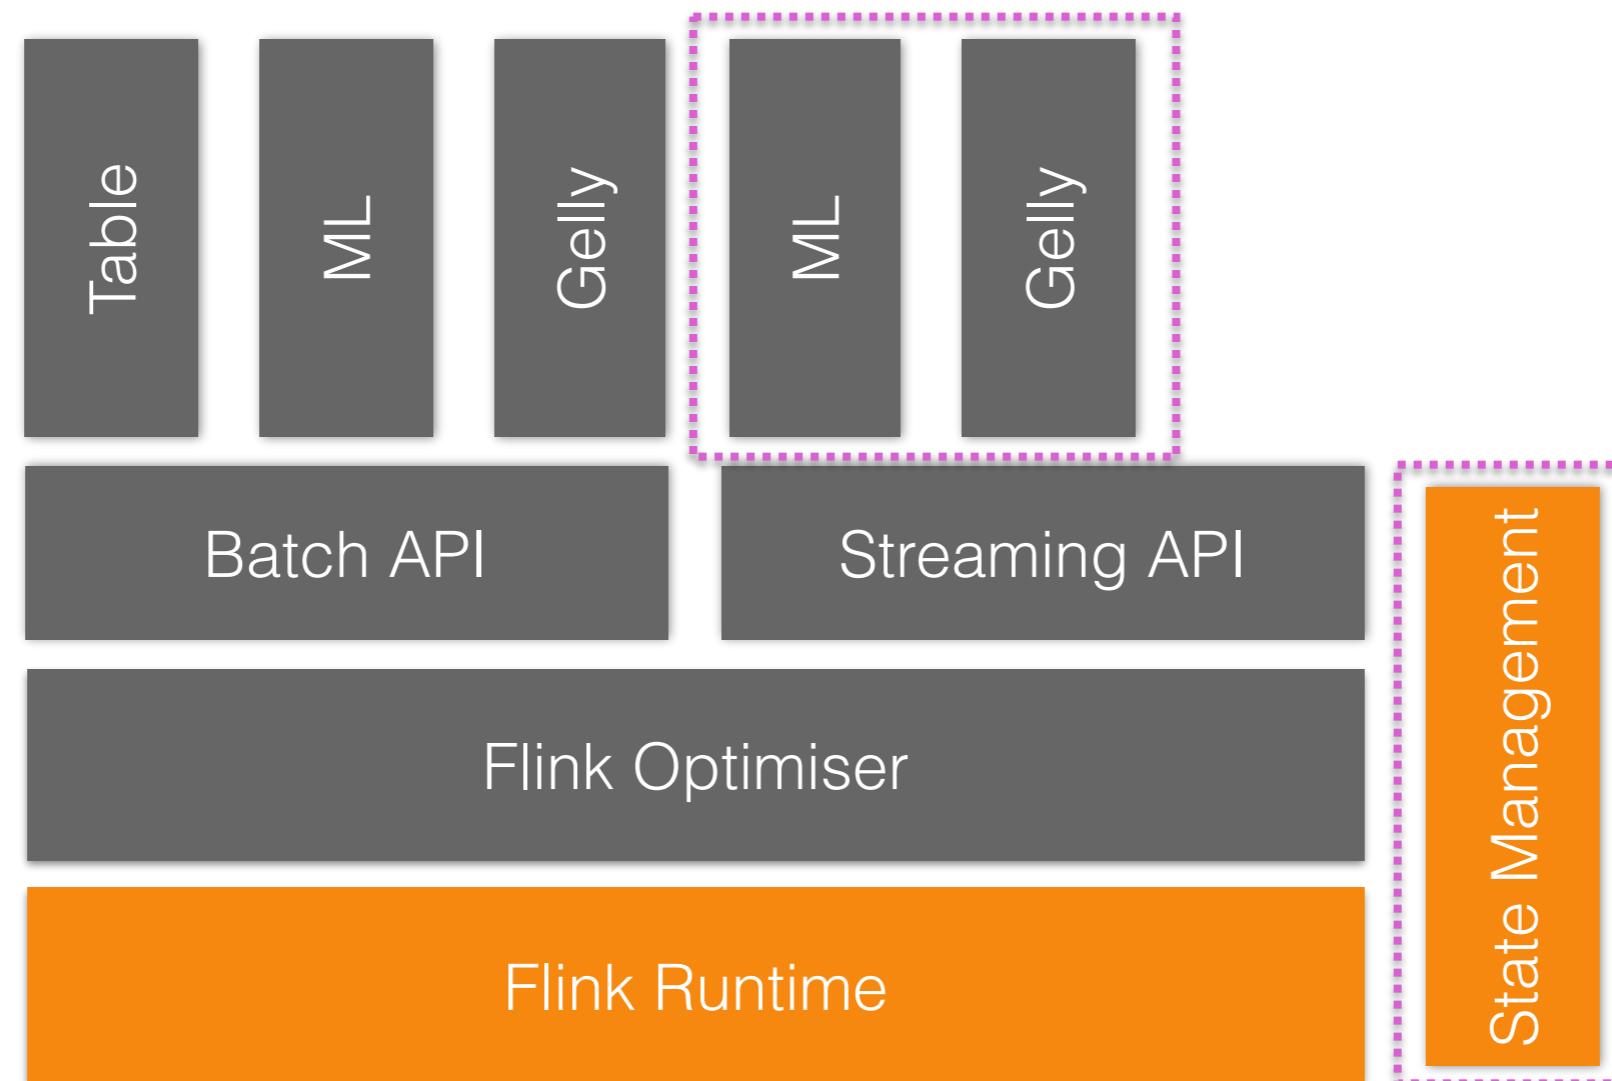
Resiliency and Consistency

Paris Carbone - PhD Candidate
KTH Royal Institute of Technology
parisc@kth.se, senorcarbone@apache.org

Overview

- The Flink Execution Engine Architecture
- Exactly-once-processing challenges
- Using Snapshots for Recovery
- The ABS Algorithm for DAGs and Cyclic Dataflows
- Recovery, Cost and Performance
- Job Manager High Availability

Current Focus

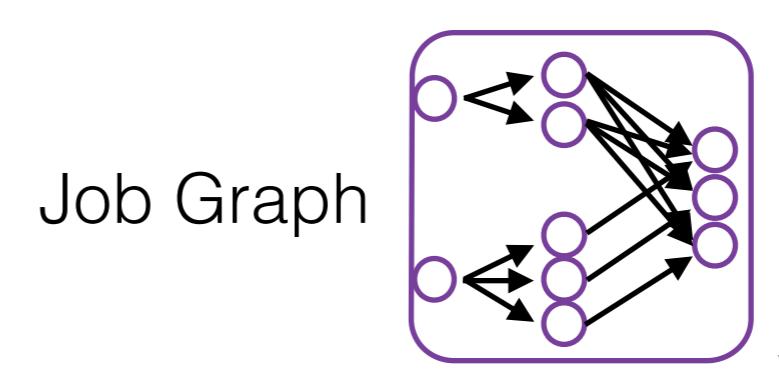


Executing Pipelines

```
val lines: DataStream[String] = env.fromSocketStream(...)  
lines.flatMap {line => line.split(" ")  
    .map(word => Word(word,1))  
    .window(Time.of(5,SECONDS)).every(Time.of(1,SECONDS))  
    .groupByKey().sum("frequency")  
    .print()}
```

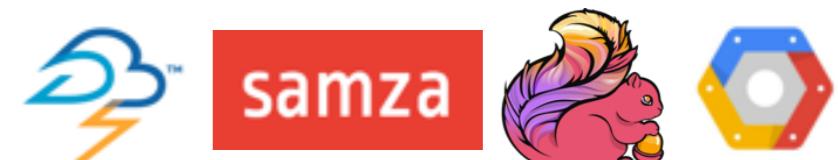


- Streaming pipelines translate into **job graphs**

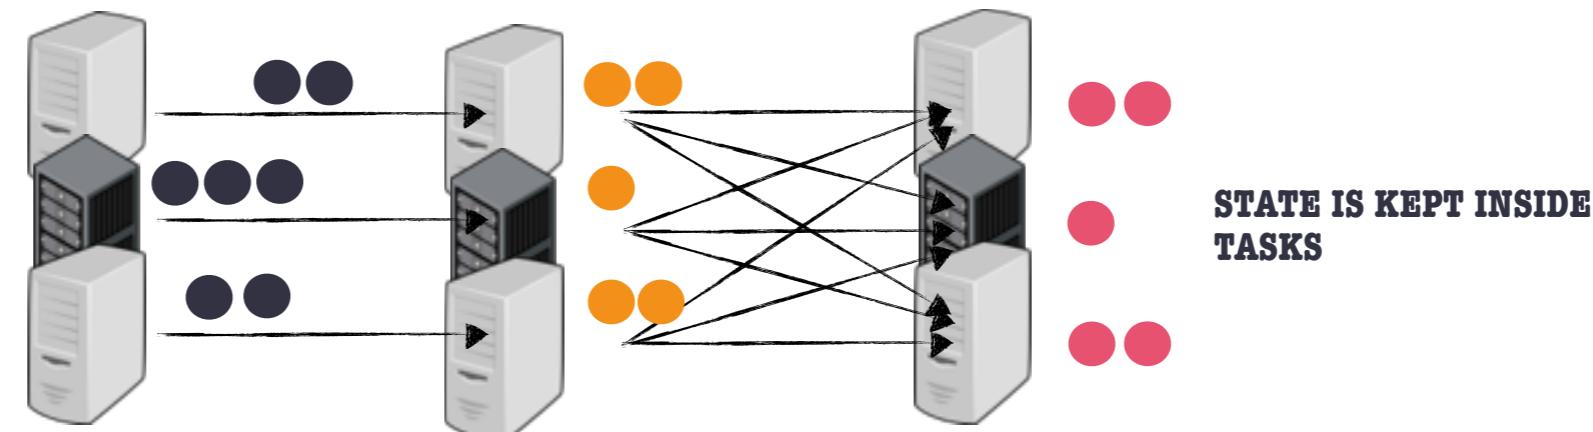


Unbounded Data Processing Architectures

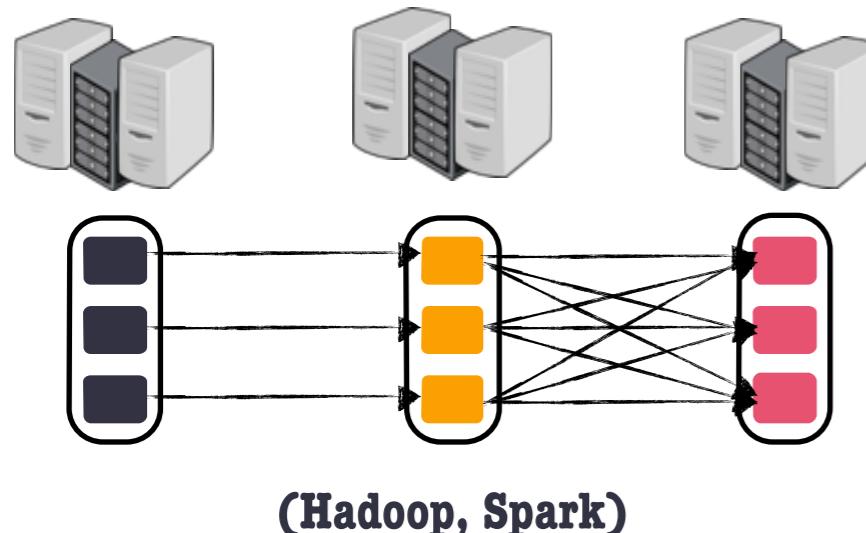
1) Streaming (Distributed Data Flow)



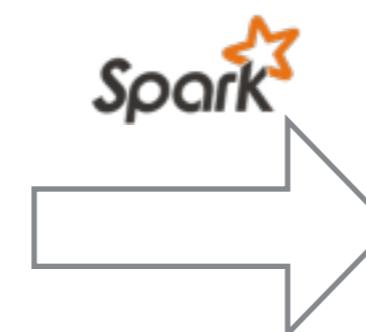
LONG-LIVED TASK EXECUTION



2) Micro-Batch



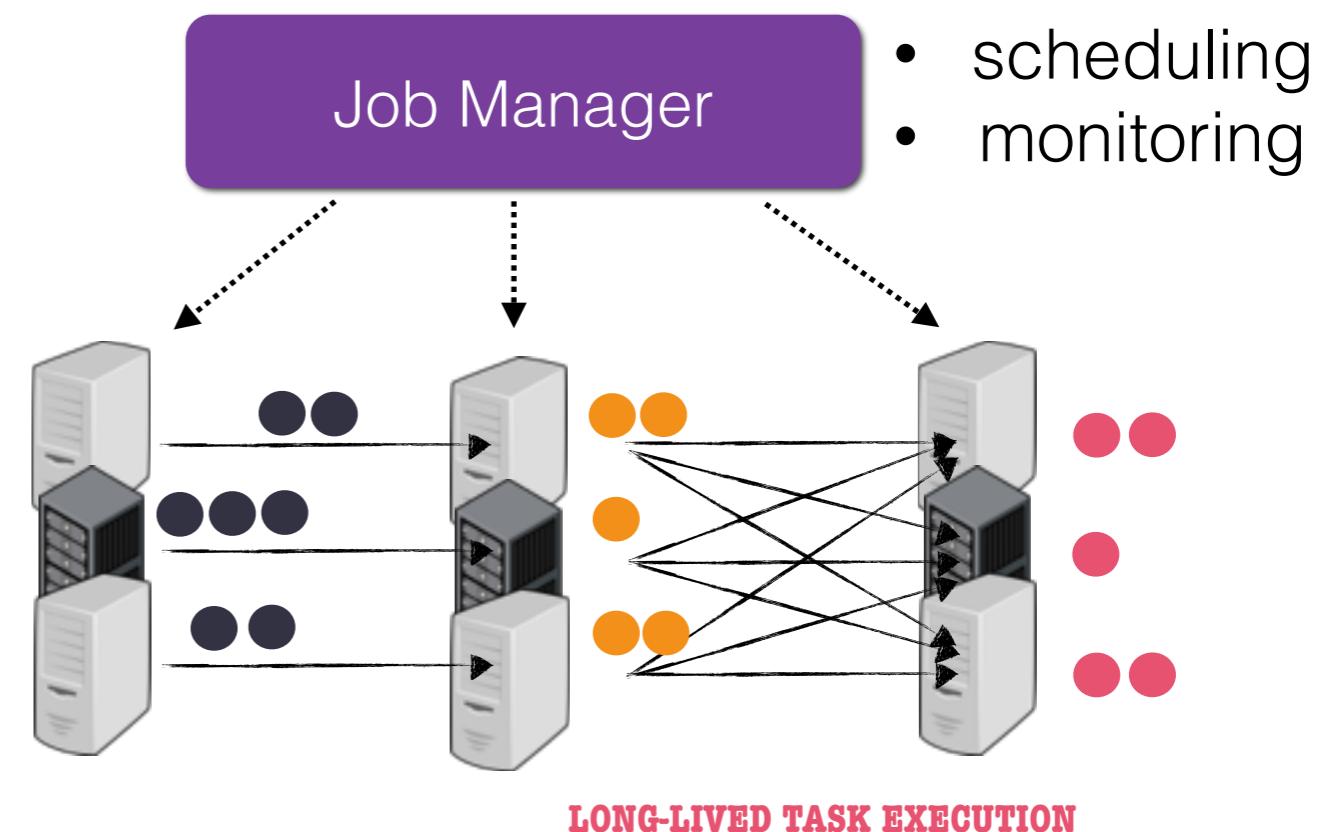
(Hadoop, Spark)



(Spark Streaming)

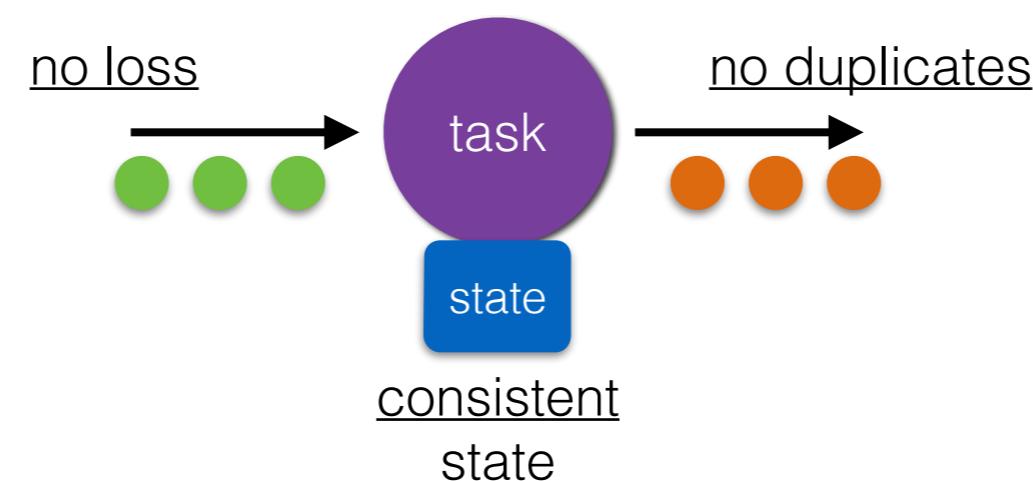
The Flink Runtime Engine

- **Tasks** run operator logic in a pipelined fashion
- They are scheduled among workers
- **State** is kept within tasks



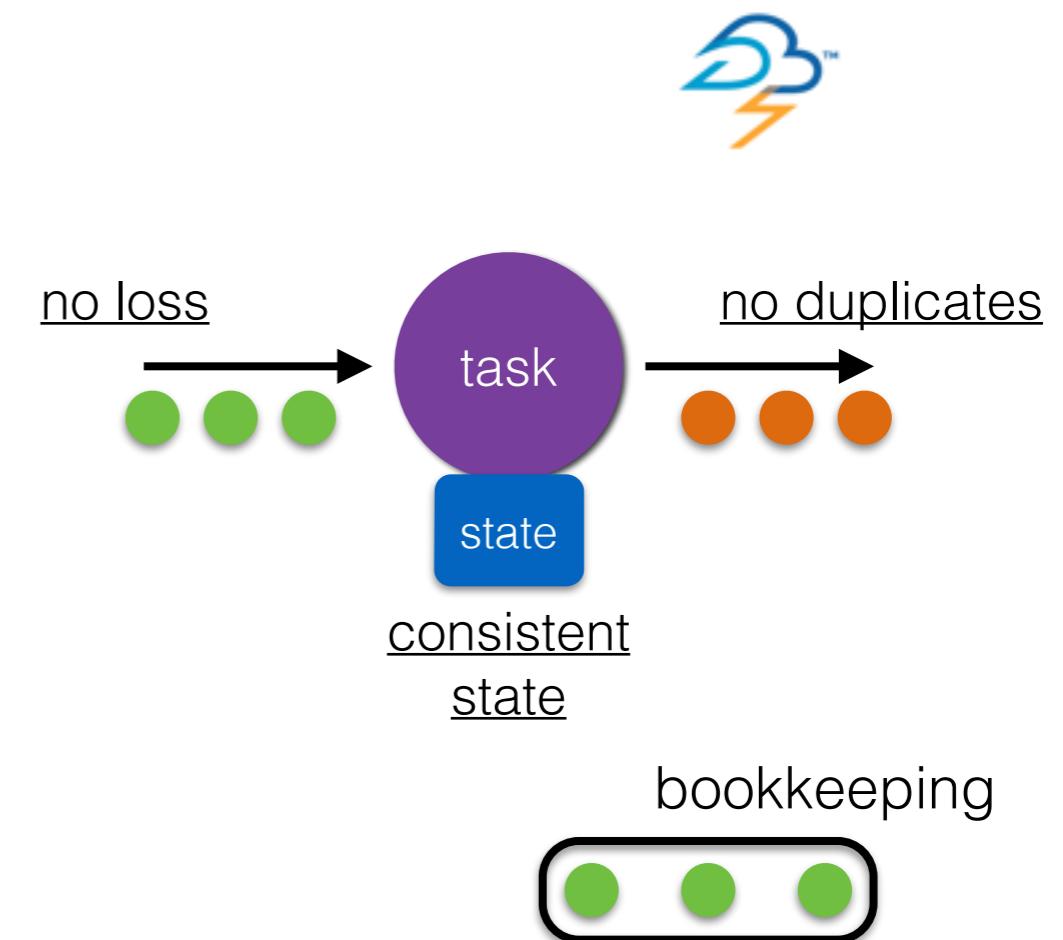
Task Failures

- Task failures are **guaranteed** to occur
- We need to make them **transparent**
- Can we simply recover a task from scratch?



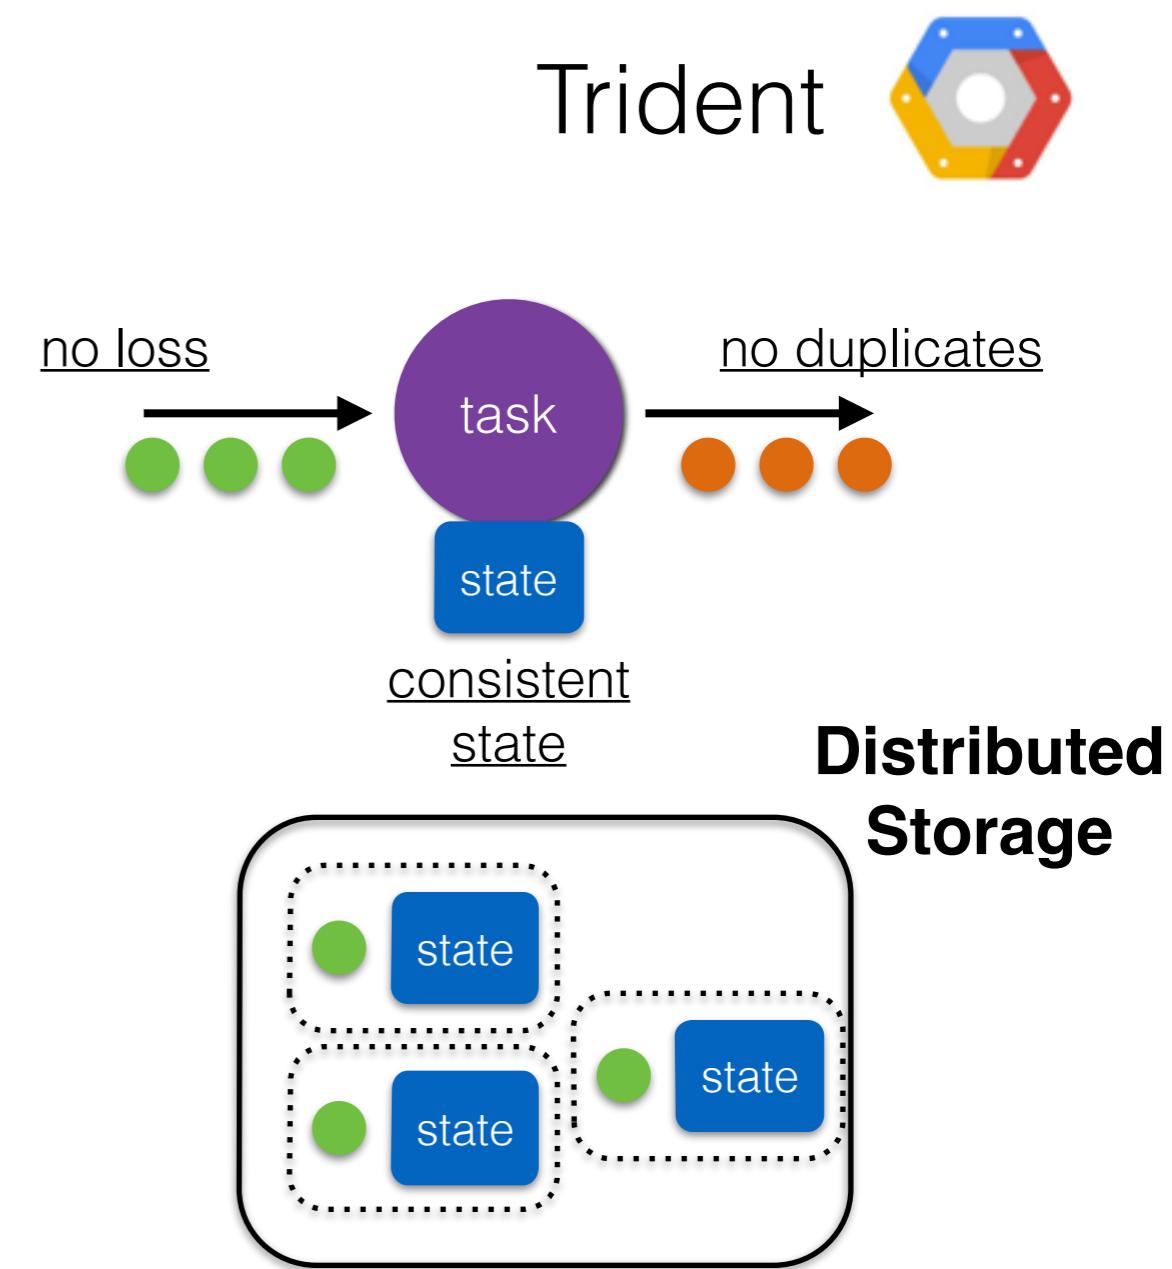
Record Acknowledgements

- We can monitor the consumption of every event
- It guarantees no loss
- Duplicates and inconsistent state are not handled

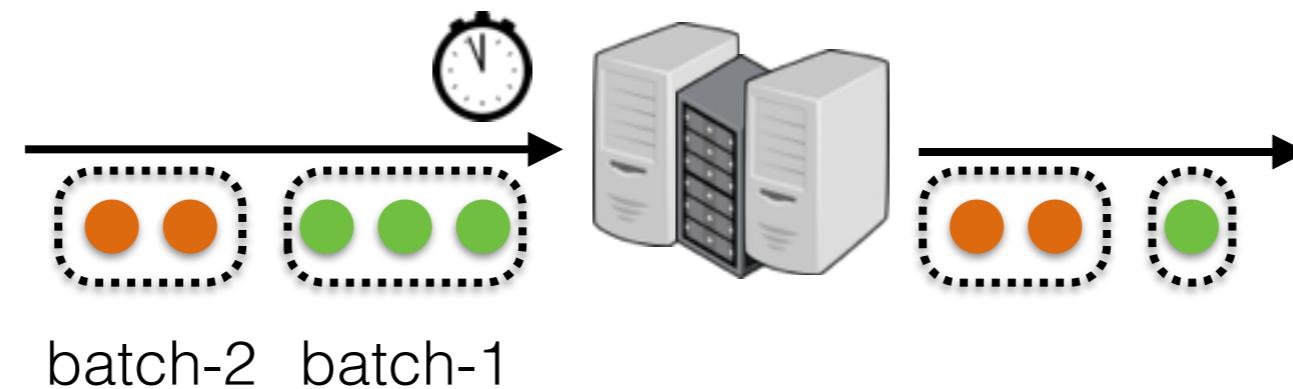


Atomic transactions per update

- It guarantees no loss and consistent state.
- No duplication can be achieved e.g. with bloom filters or caching
- Fine grained recovery
- Non-constant load in external storage can lead to unsustainable execution



Lessons Learned from Batch



- If a batch computation fails, simply repeat computation as a transaction
- Transaction rate is **constant**
- Can we apply these principles to a true streaming execution?

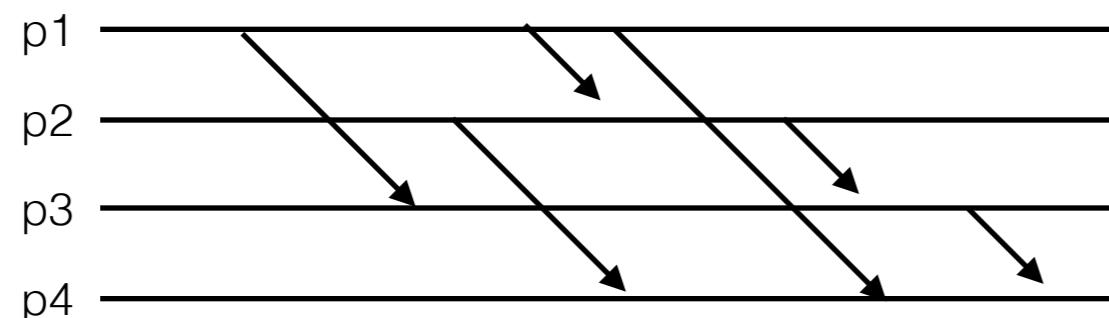
Distributed Snapshots

Distributed Snapshots

*“A collection of operator states
and records in transit (channels) that reflects a
moment at a valid execution”*

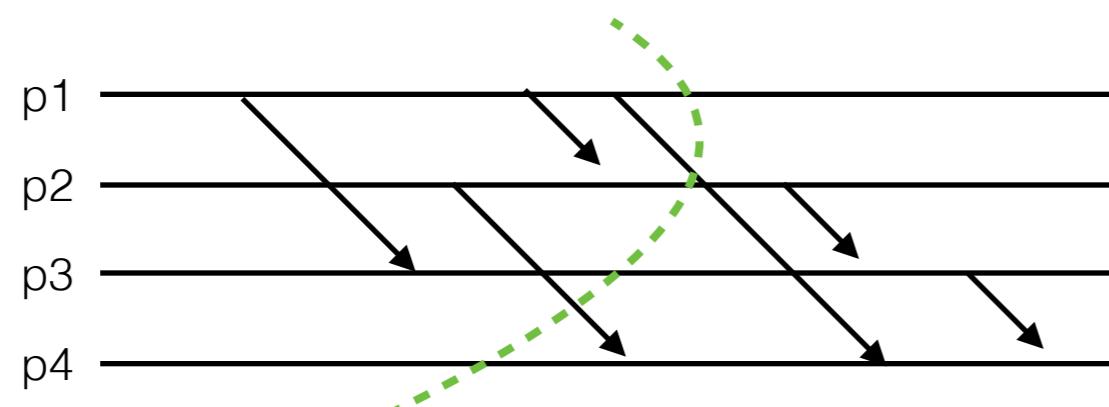
Distributed Snapshots

*“A collection of operator states
and records in transit (channels) that reflects a
moment at a valid execution”*



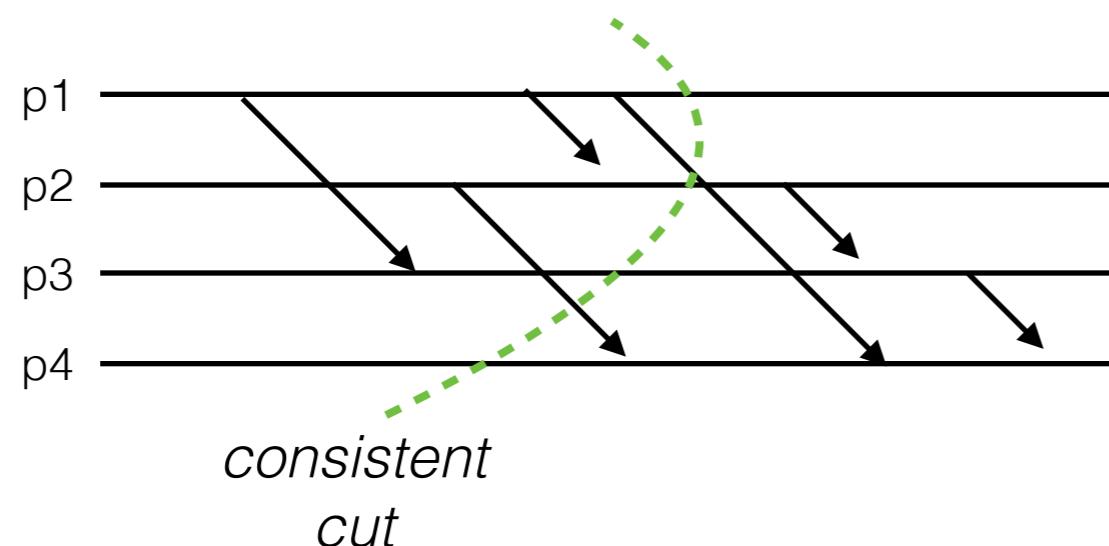
Distributed Snapshots

*“A collection of operator states
and records in transit (channels) that reflects a
moment at a valid execution”*



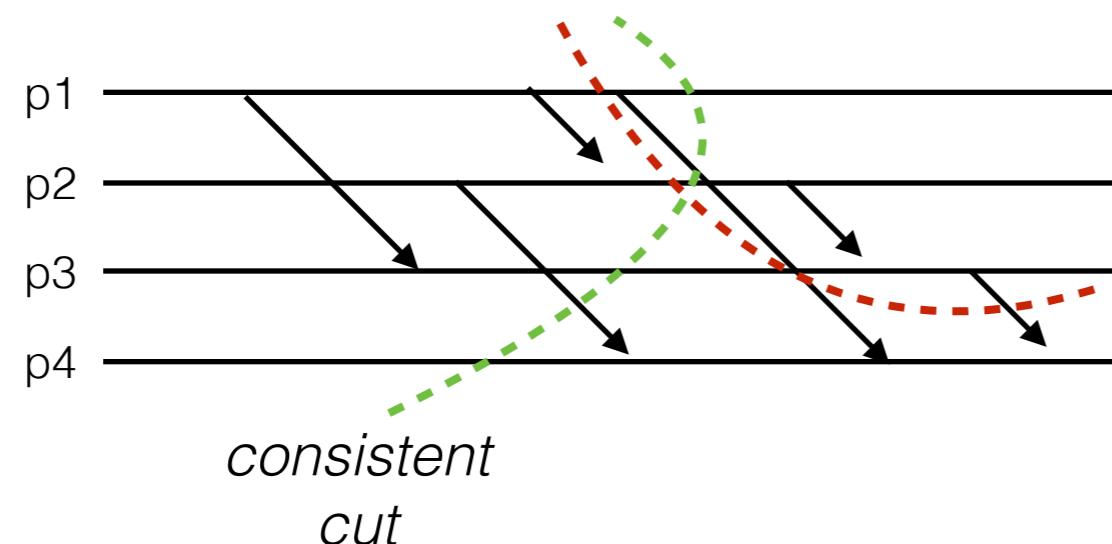
Distributed Snapshots

*“A collection of operator states
and records in transit (channels) that reflects a
moment at a valid execution”*



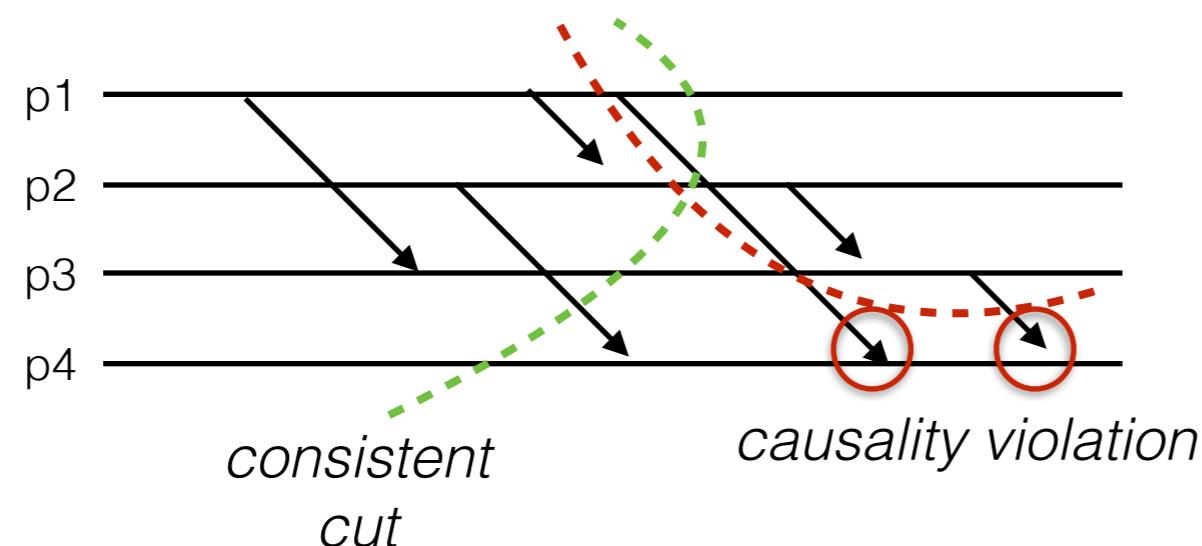
Distributed Snapshots

*“A collection of operator states
and records in transit (channels) that reflects a
moment at a valid execution”*



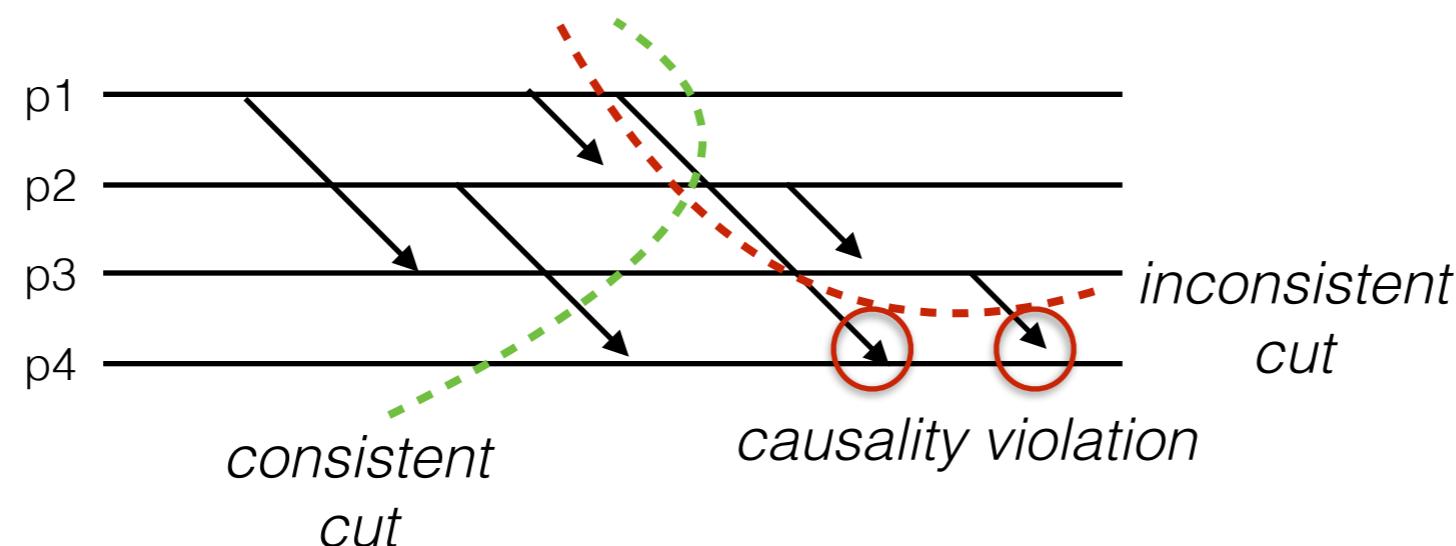
Distributed Snapshots

“A collection of operator states and records in transit (channels) that reflects a moment at a valid execution”



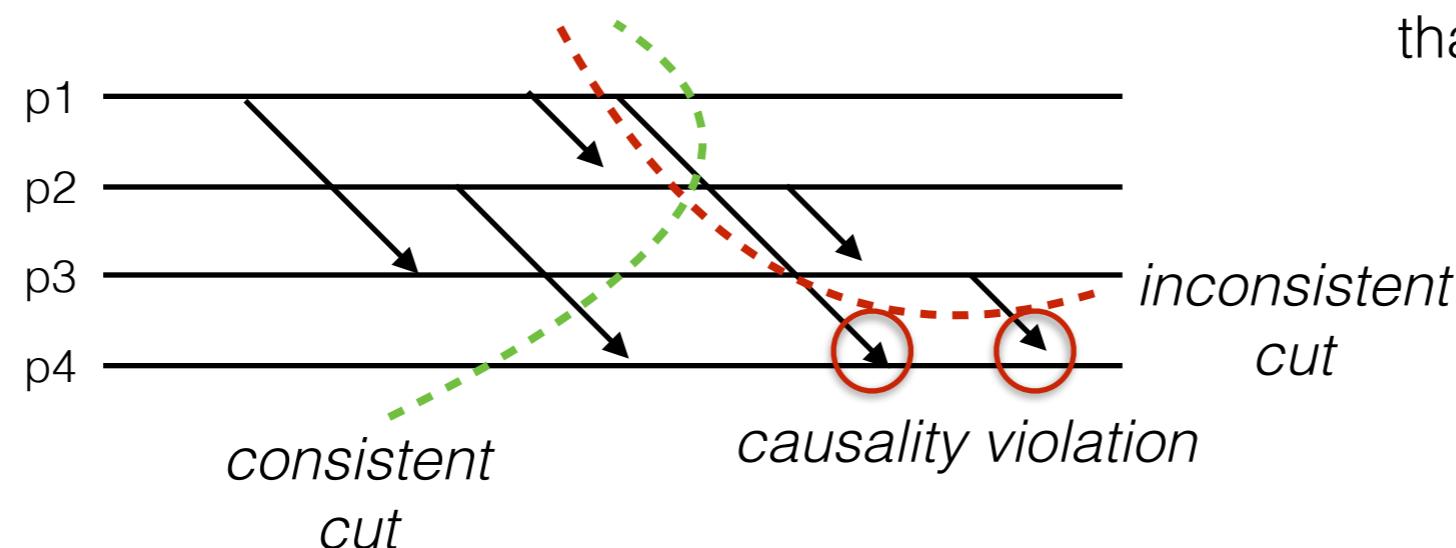
Distributed Snapshots

“A collection of operator states and records in transit (channels) that reflects a moment at a valid execution”



Distributed Snapshots

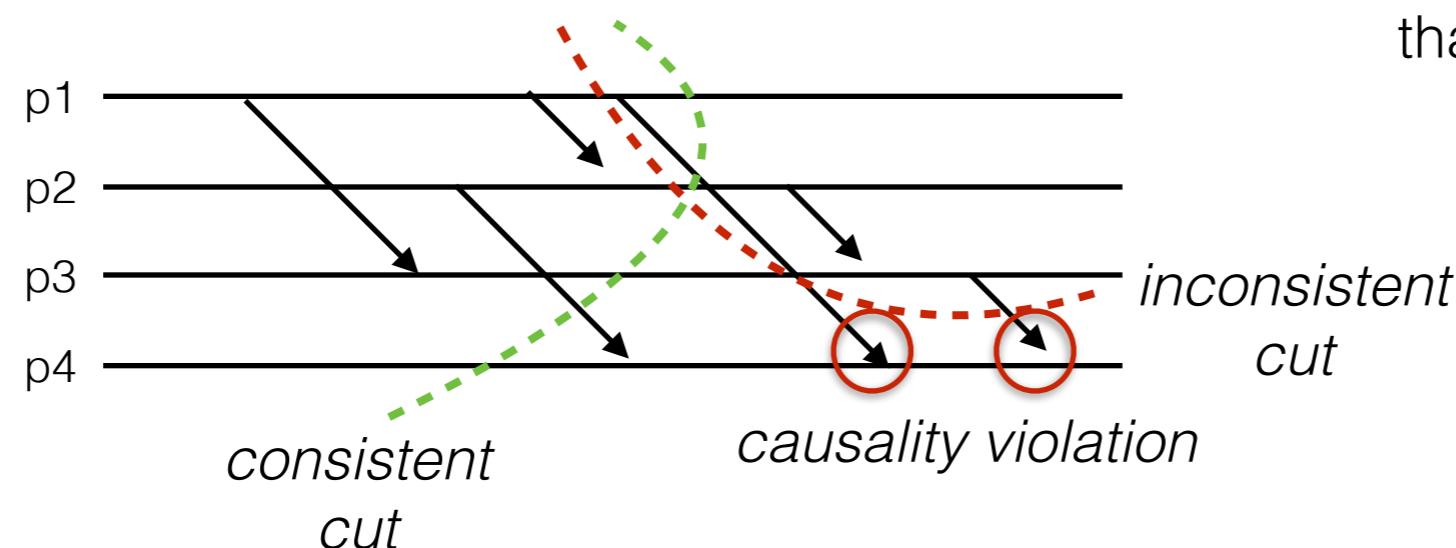
“A collection of operator states and records in transit (channels) that reflects a moment at a valid execution”



Idea: We can resume from a snapshot that defines a consistent cut

Distributed Snapshots

“A collection of operator states and records in transit (channels) that reflects a moment at a valid execution”



Idea: We can resume from a snapshot that defines a consistent cut

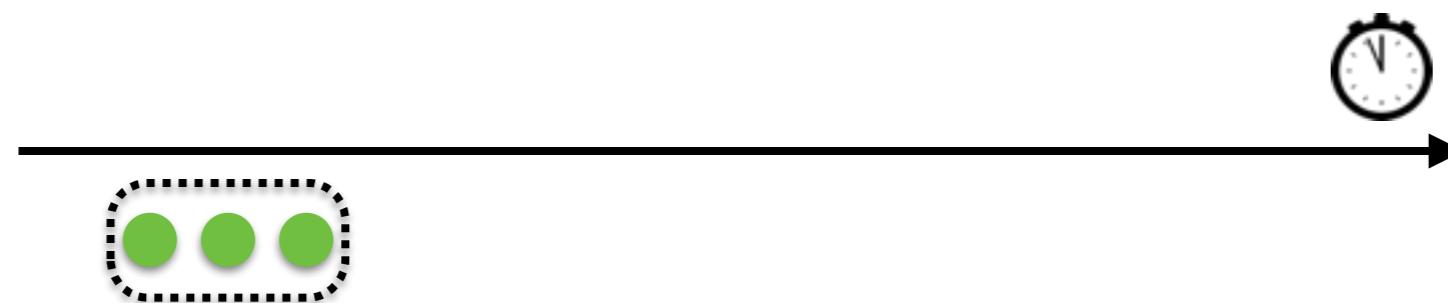
Assumptions

- repeatable sources
- reliable FIFO channels

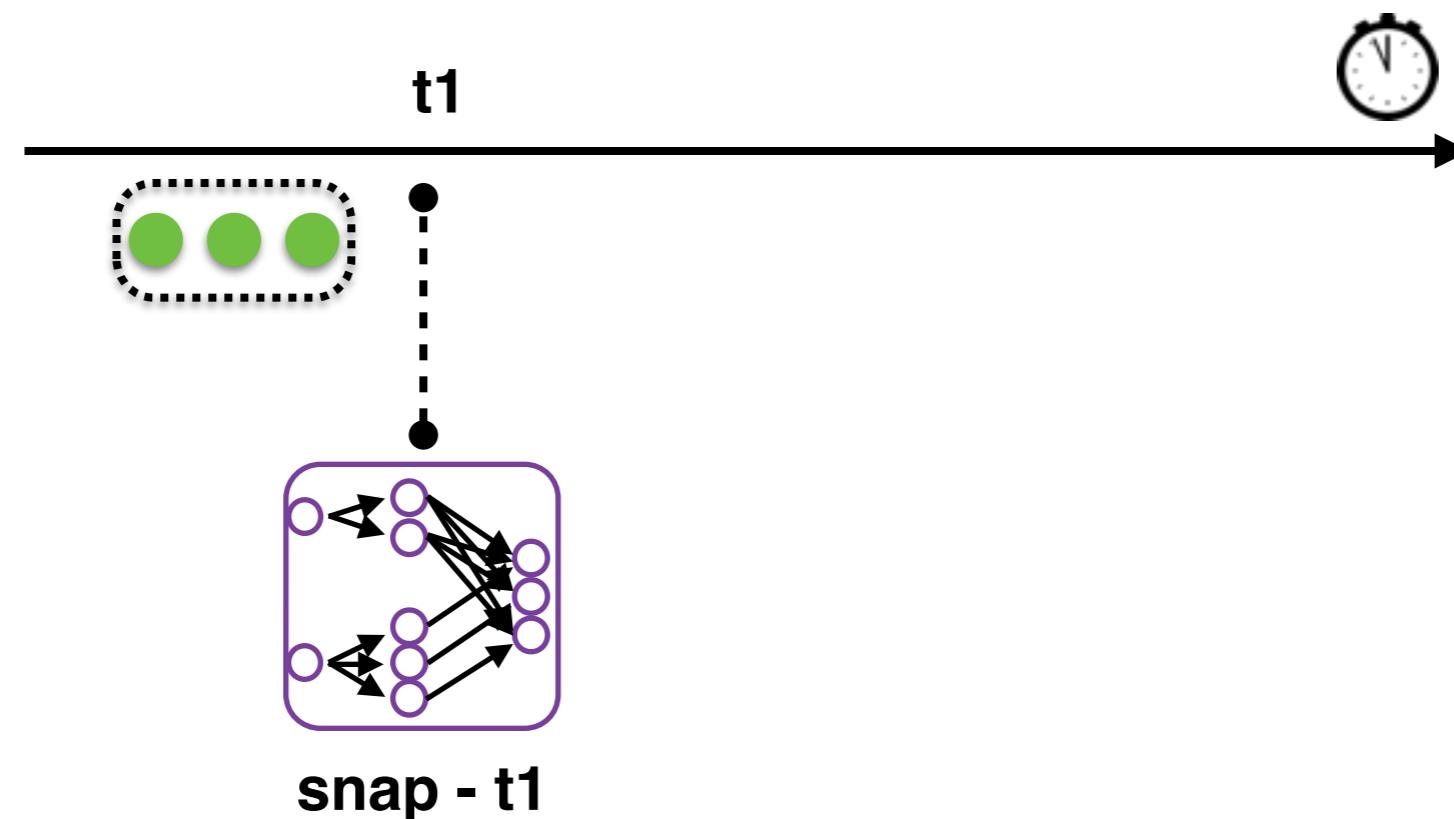
Distributed Snapshots



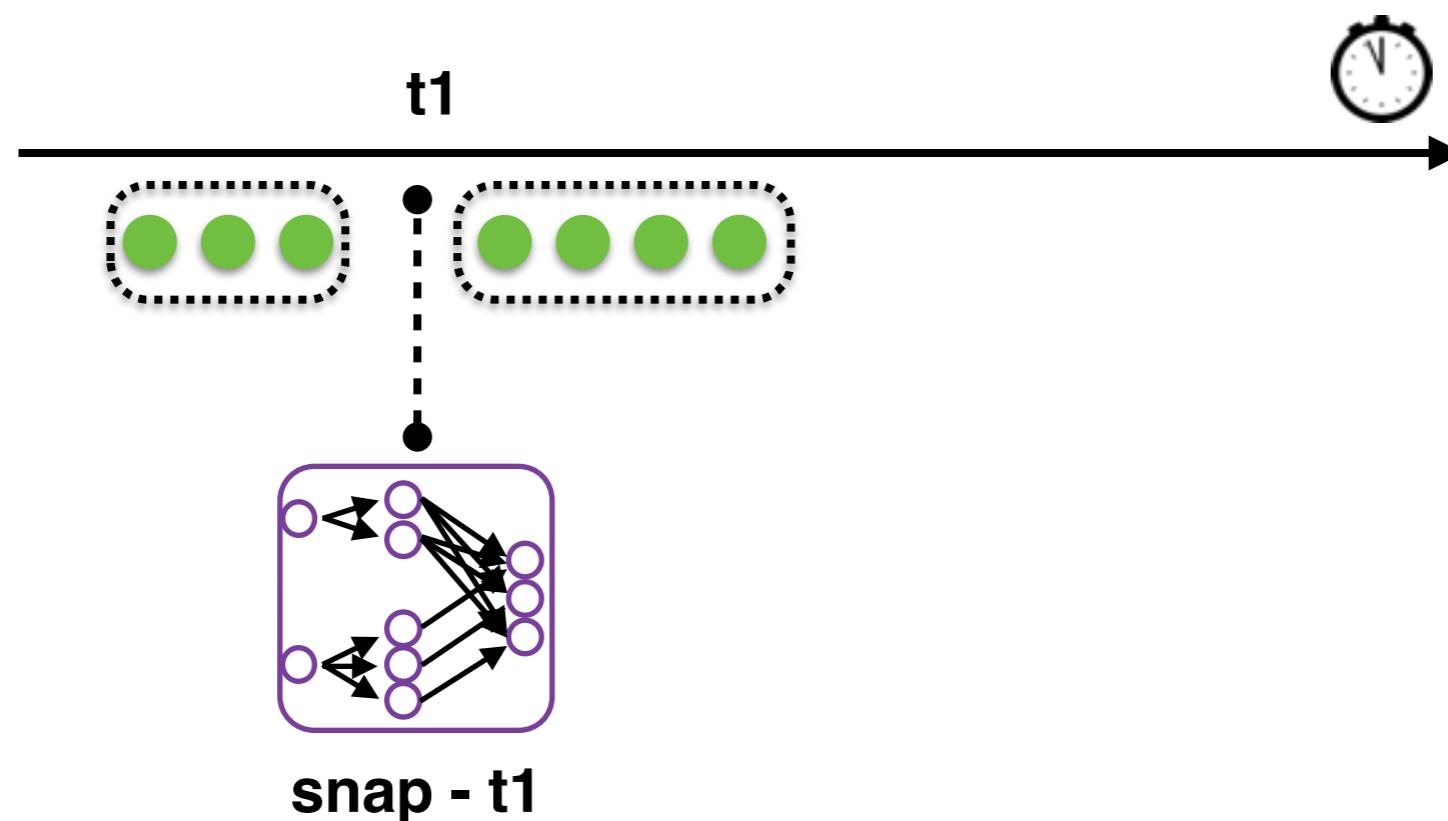
Distributed Snapshots



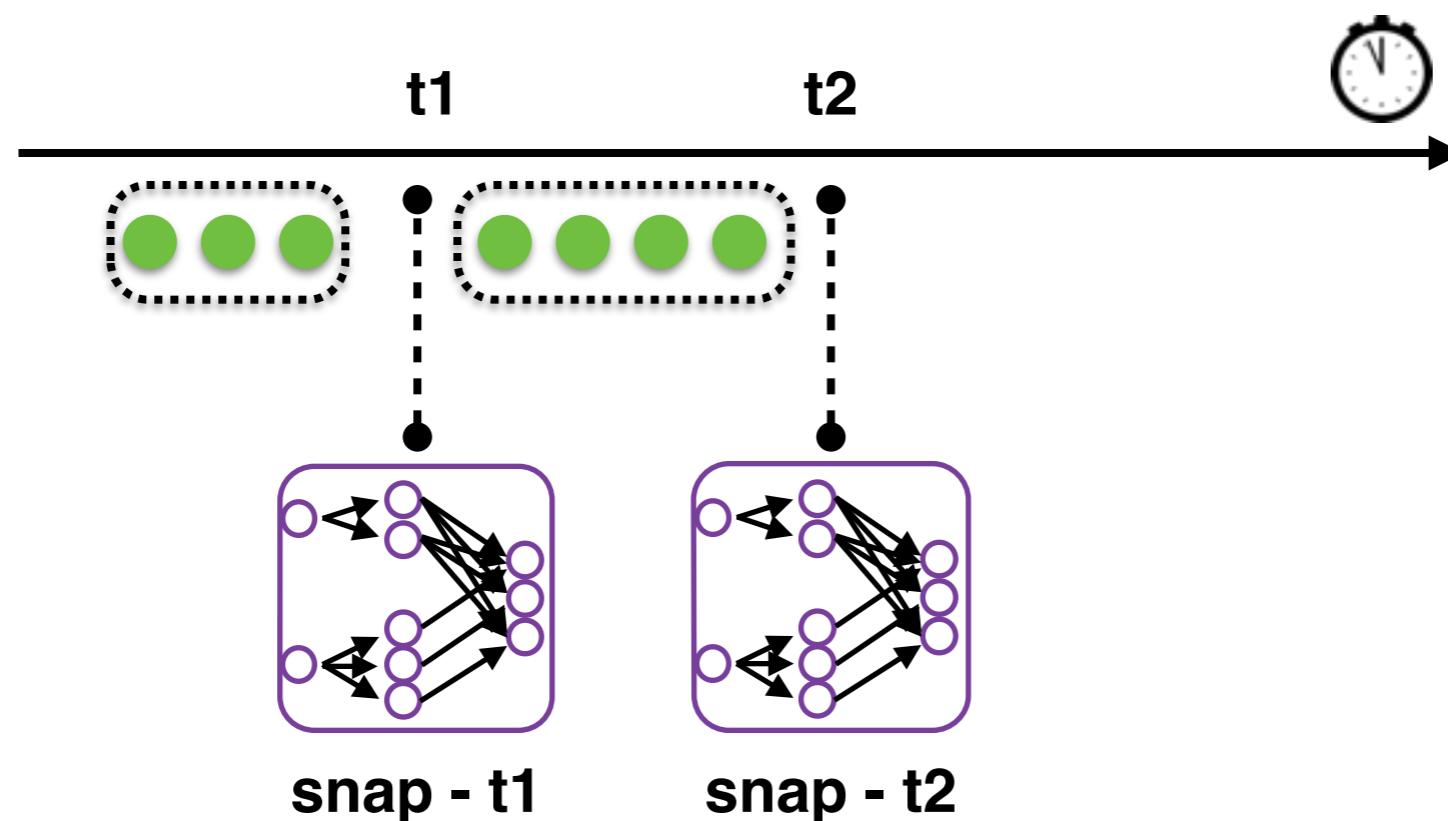
Distributed Snapshots



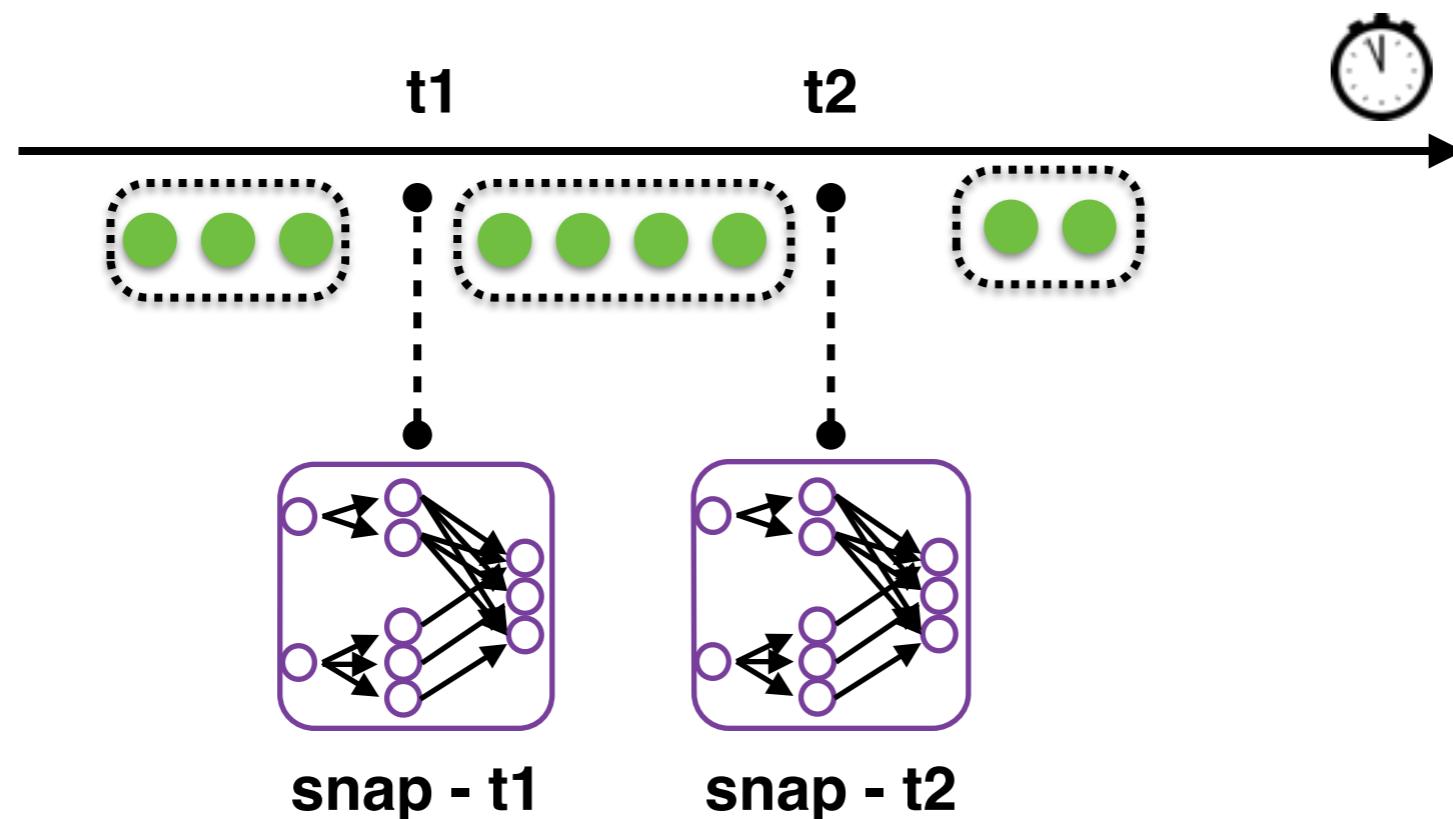
Distributed Snapshots



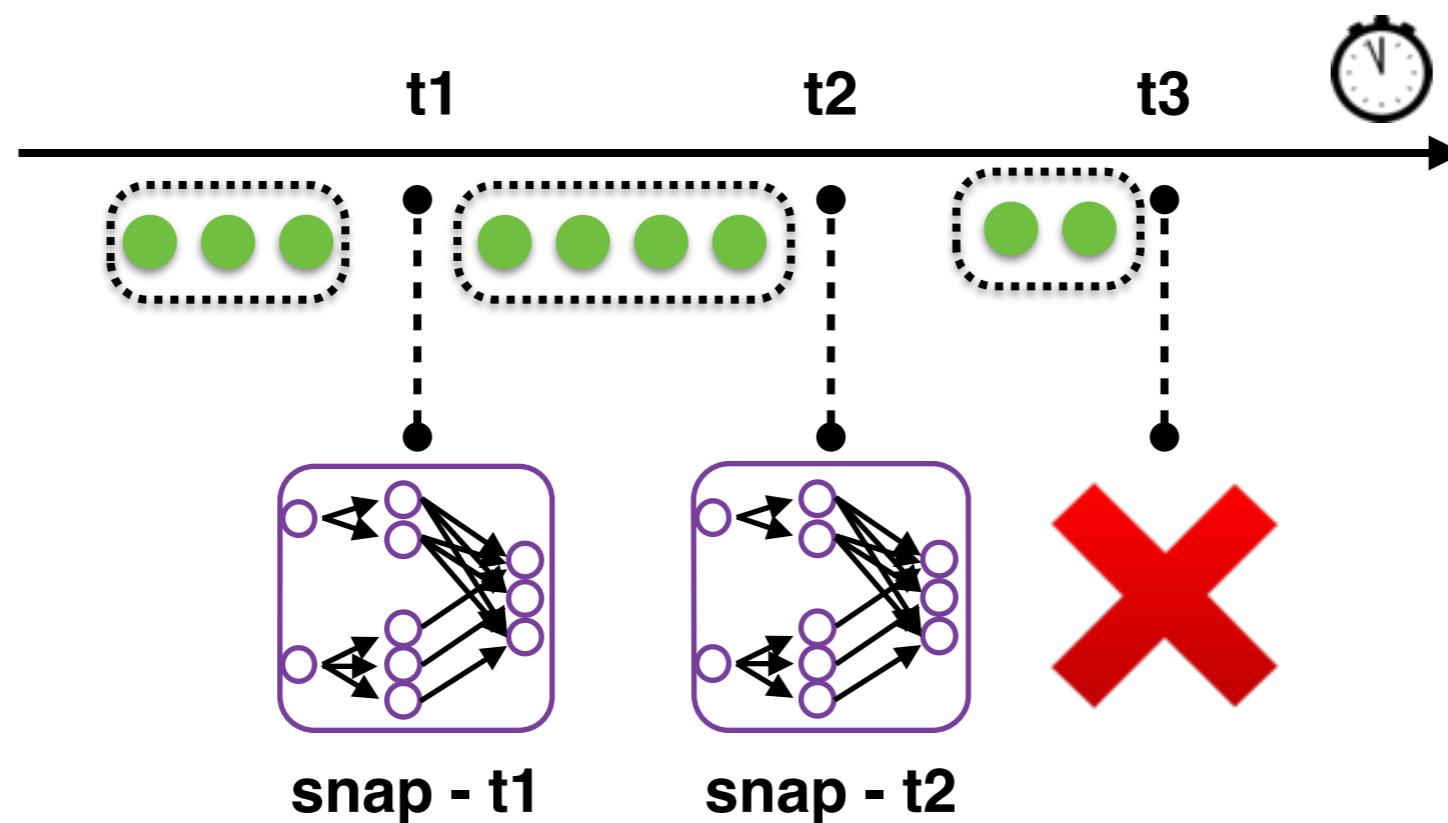
Distributed Snapshots



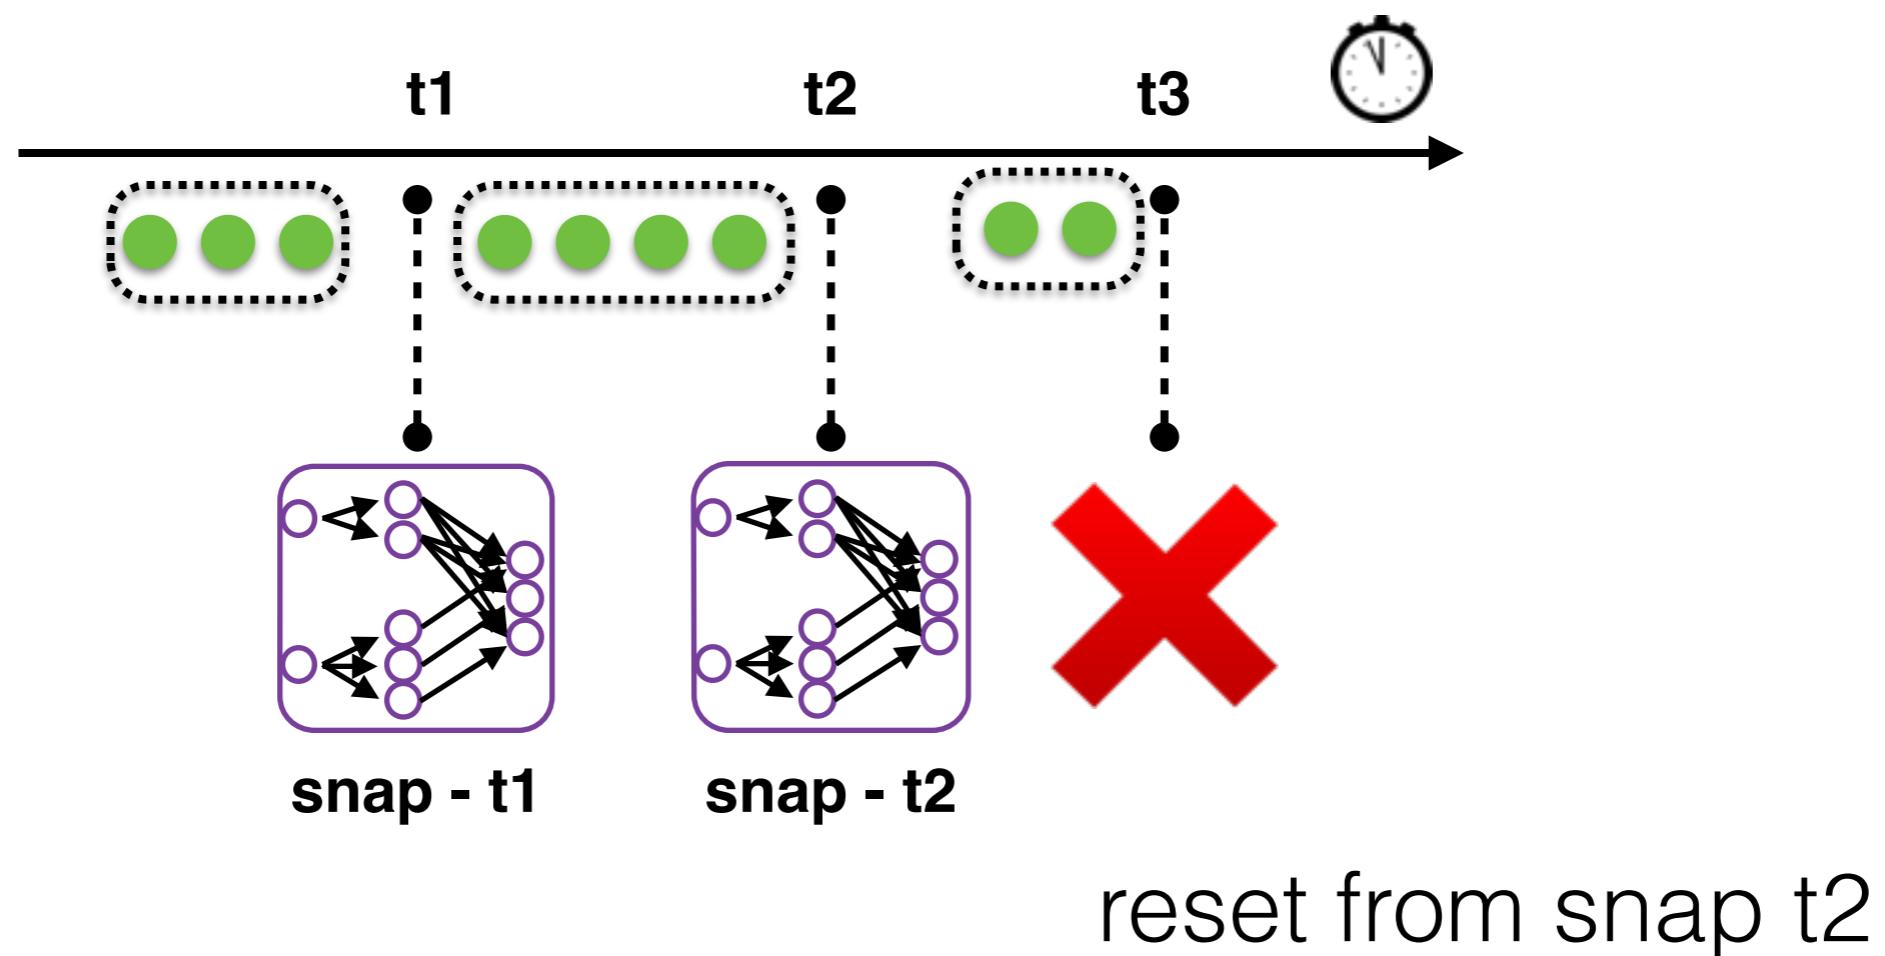
Distributed Snapshots



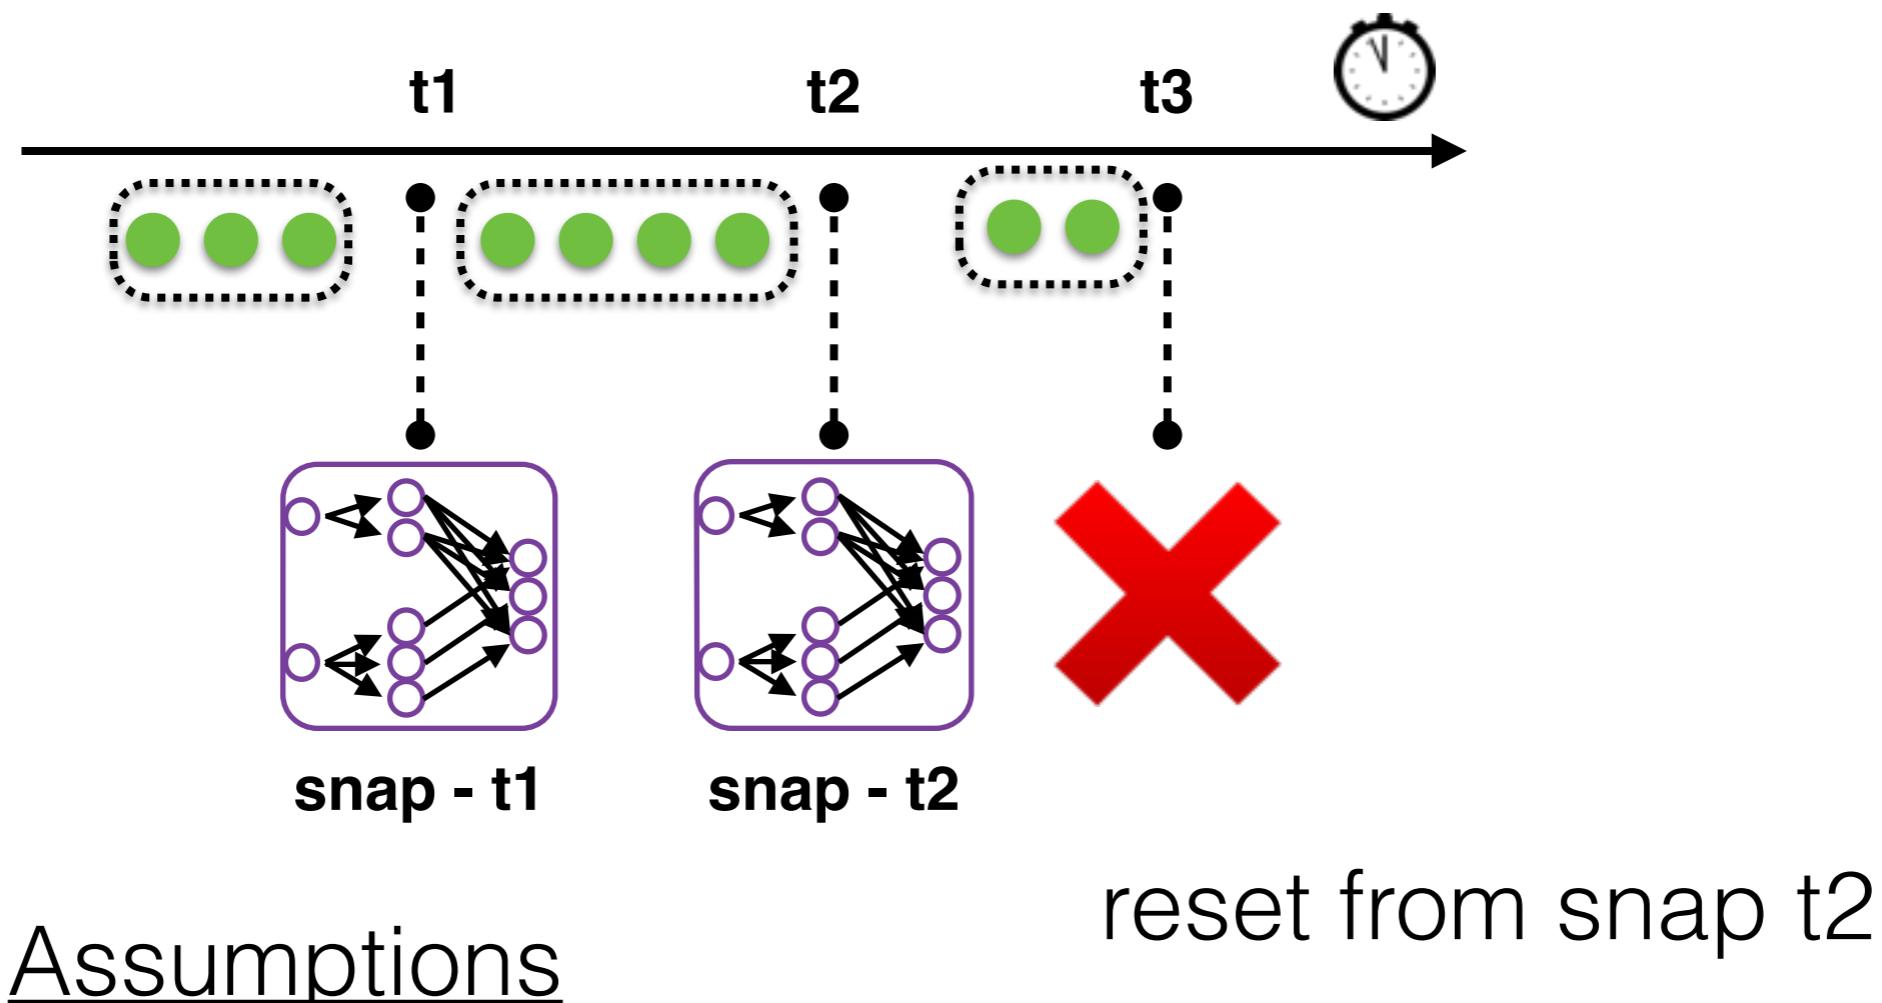
Distributed Snapshots



Distributed Snapshots



Distributed Snapshots

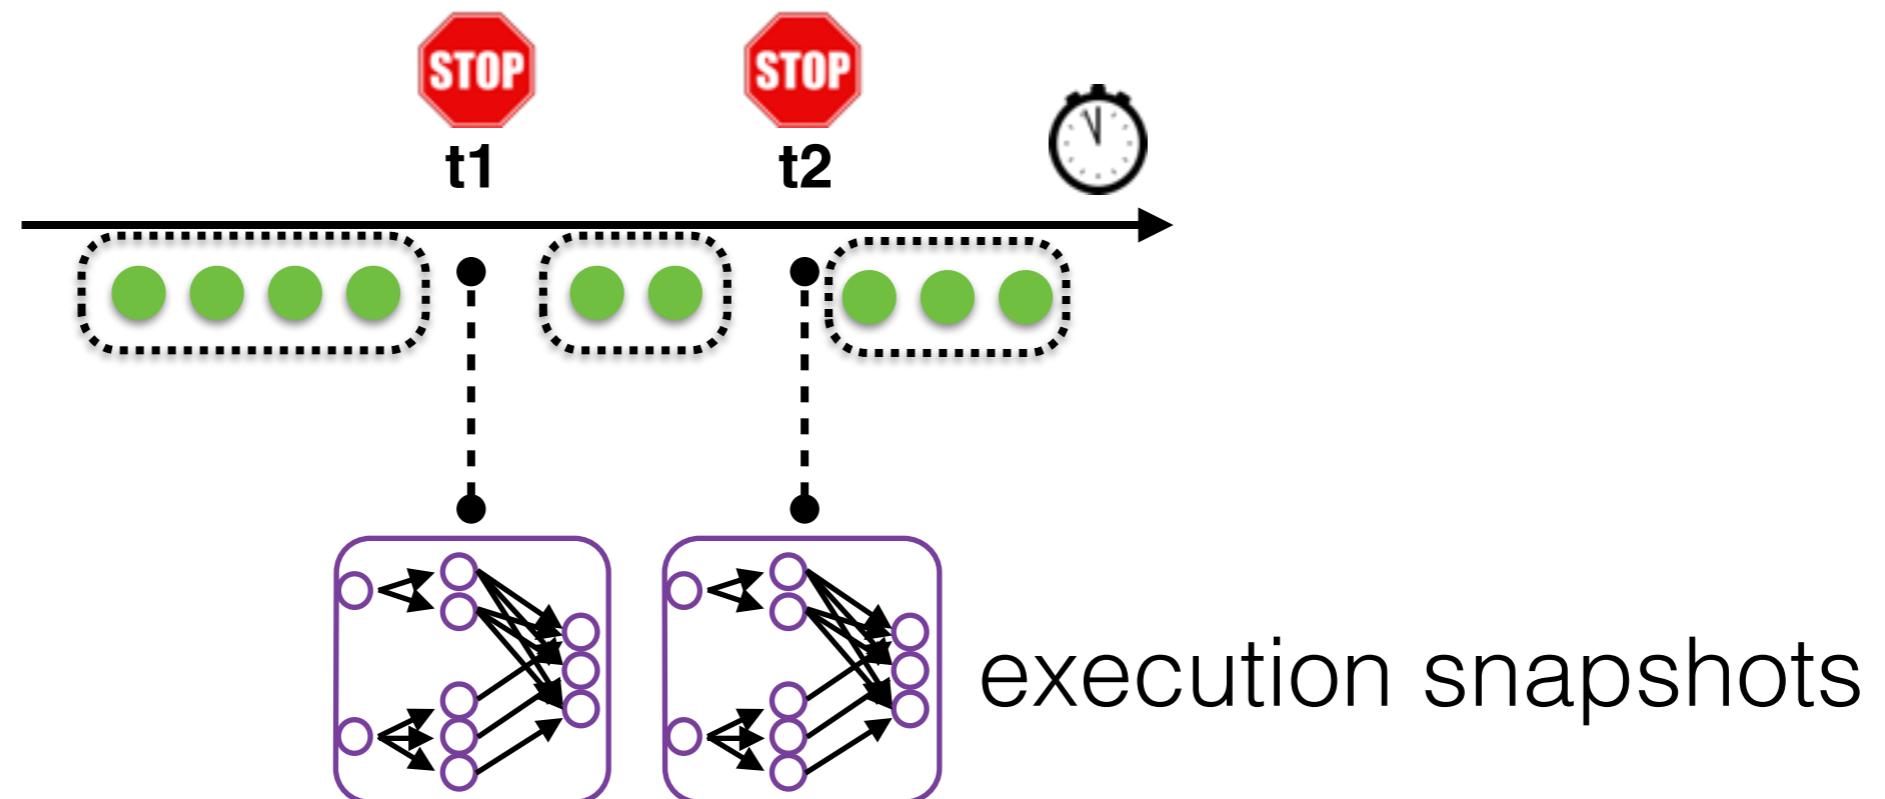


Assumptions

- repeatable sources
- reliable FIFO channels

reset from snap t2

Taking Snapshots



Initial approach (see Naiad)

- Pause execution on t_1, t_2, \dots
- Collect state
- Restore execution

Lamport On the Rescue

“The global-state-detection algorithm is to be superimposed on the underlying computation:

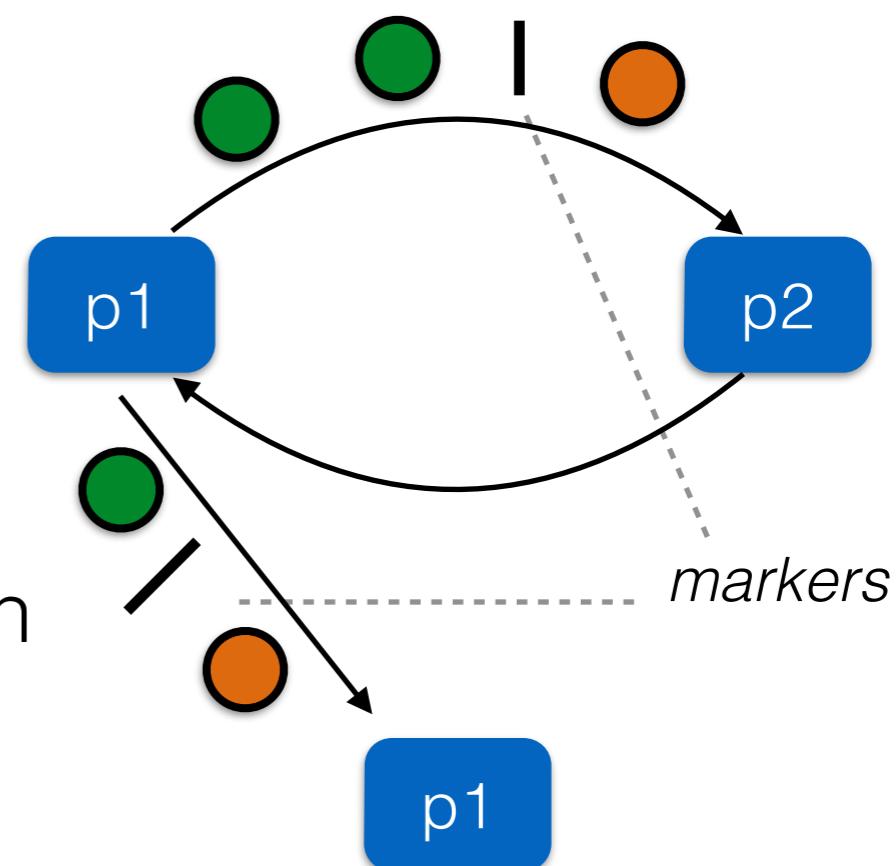
it must run concurrently with, but not alter, this underlying computation”

Chandy-Lamport Snapshots

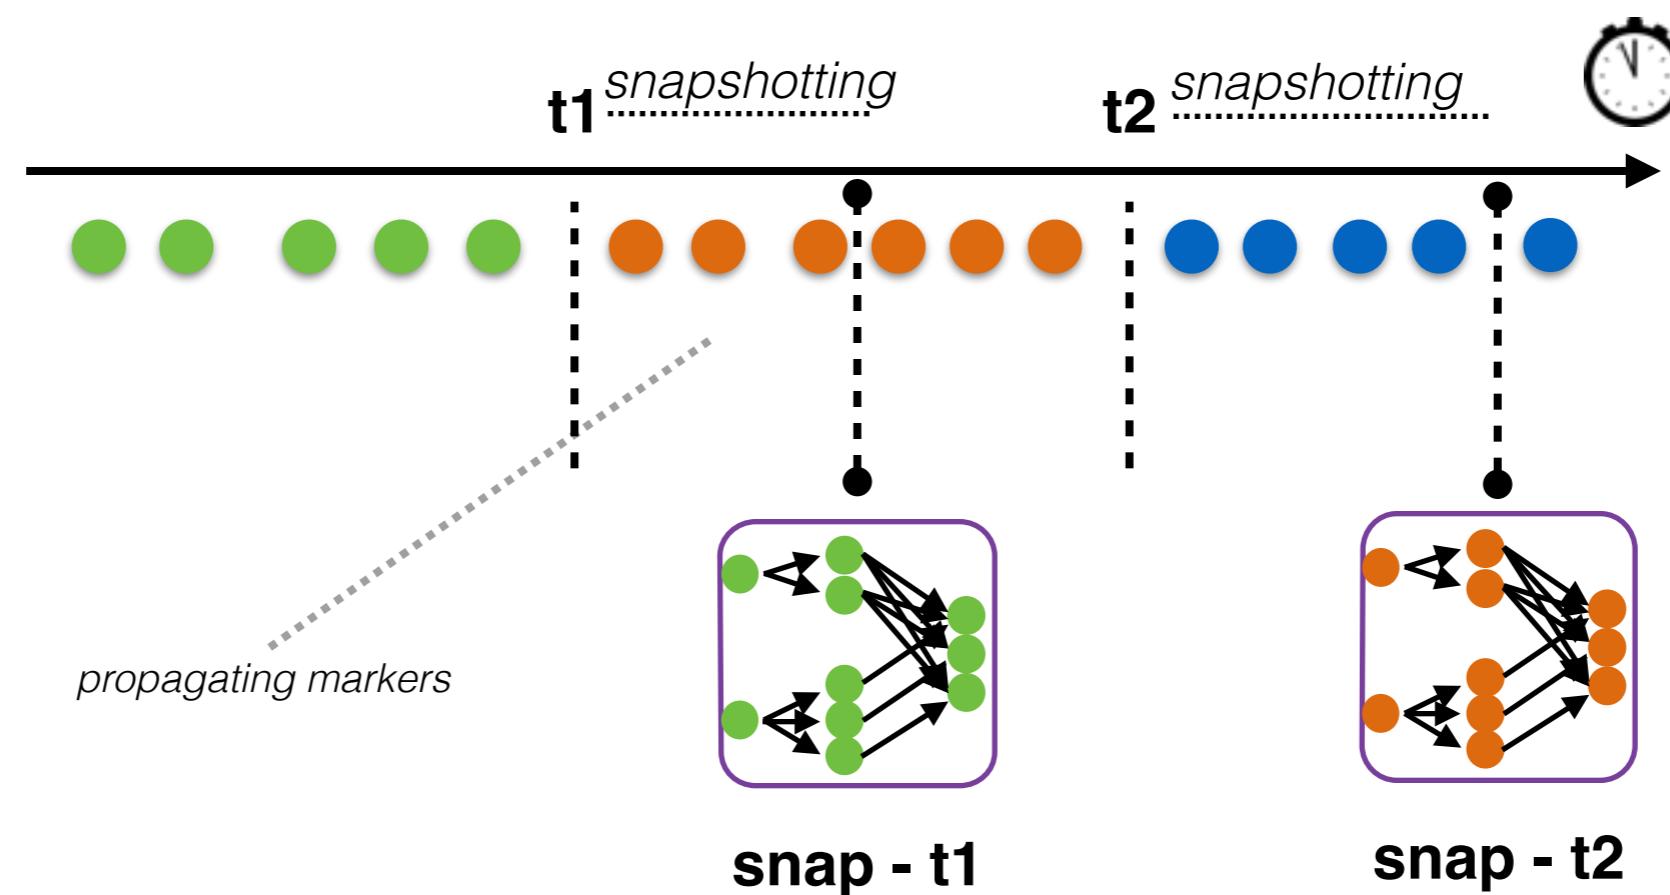
Using Markers/Barriers

- Triggers Snapshots
- Separates preshot-postshot events
- Leads to a consistent execution cut

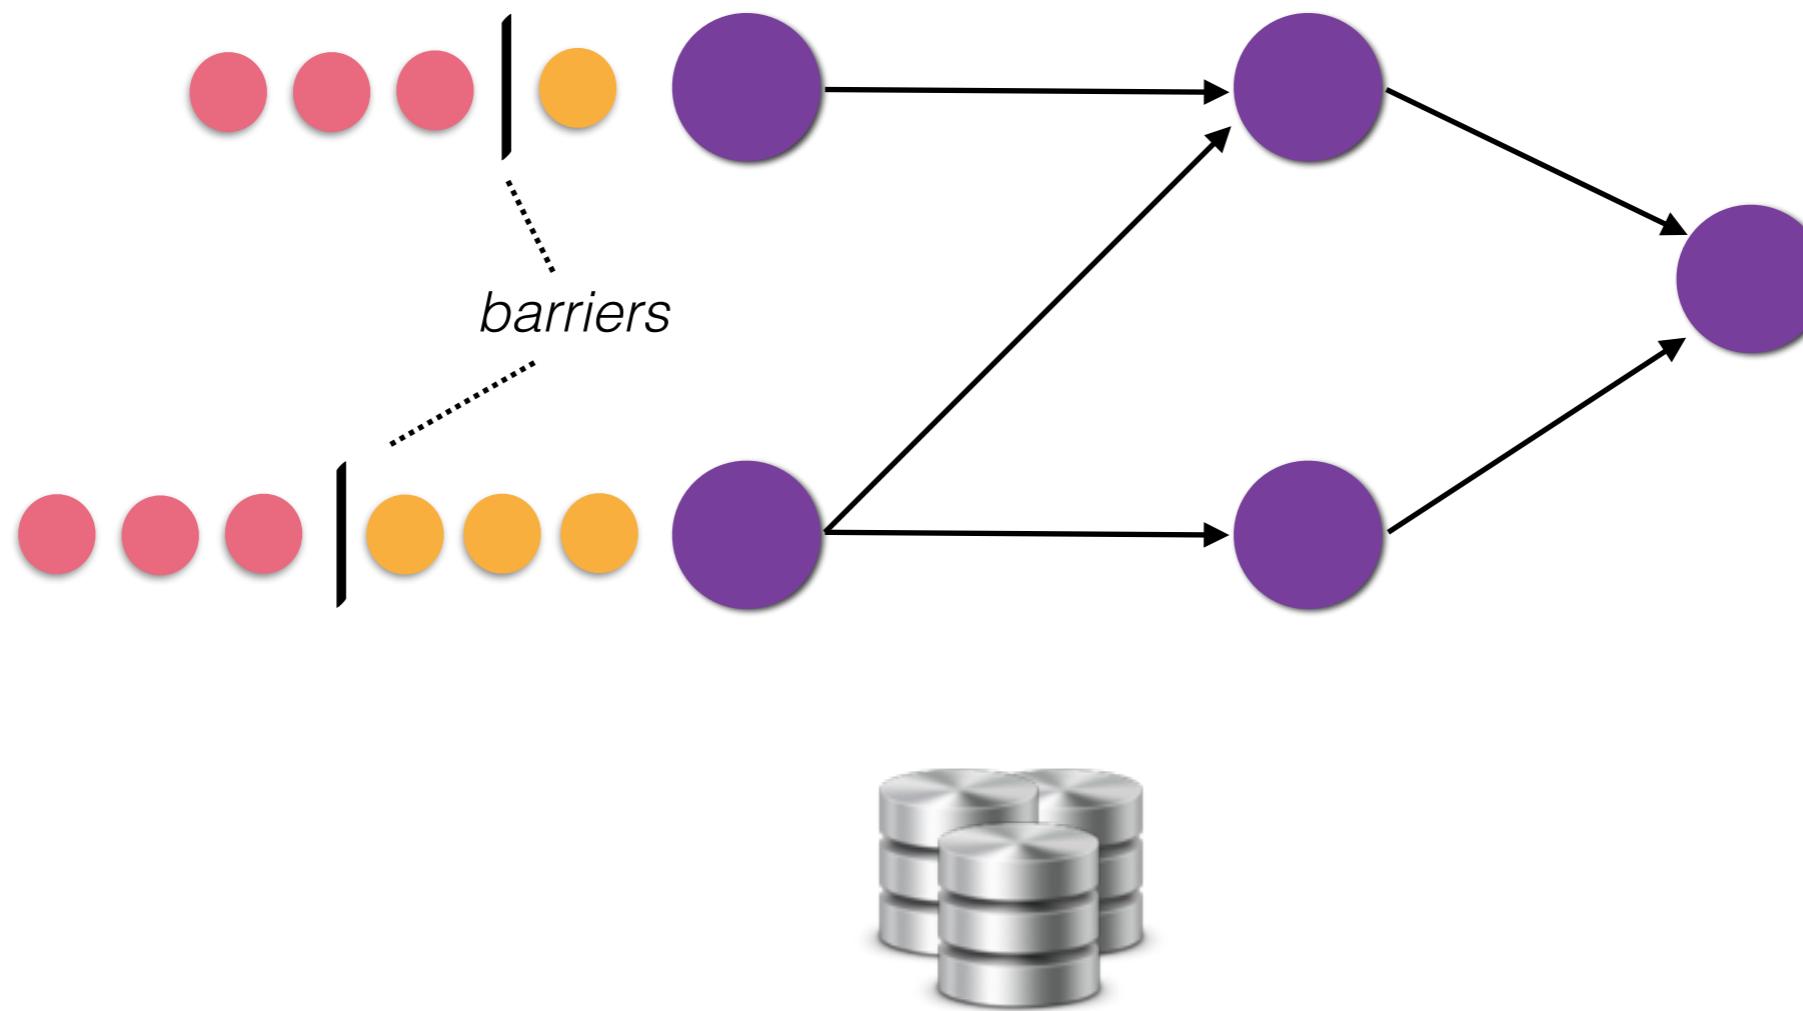
markers propagate the snapshot execution under continuous ingestion



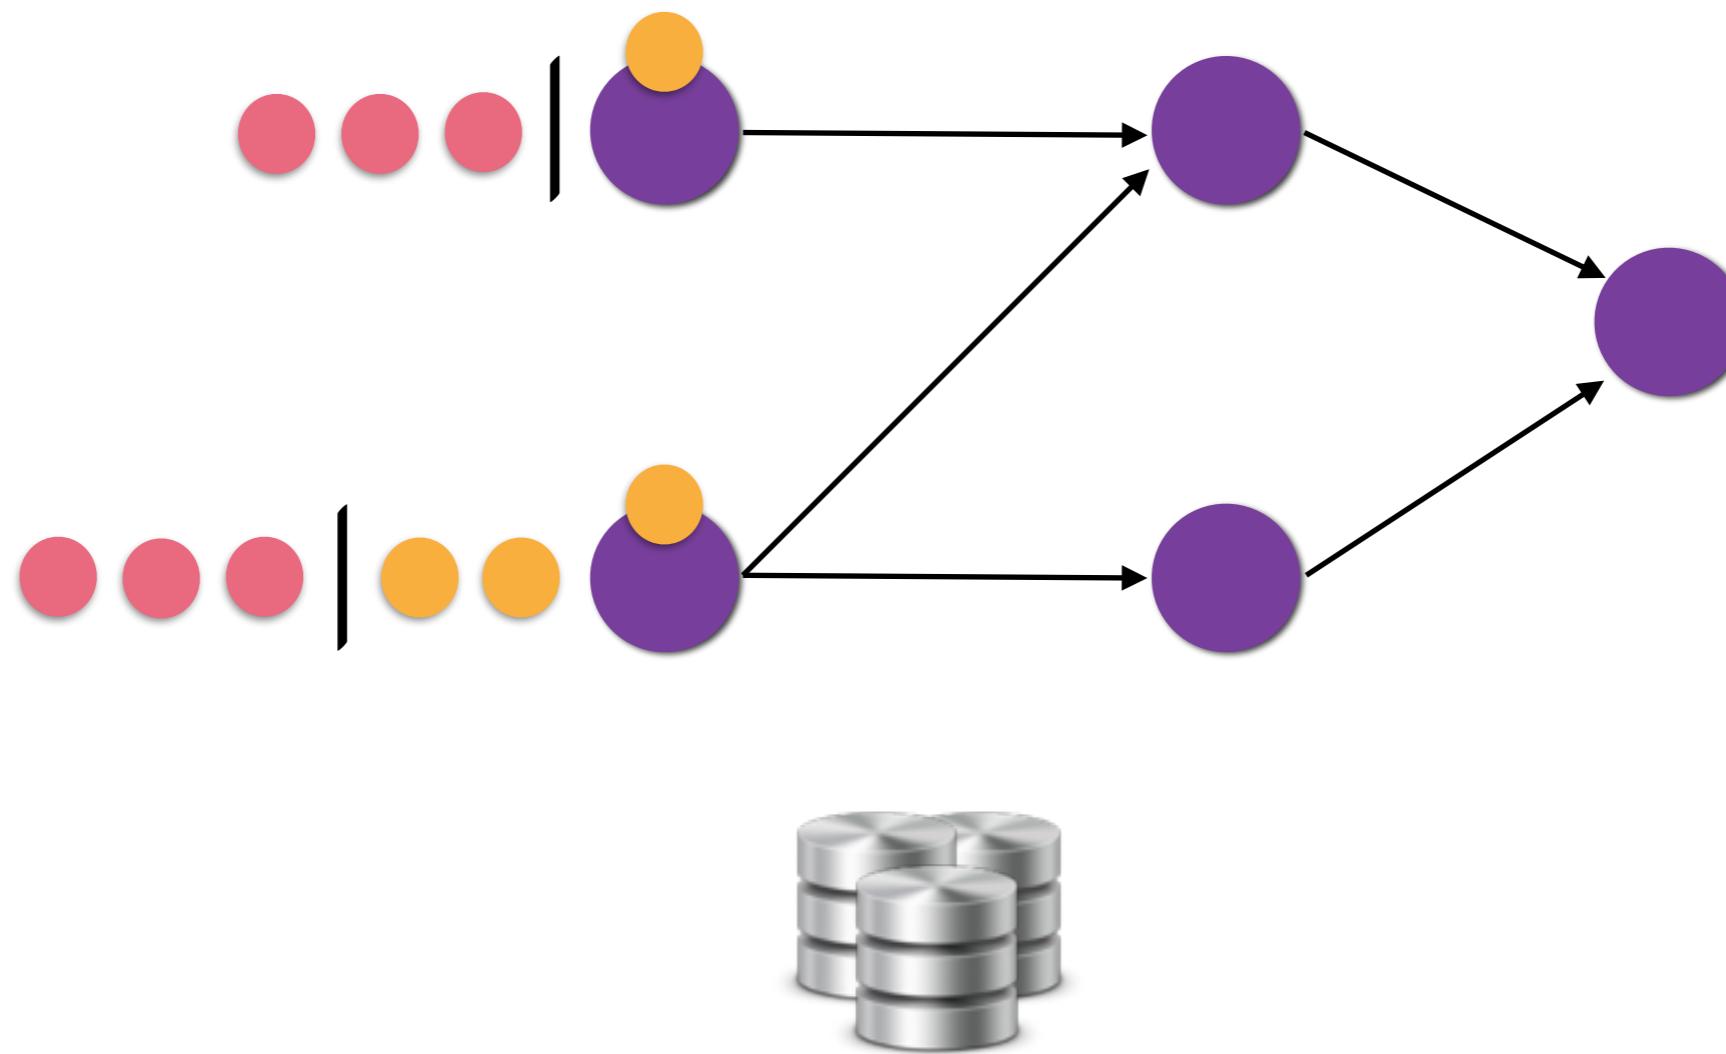
Asynchronous Snapshots



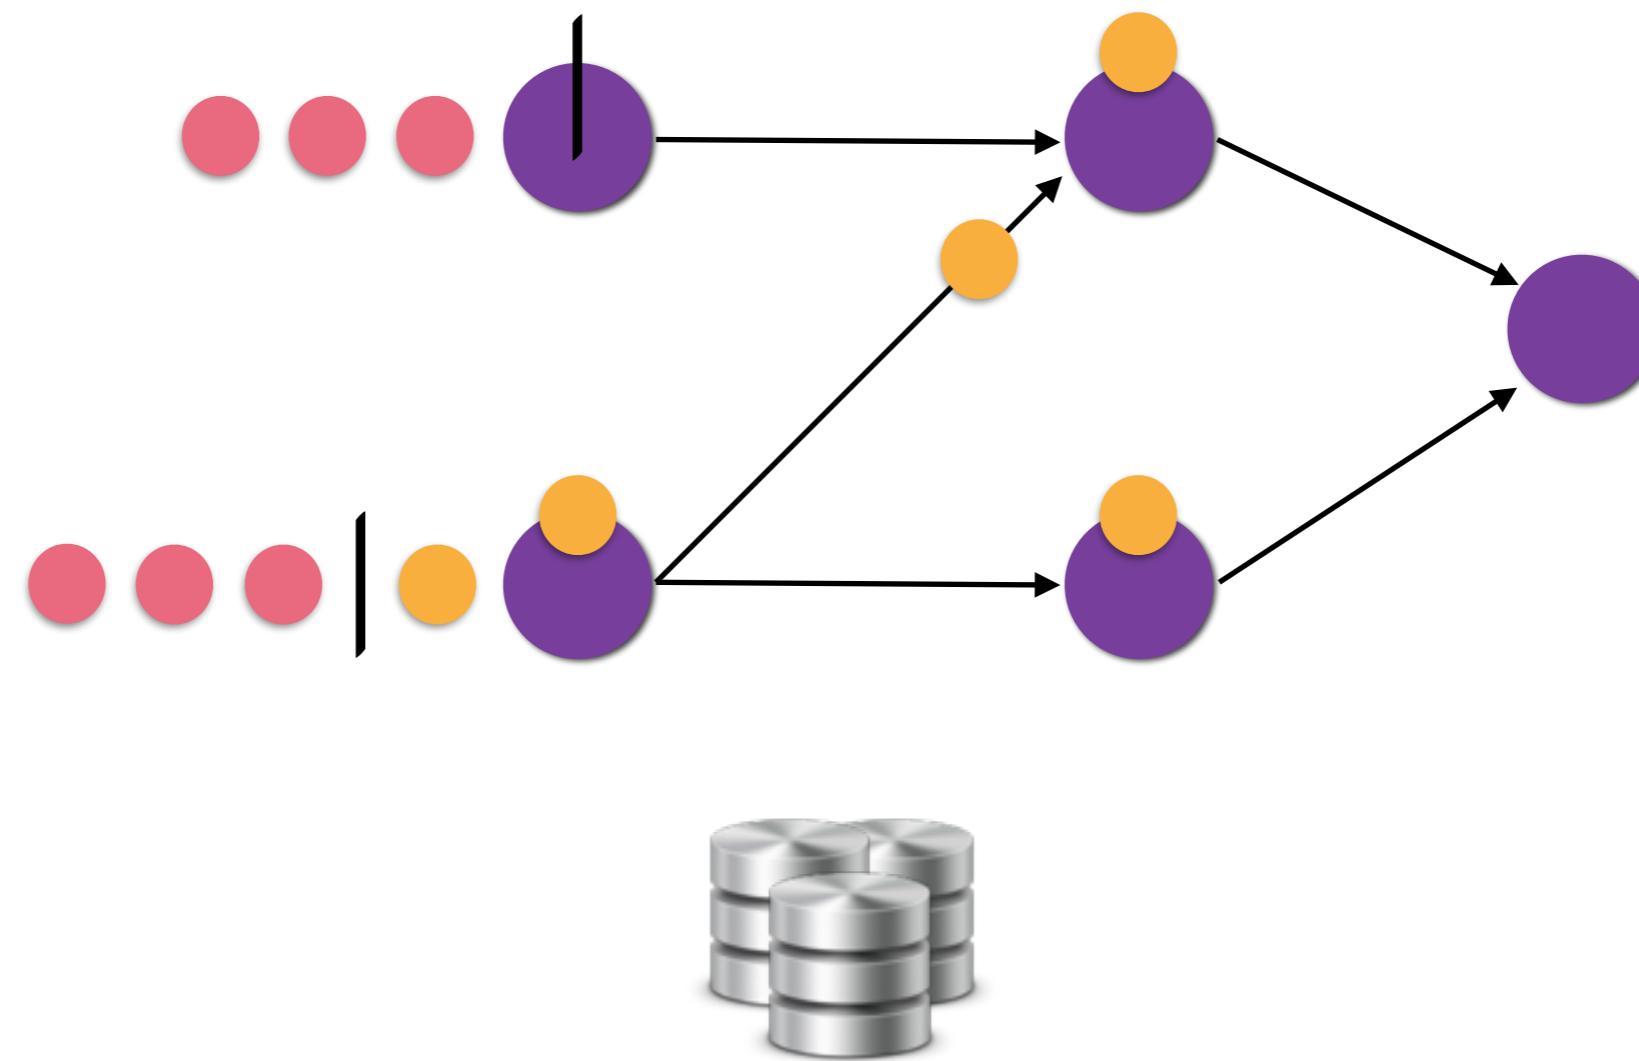
Asynchronous Barrier Snapshotting



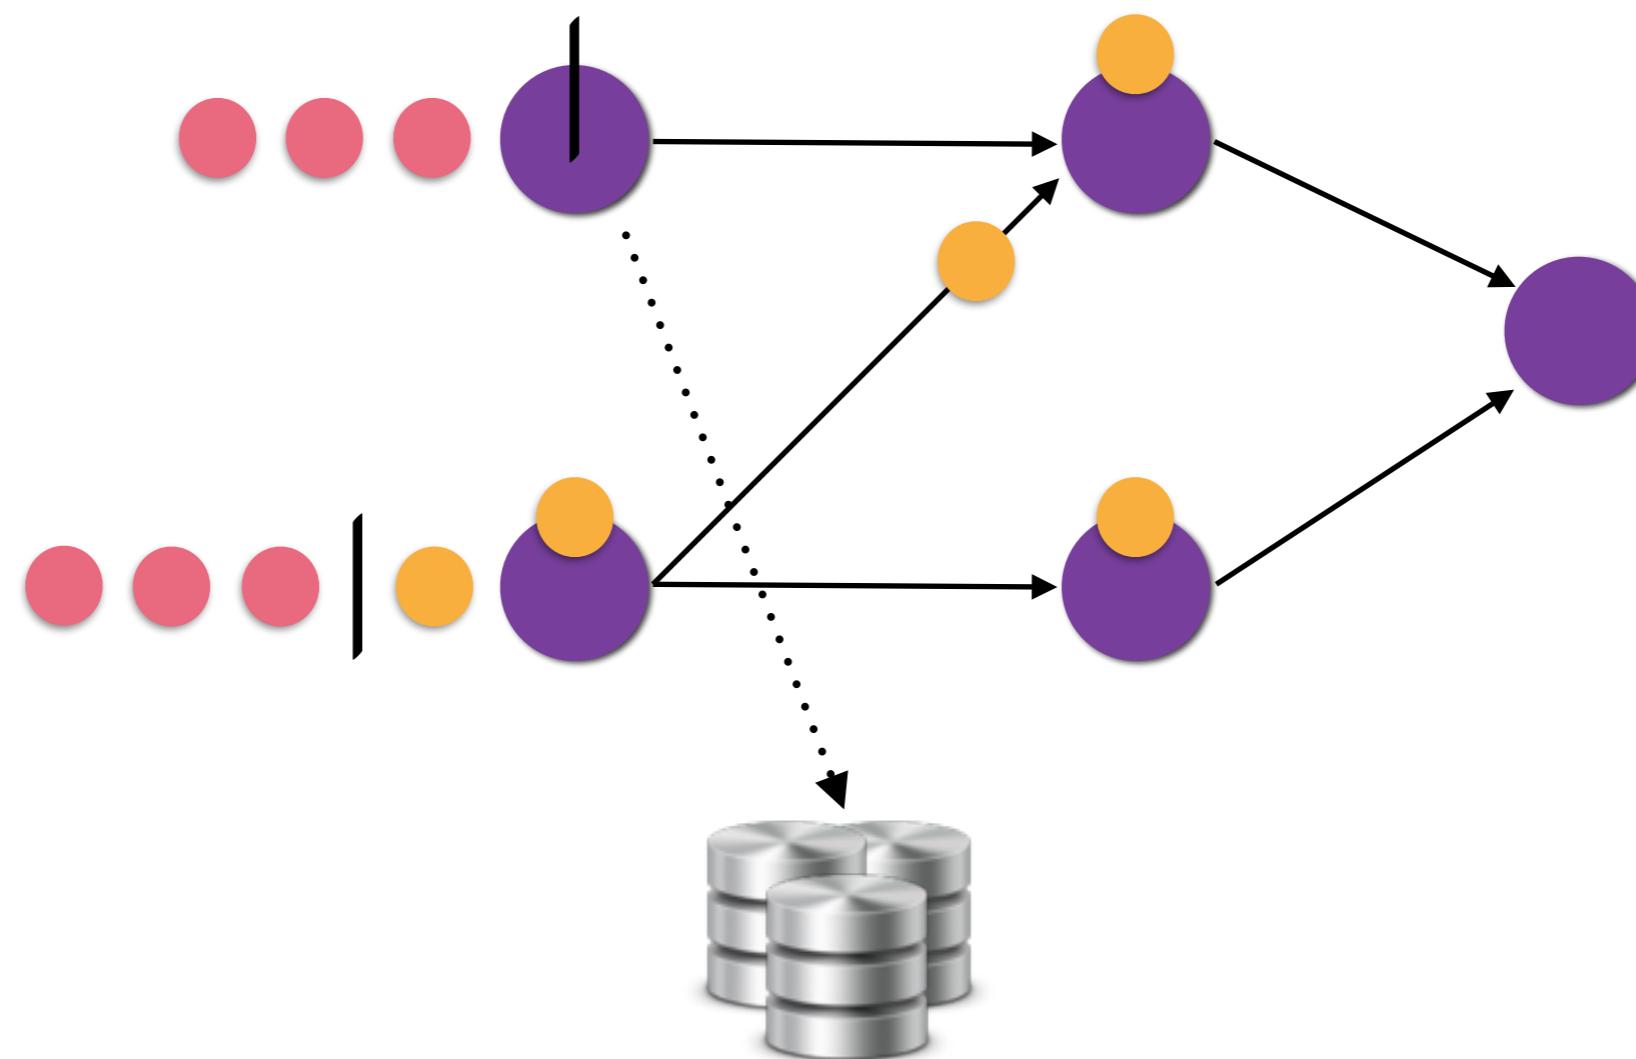
Asynchronous Barrier Snapshotting



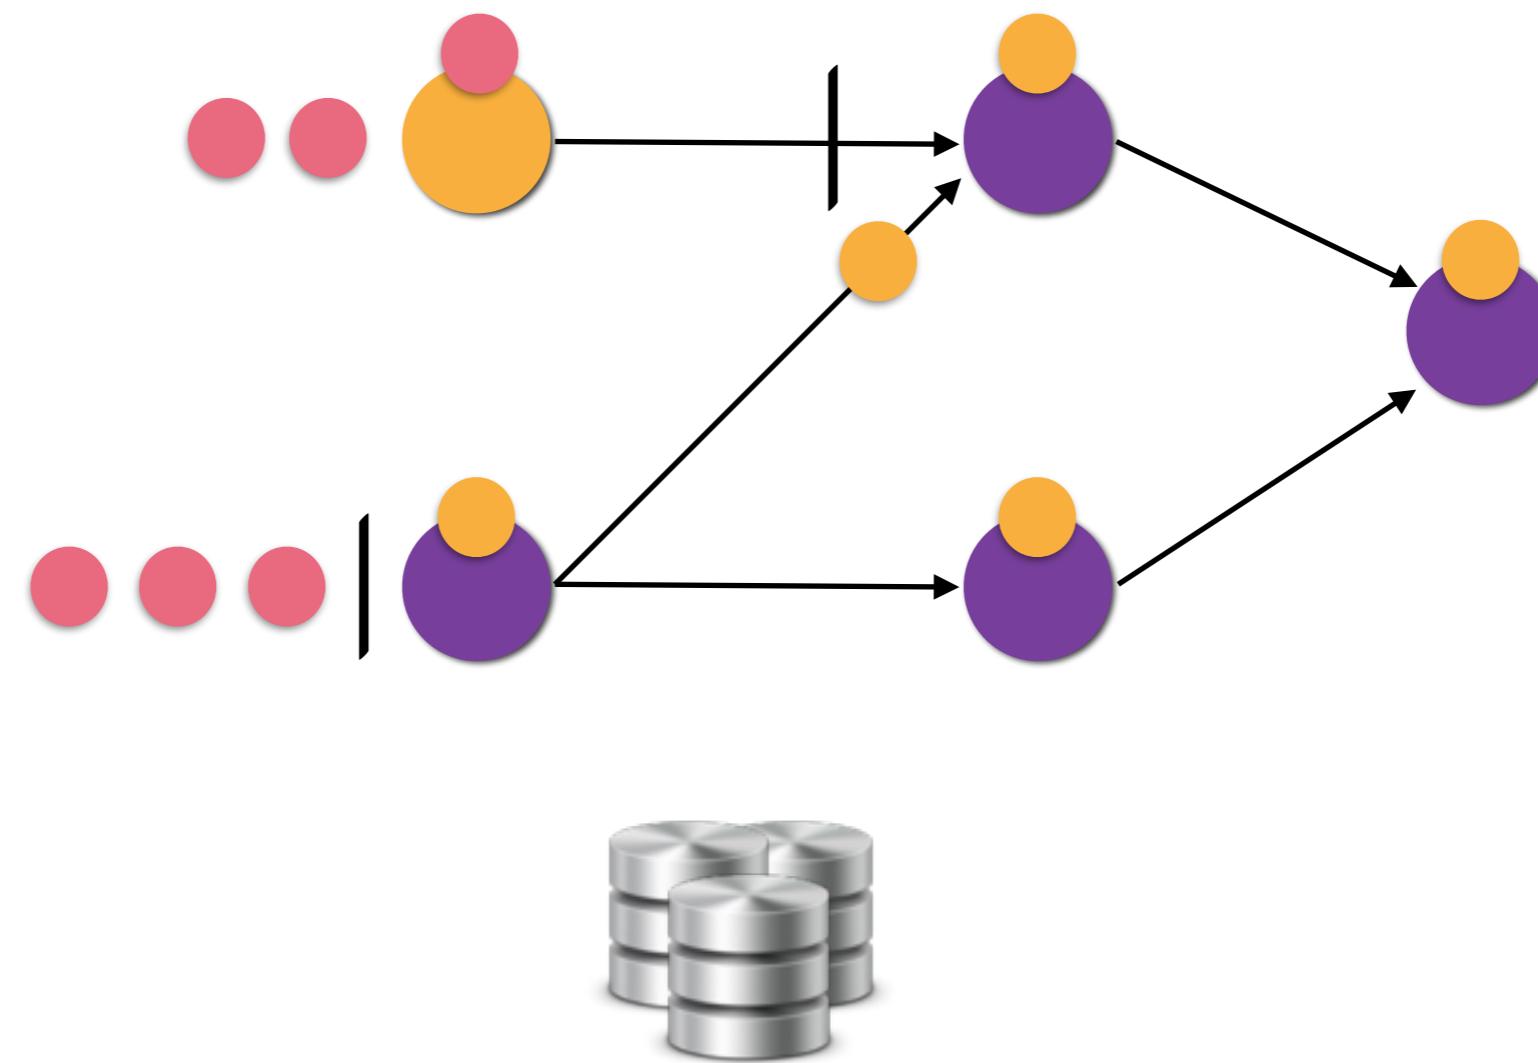
Asynchronous Barrier Snapshotting



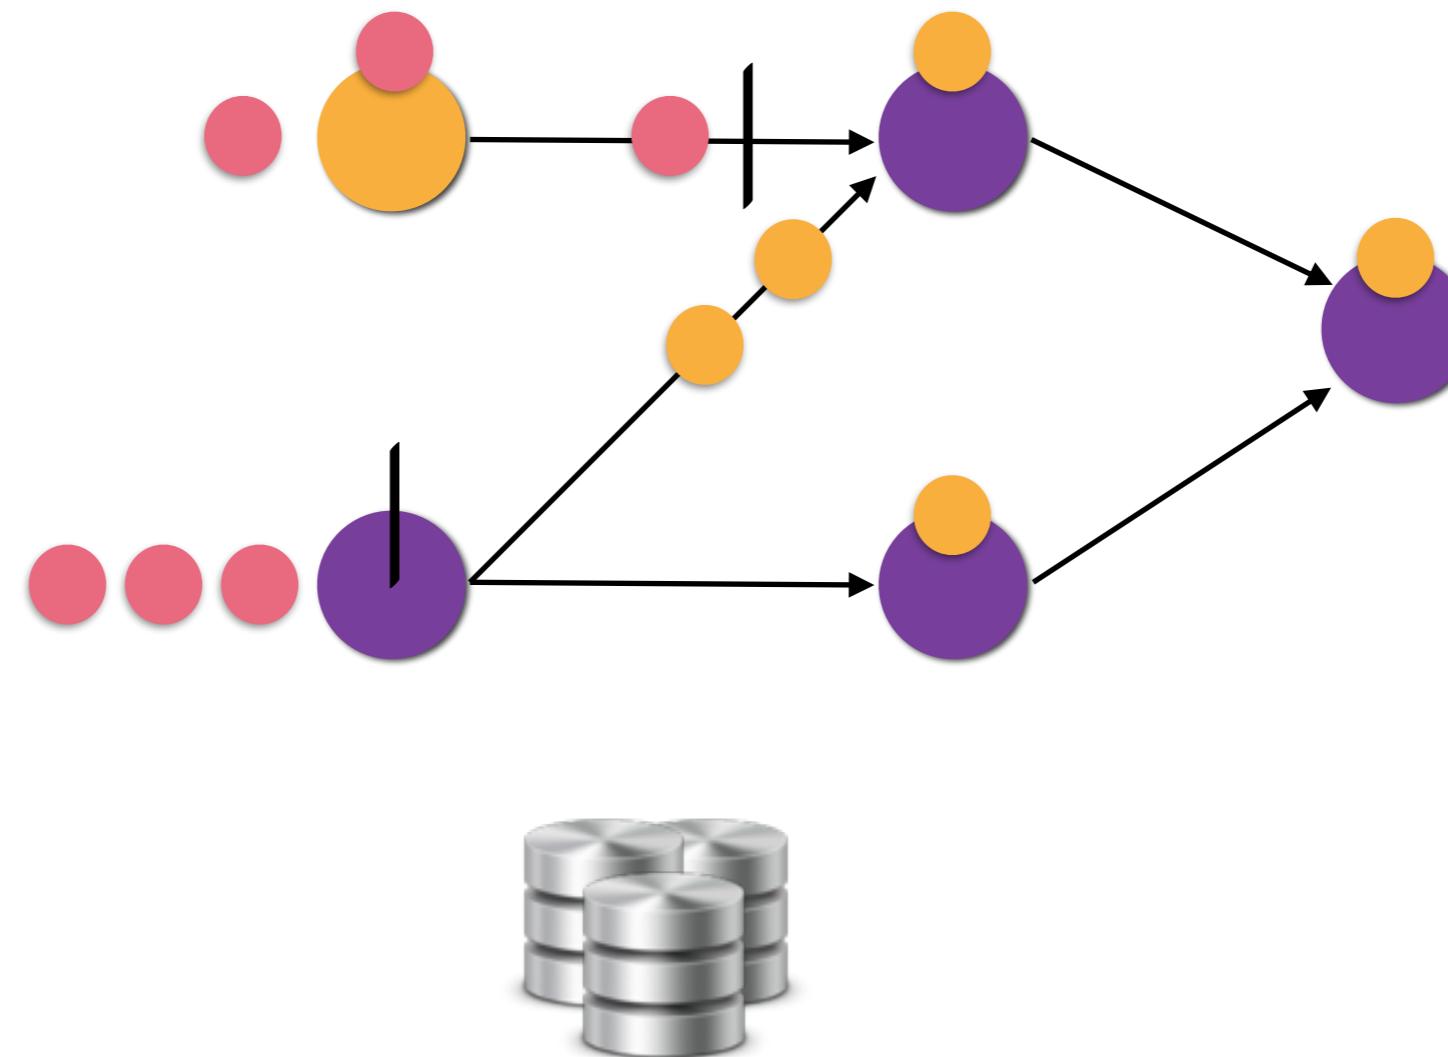
Asynchronous Barrier Snapshotting



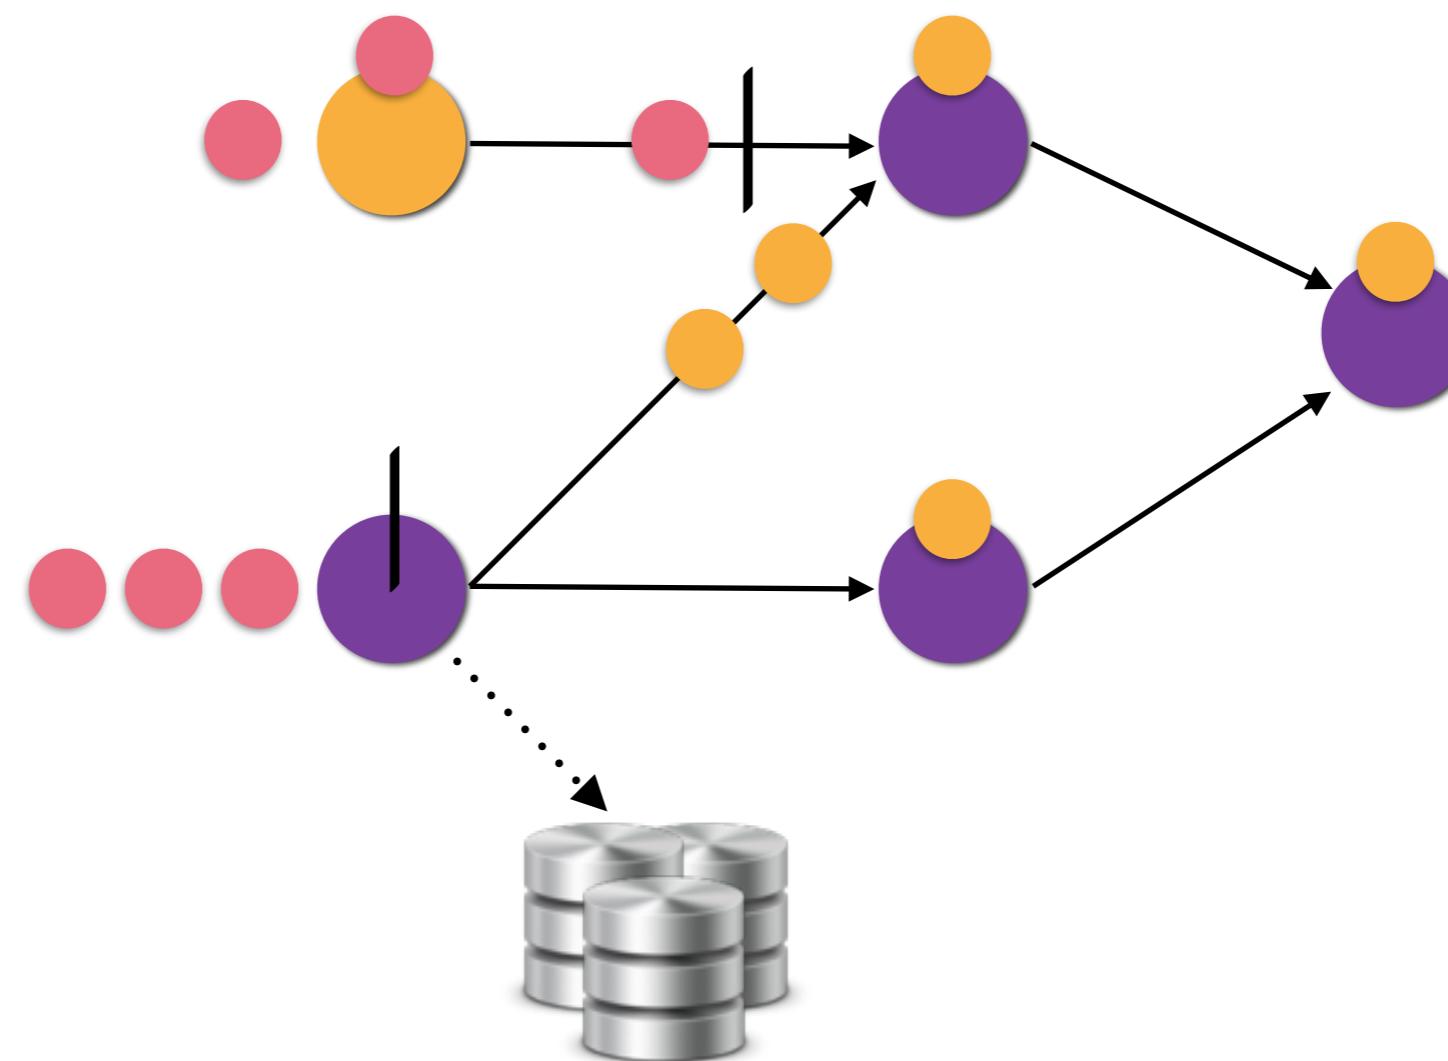
Asynchronous Barrier Snapshotting



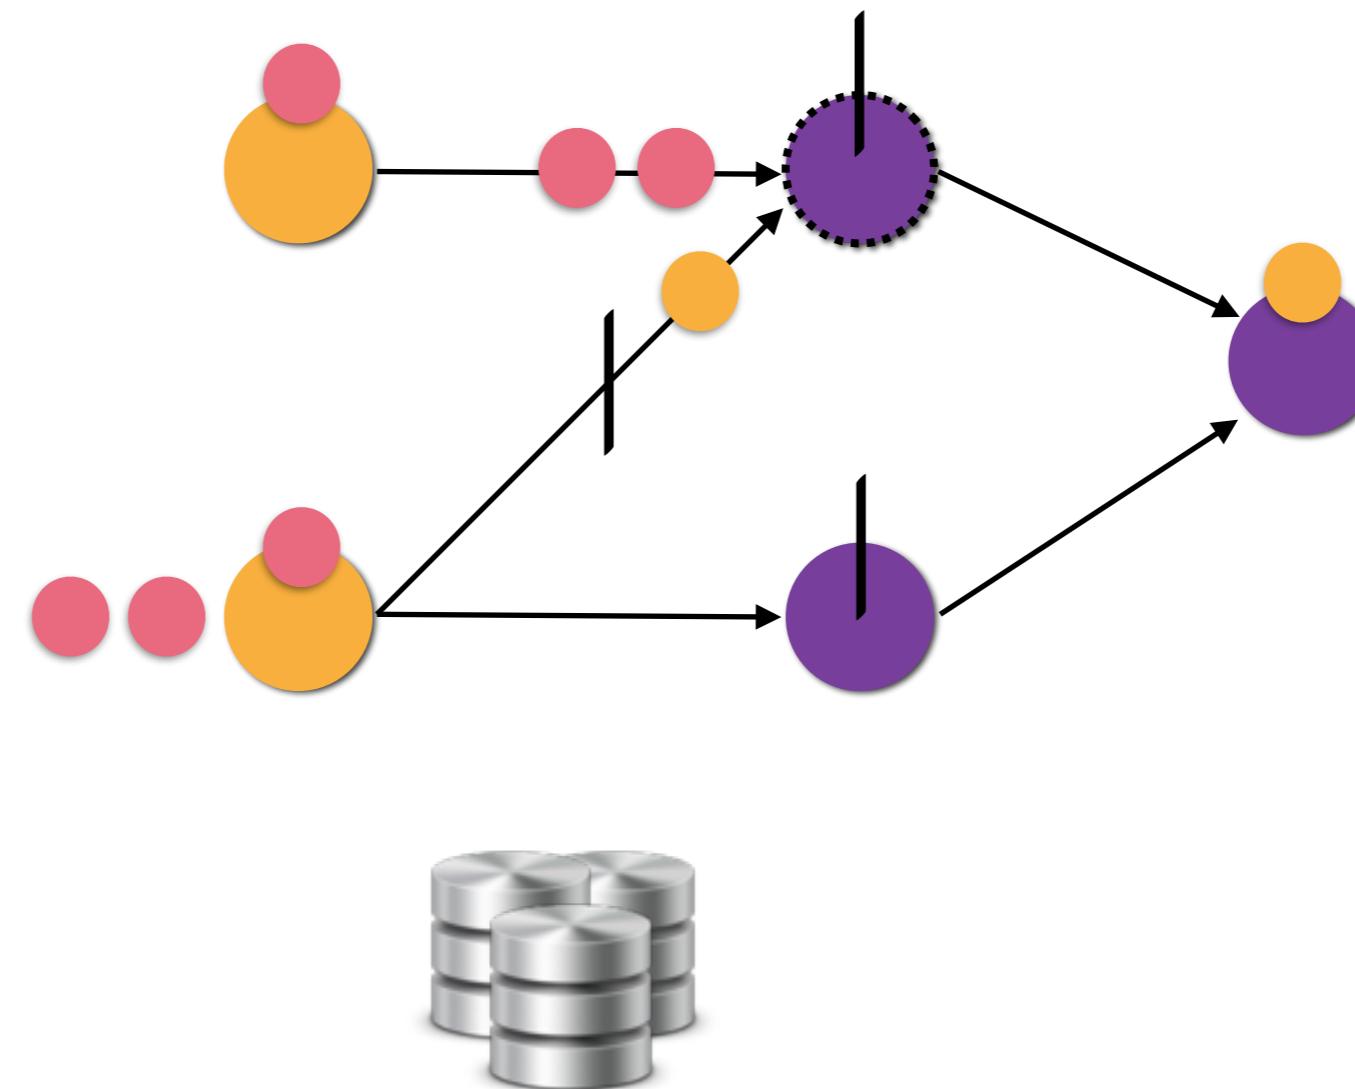
Asynchronous Barrier Snapshotting



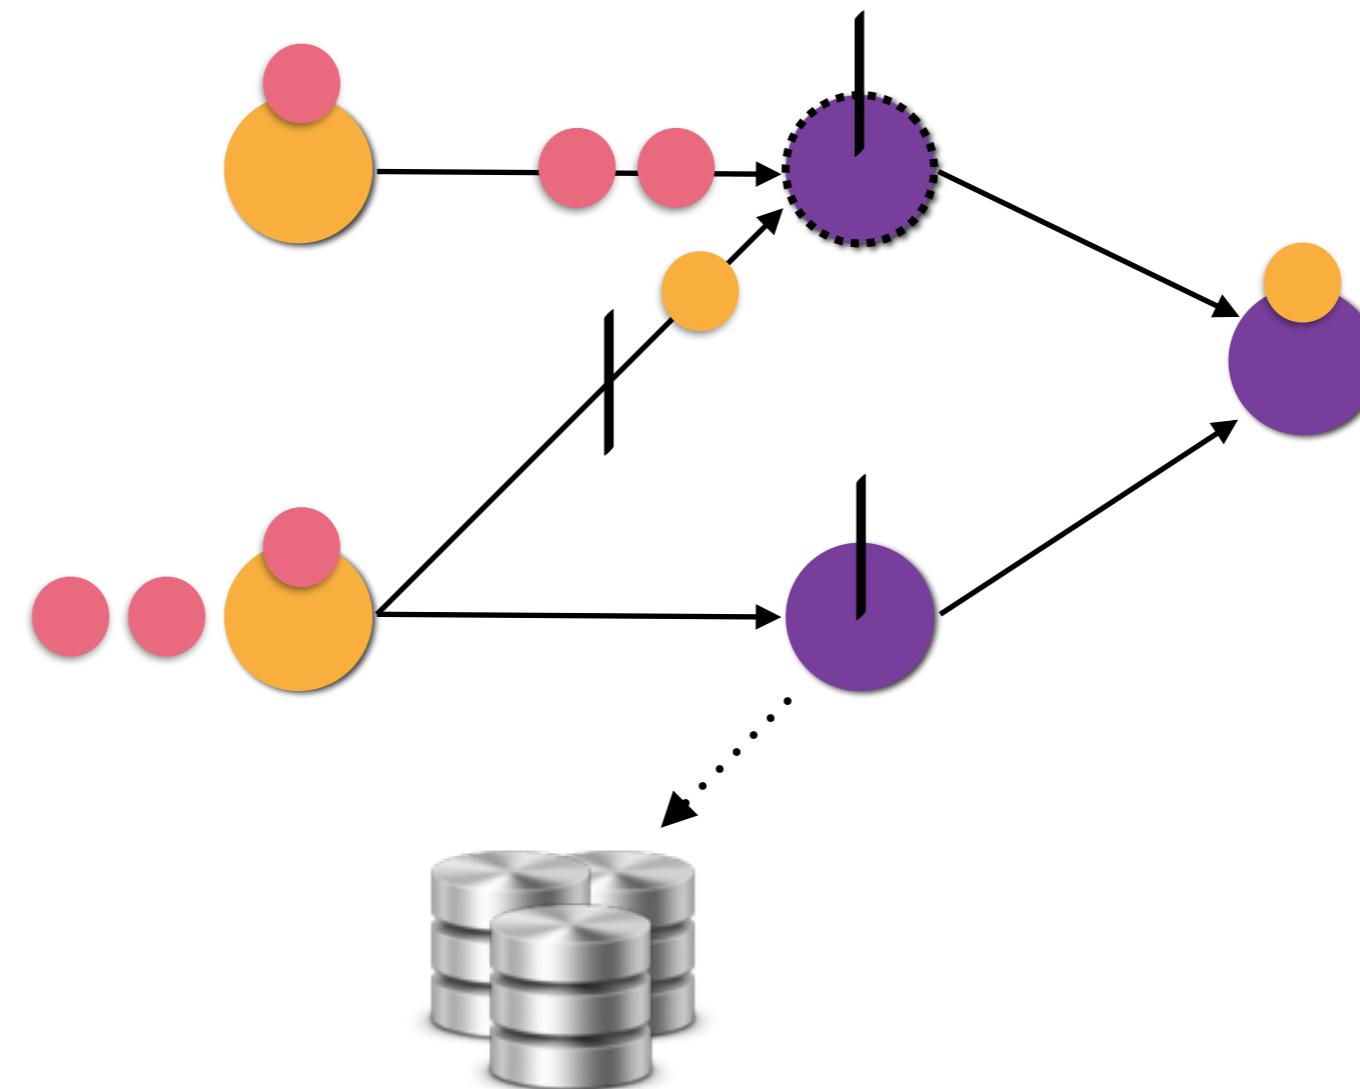
Asynchronous Barrier Snapshotting



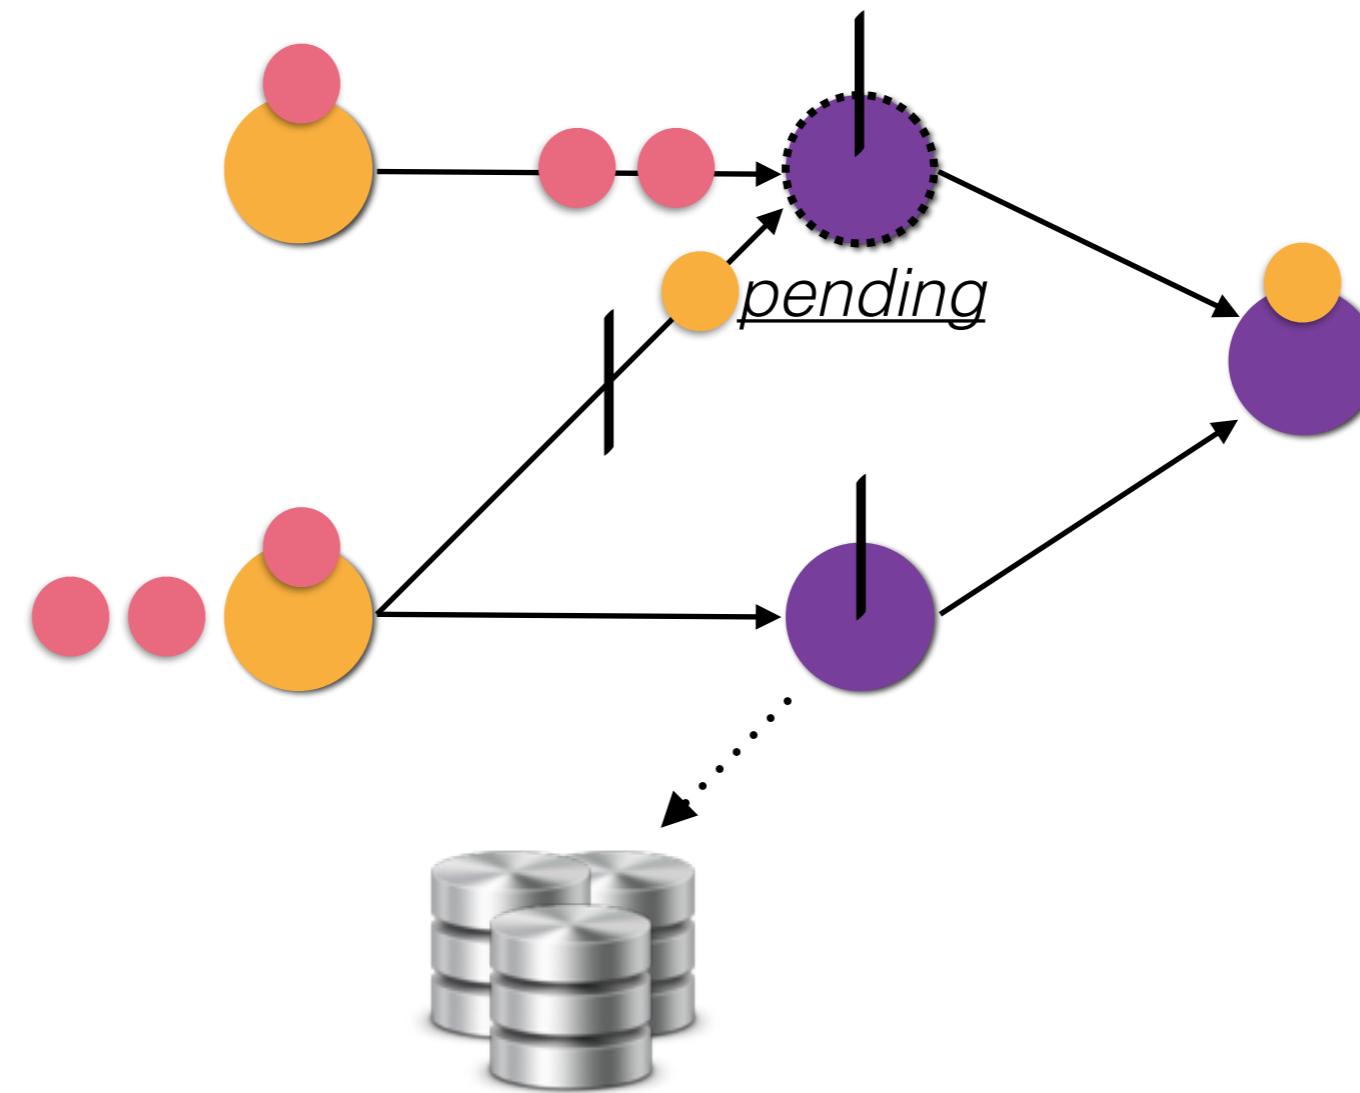
Asynchronous Barrier Snapshotting



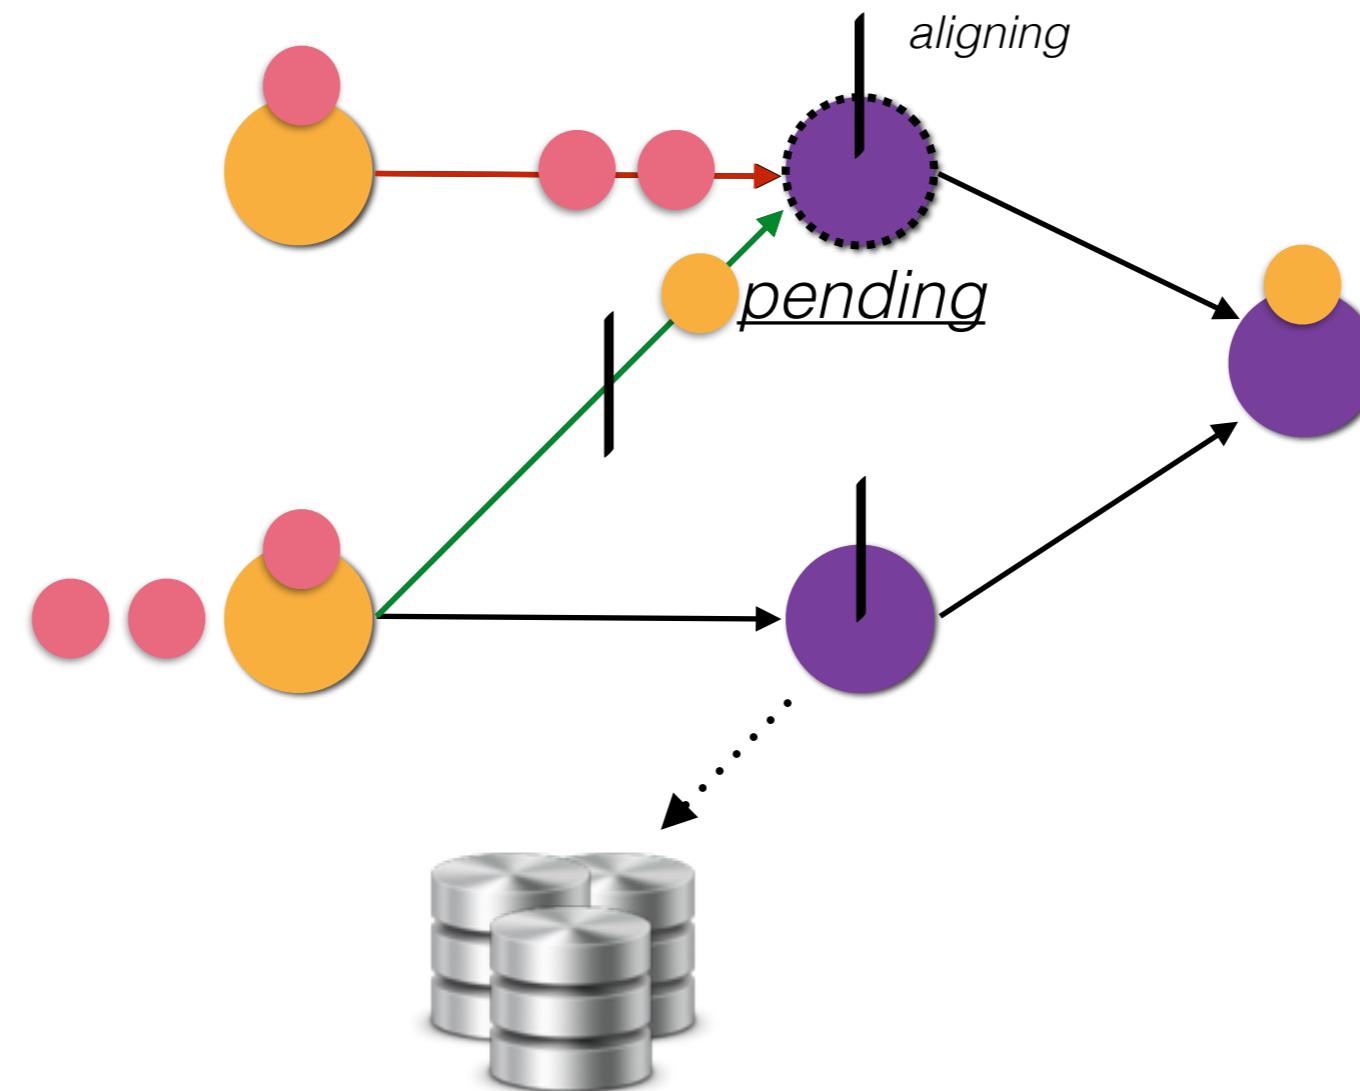
Asynchronous Barrier Snapshotting



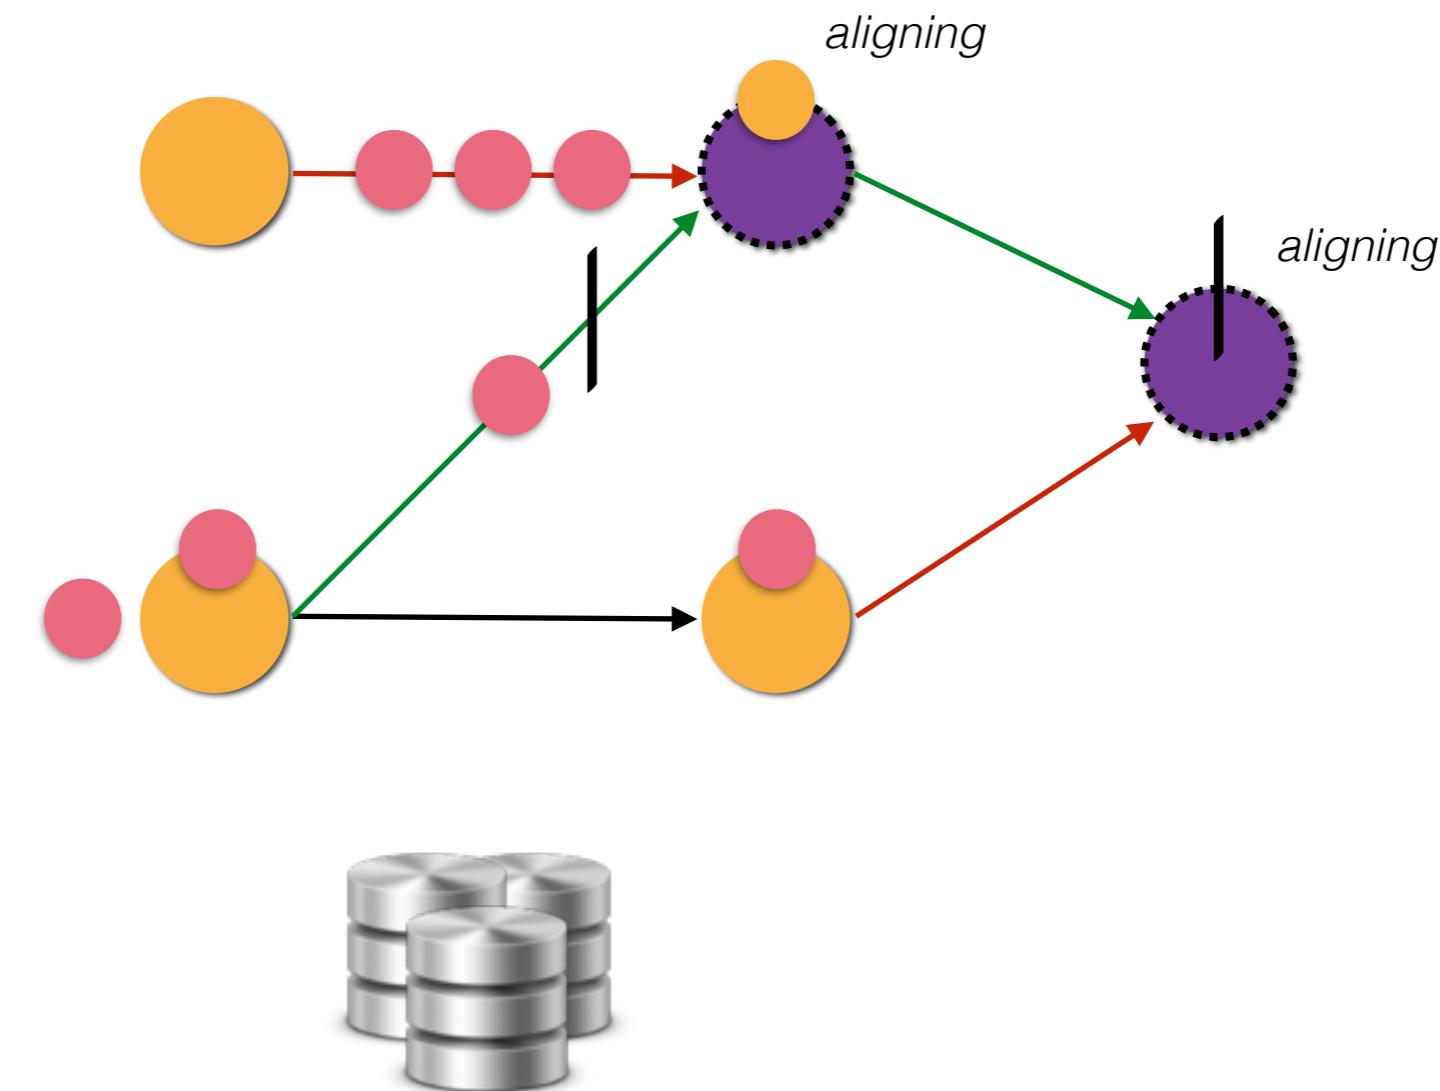
Asynchronous Barrier Snapshotting



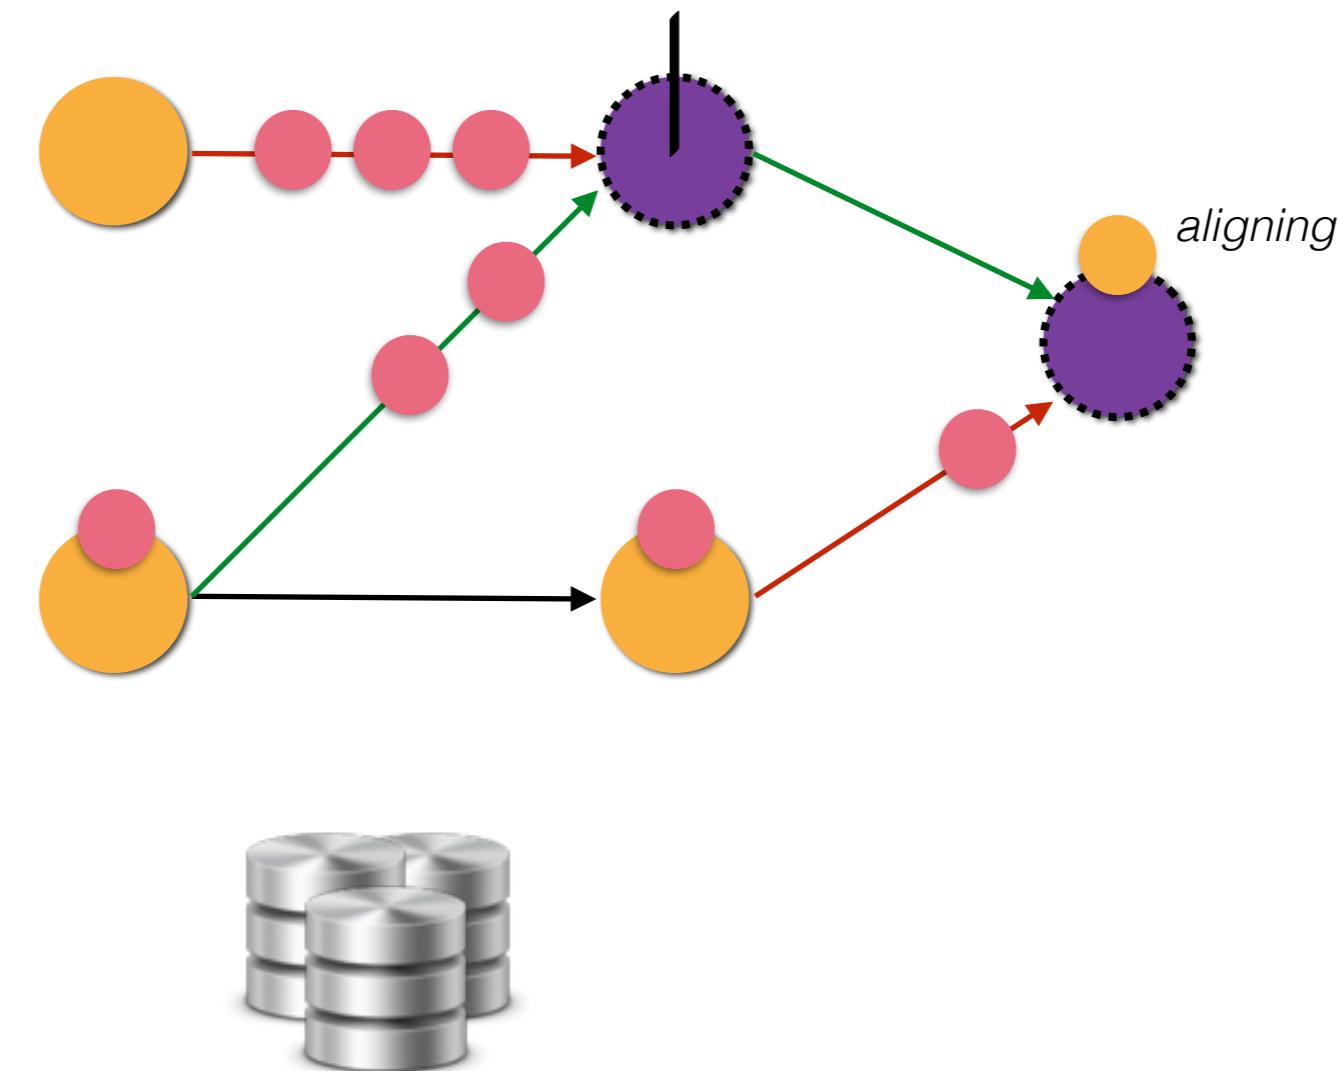
Asynchronous Barrier Snapshotting



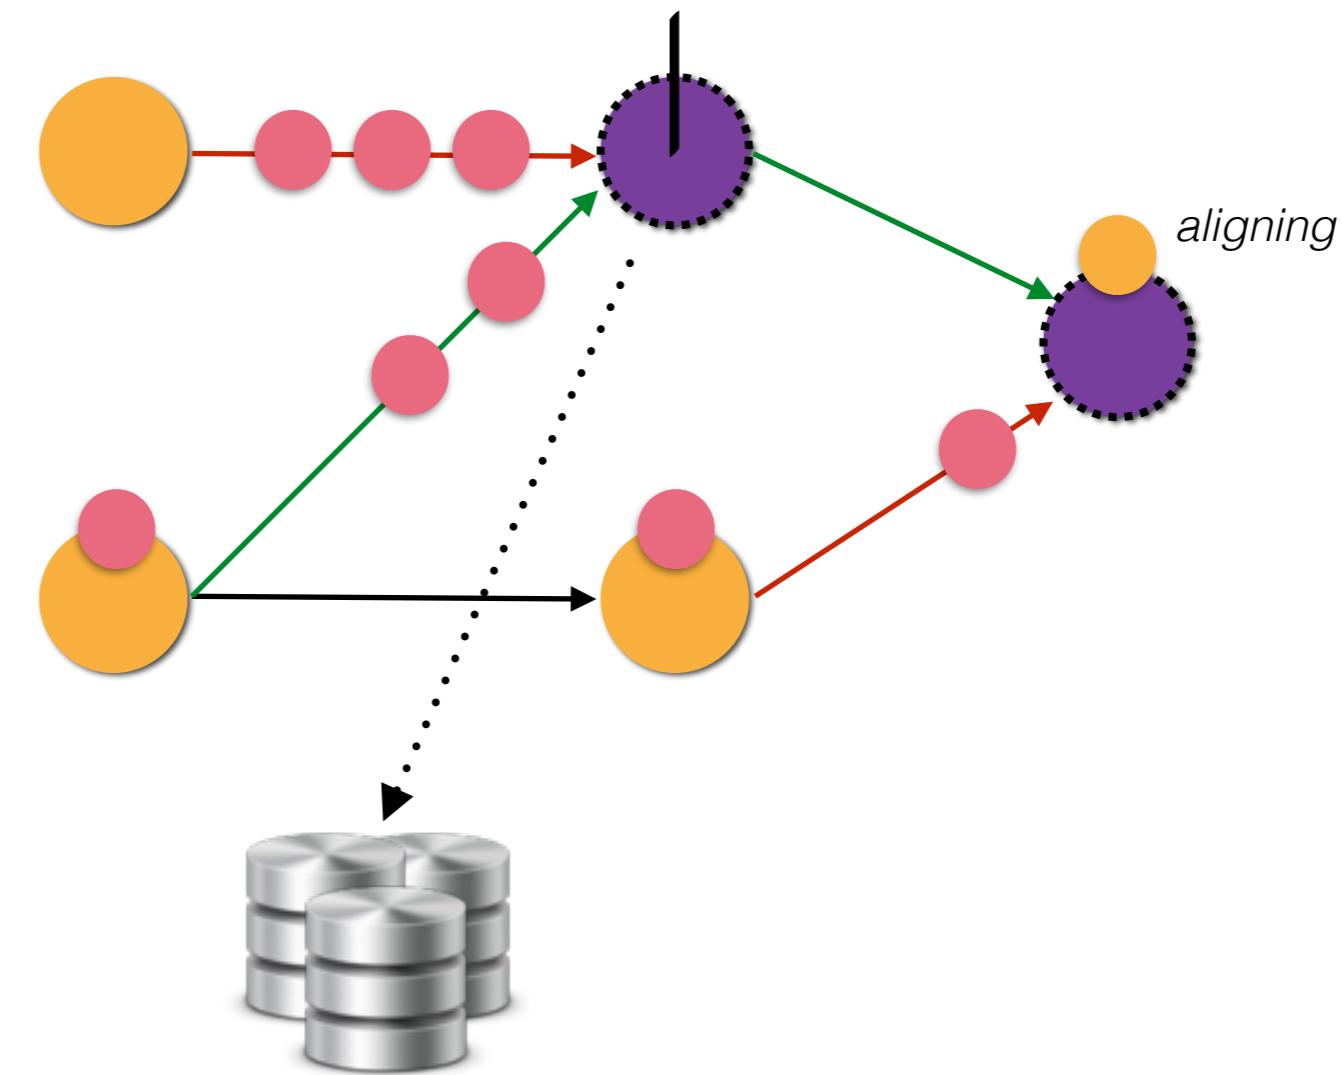
Asynchronous Barrier Snapshotting



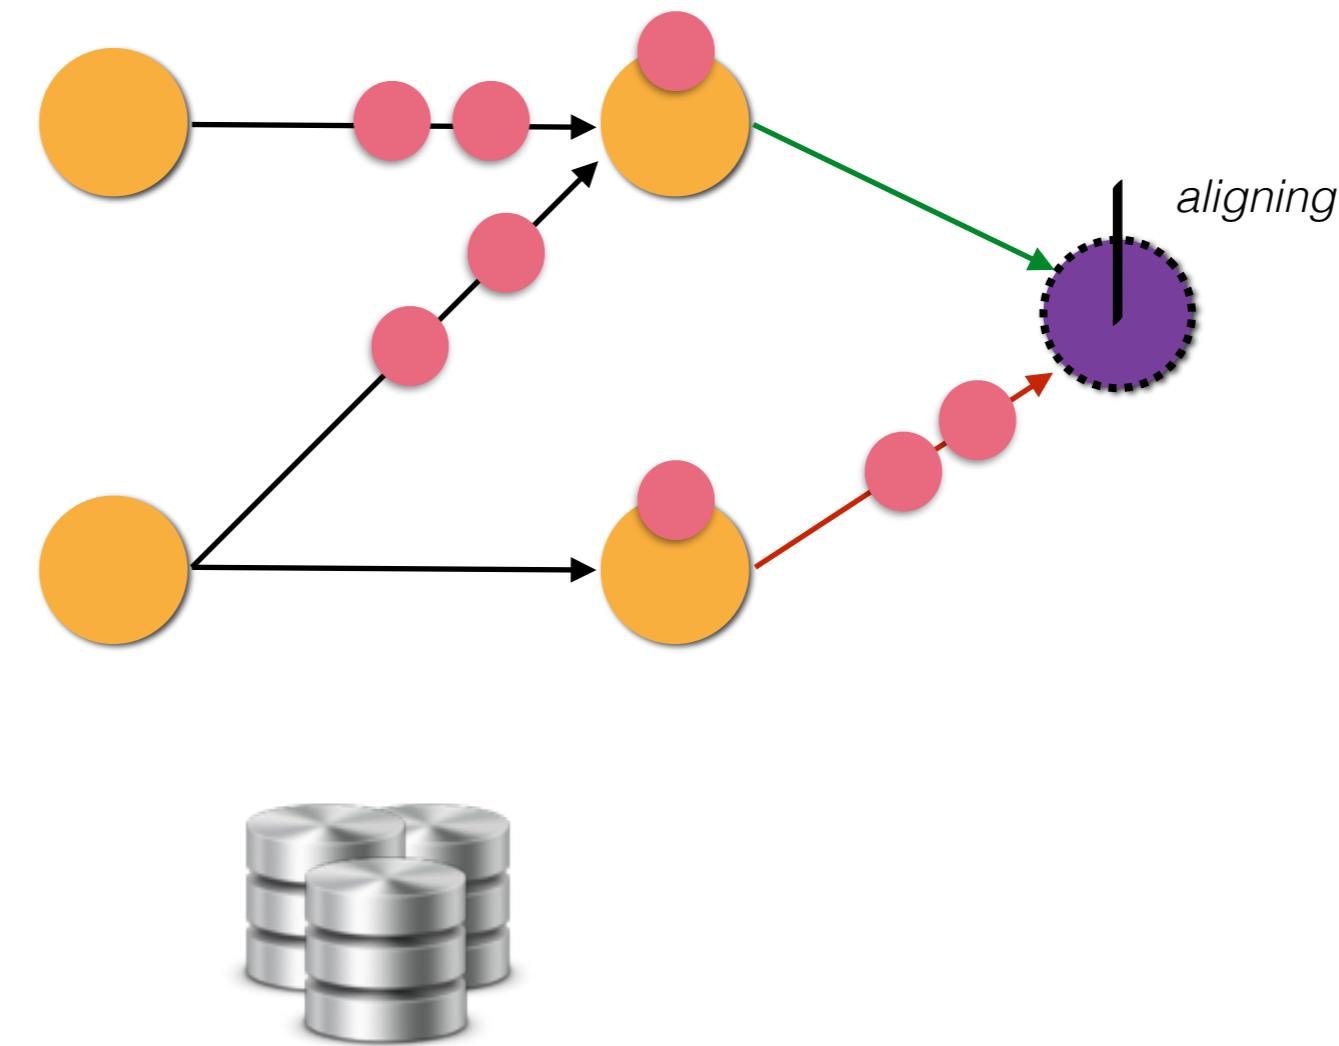
Asynchronous Barrier Snapshotting



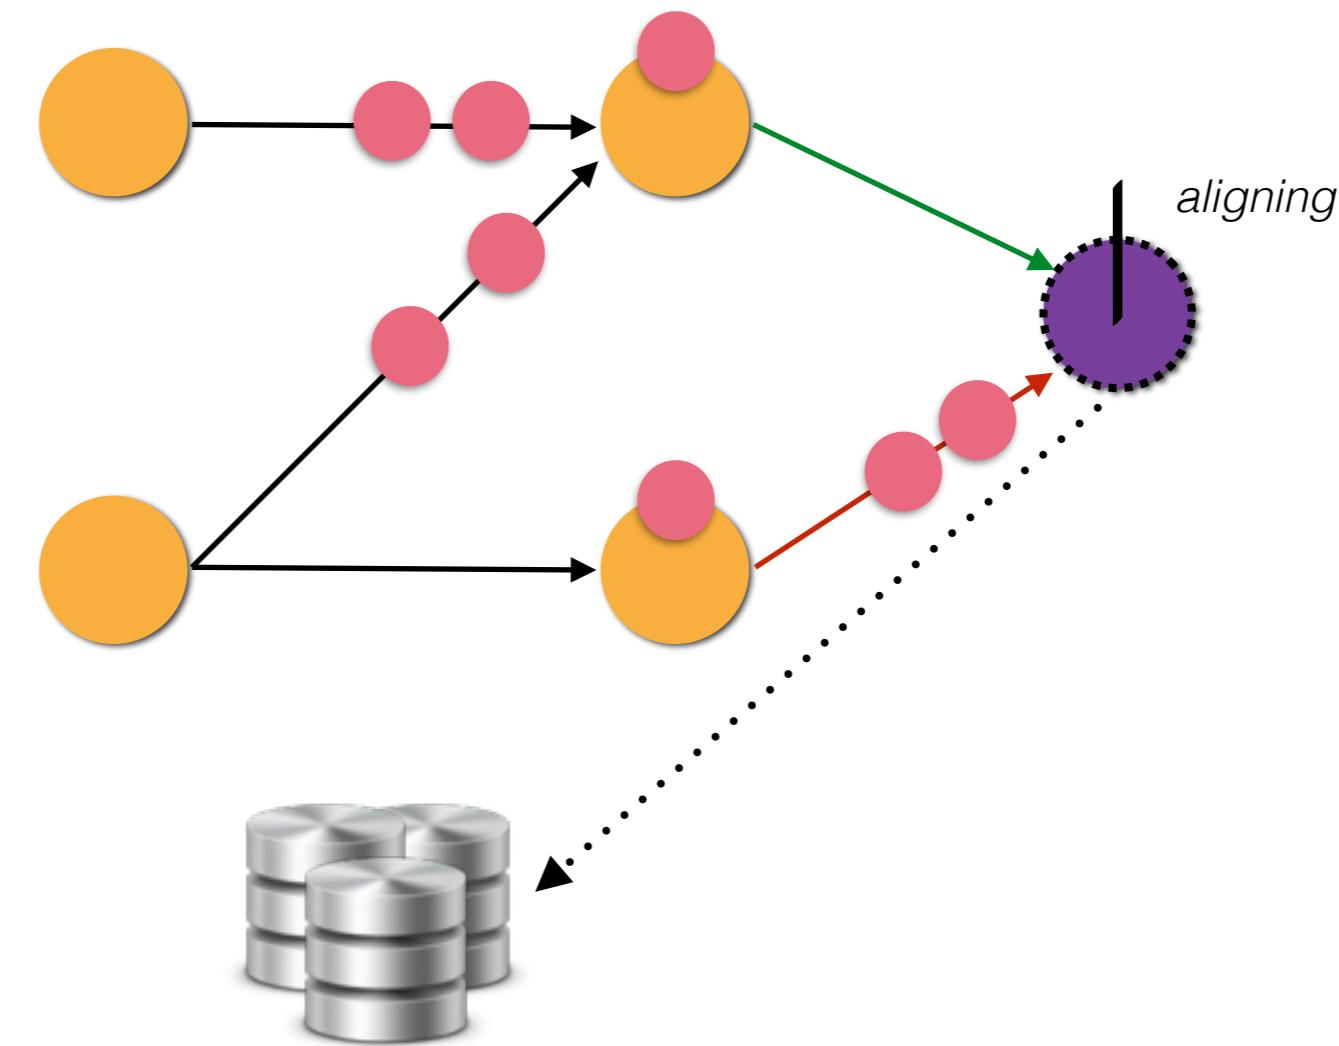
Asynchronous Barrier Snapshotting



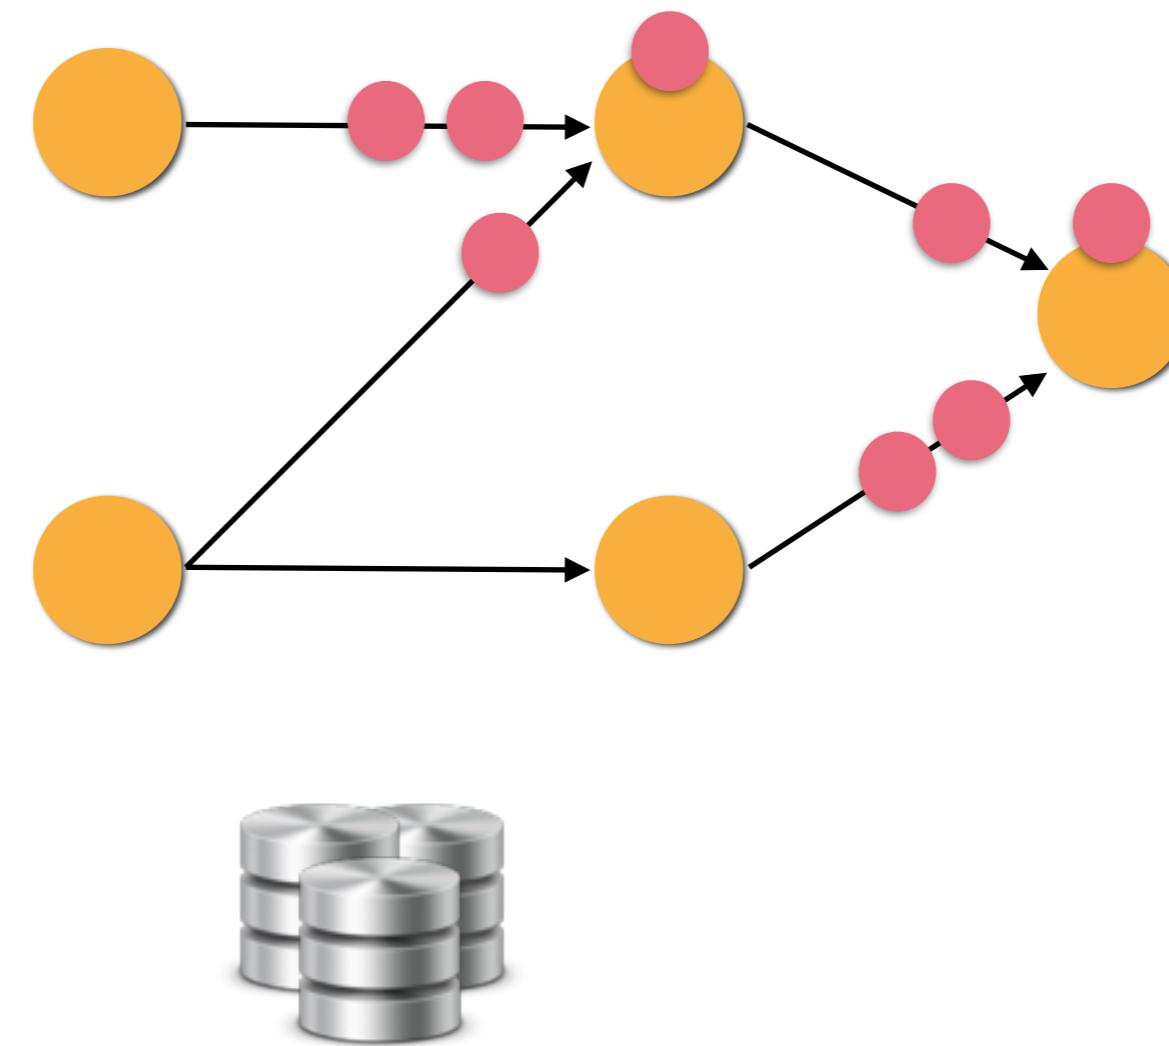
Asynchronous Barrier Snapshotting



Asynchronous Barrier Snapshotting



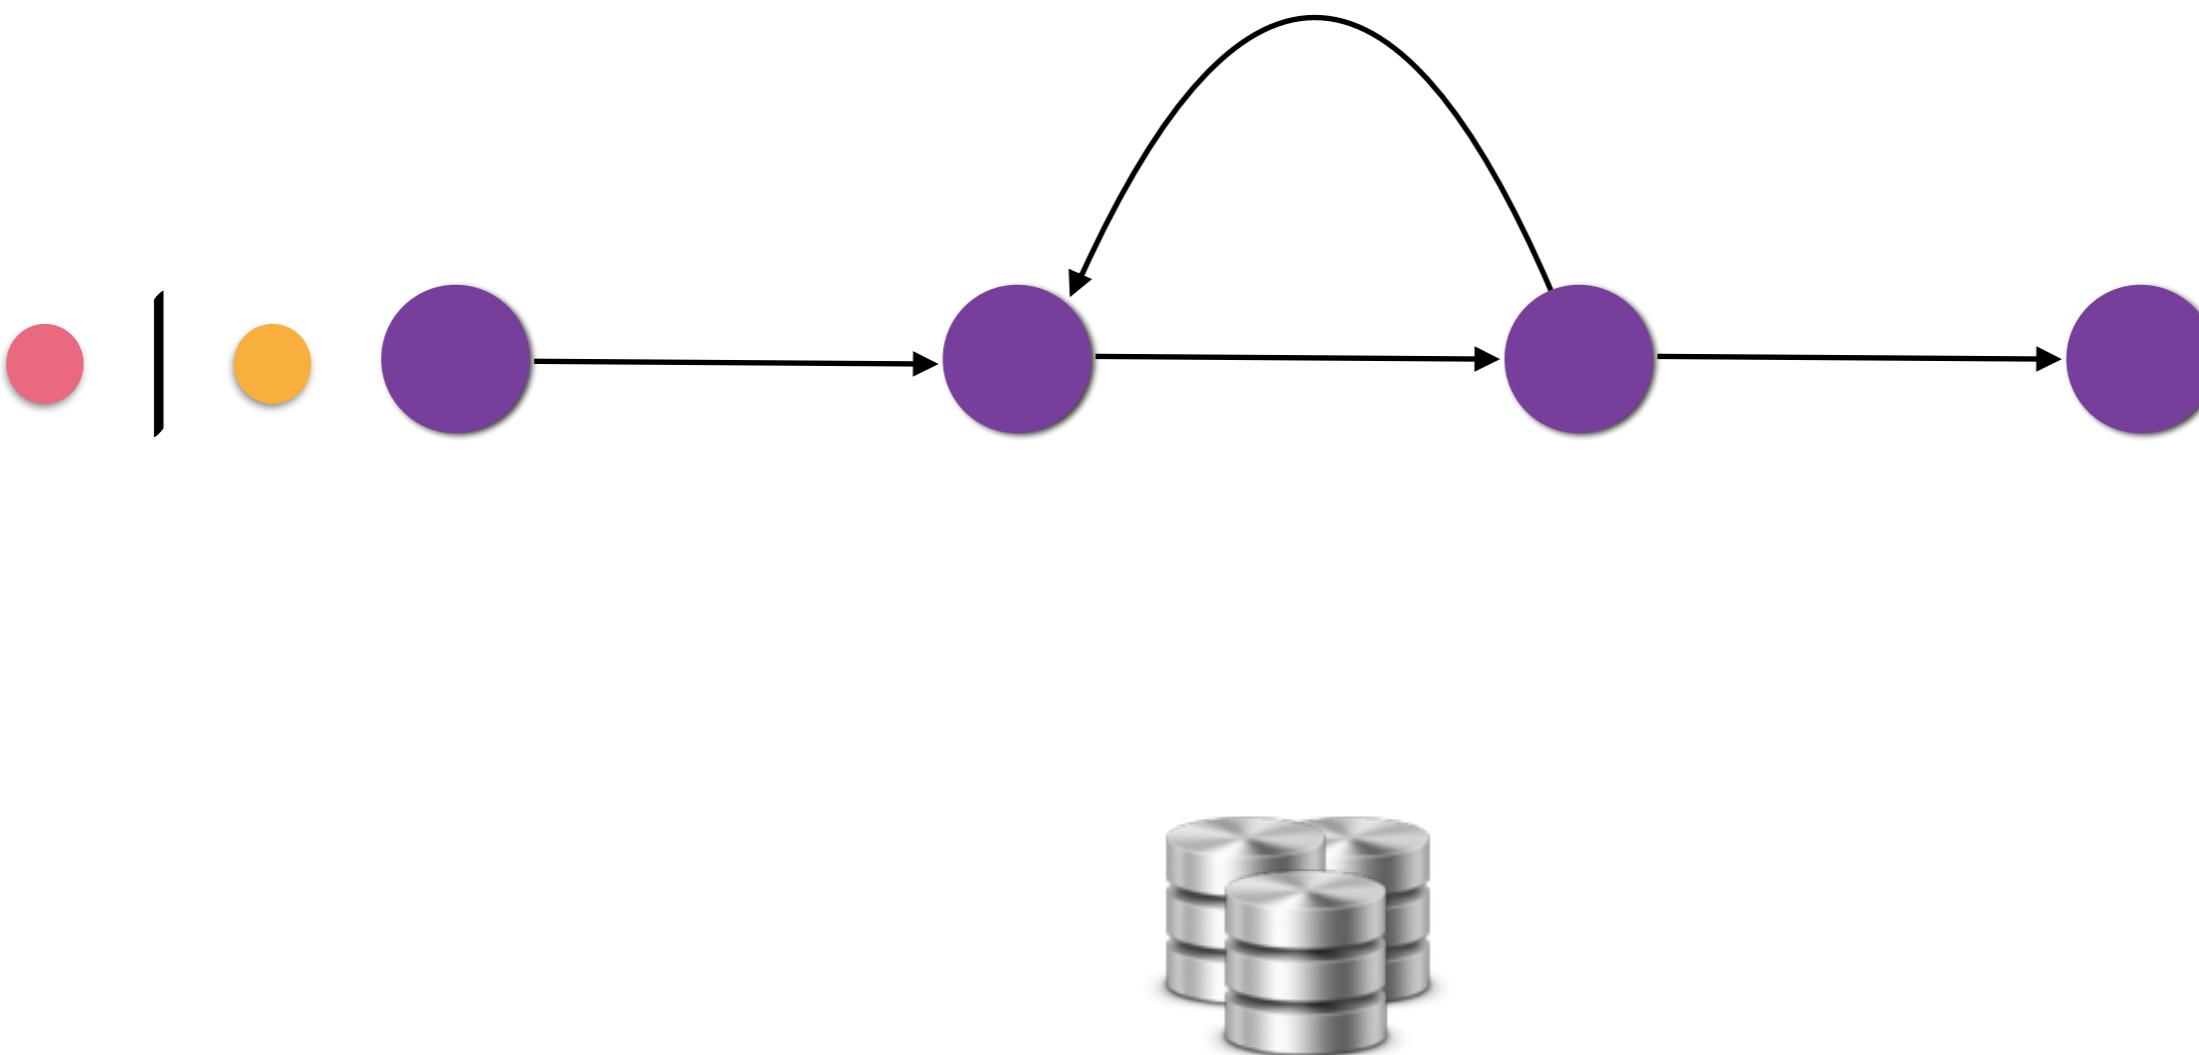
Asynchronous Barrier Snapshotting



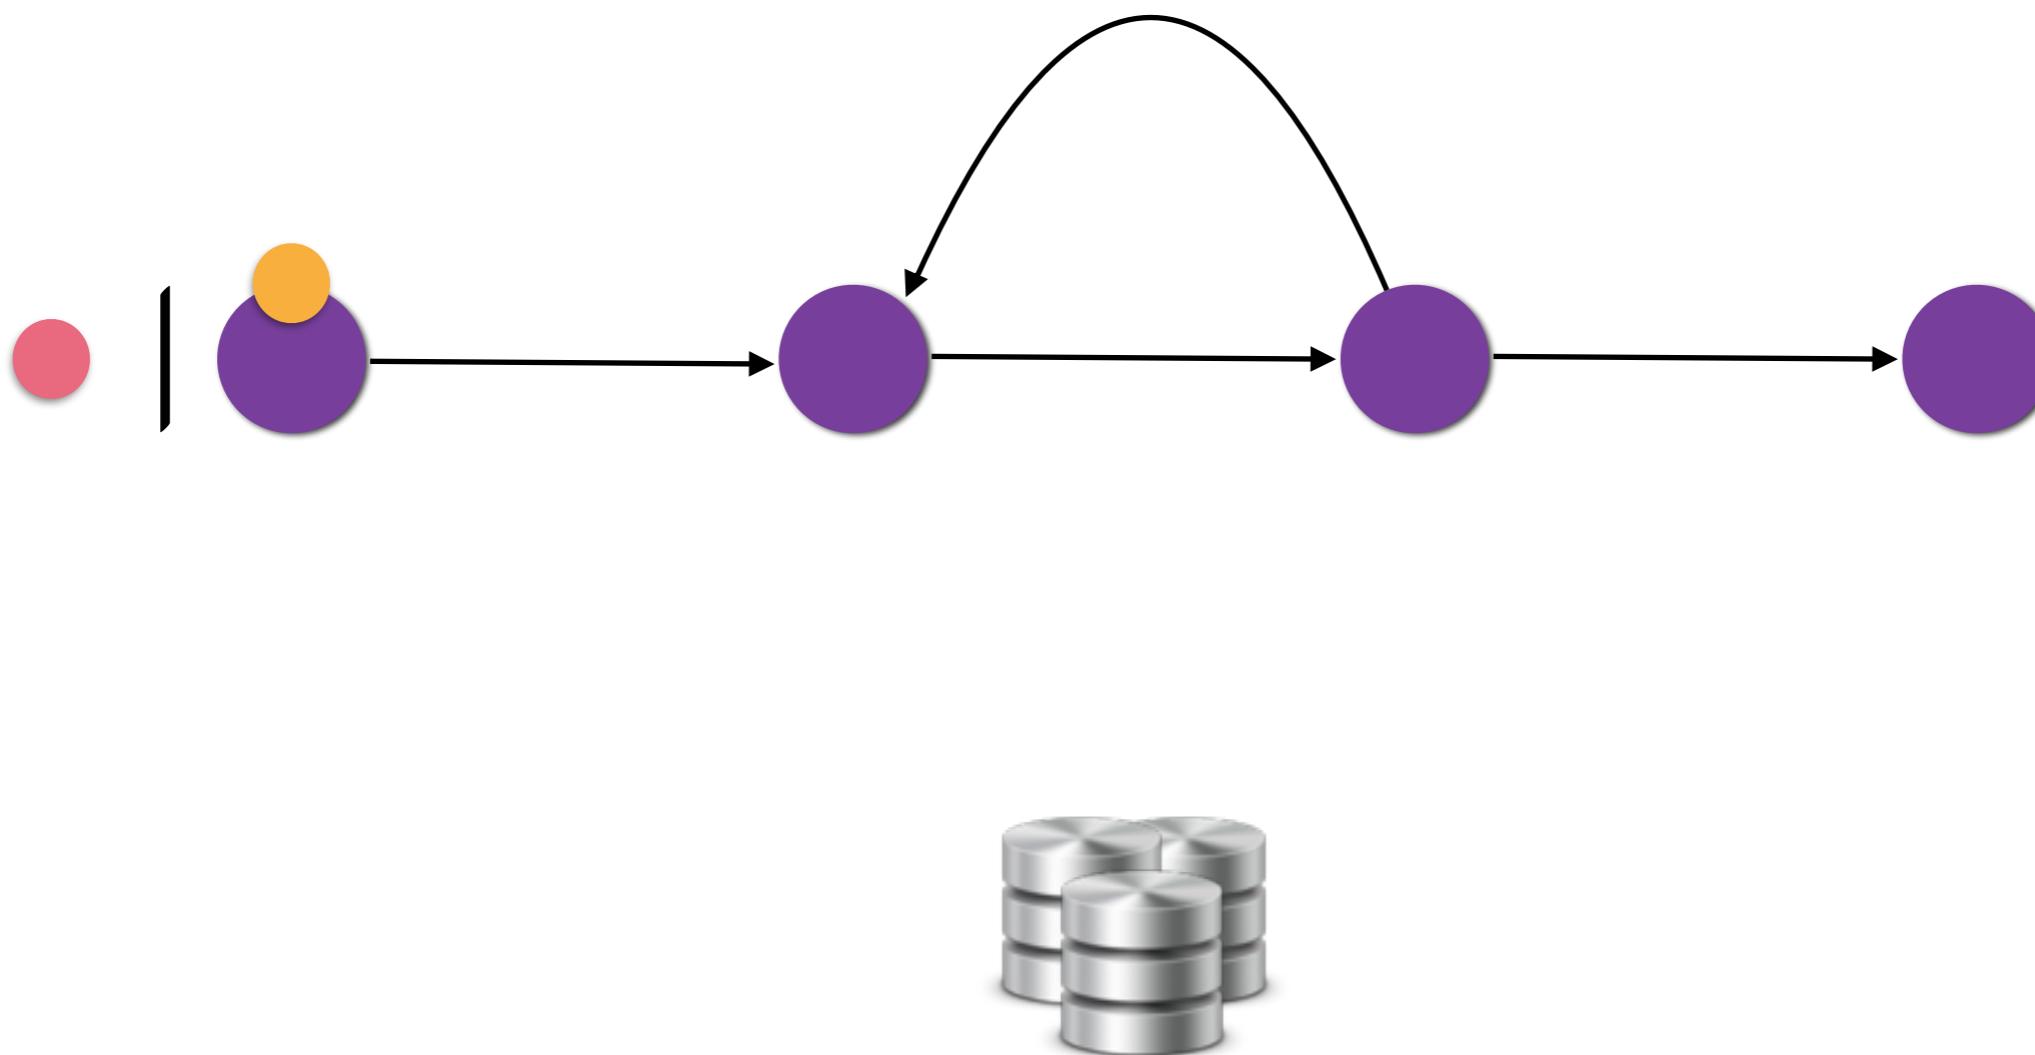
Asynchronous Barrier Snapshotting Benefits

- Taking advantage of the execution graph structure
- No records in transit included in the snapshot
- Aligning has lower impact than halting

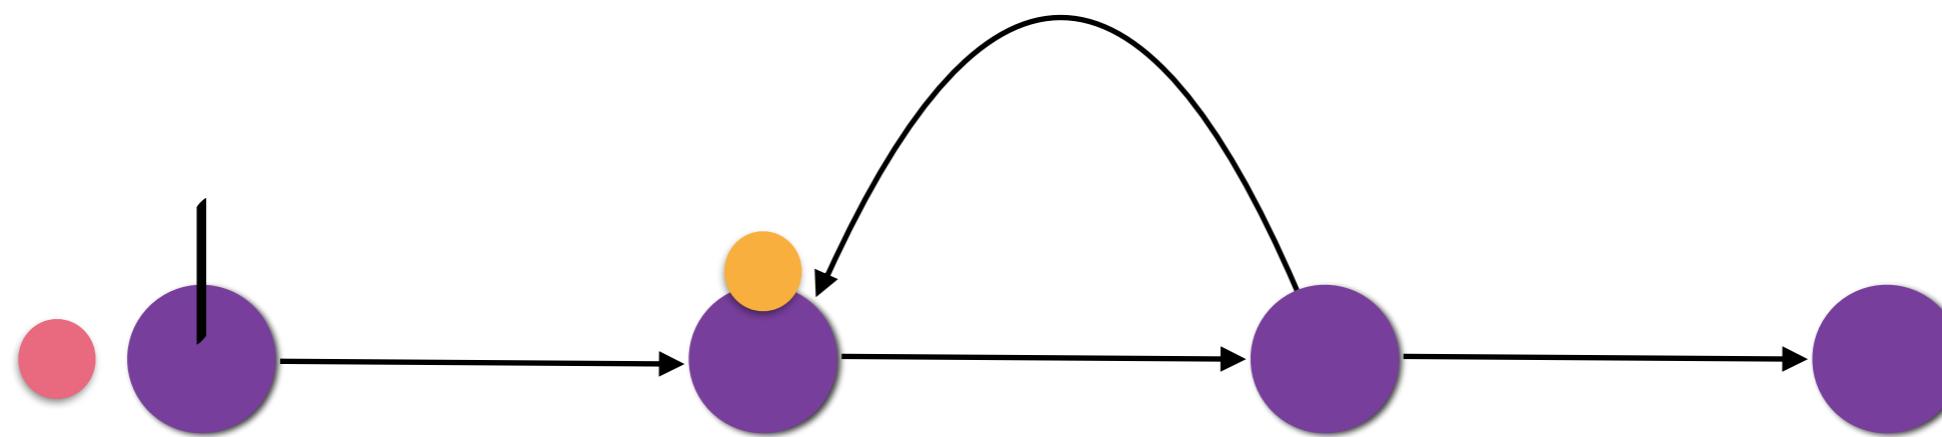
Snapshots on Cyclic Dataflows



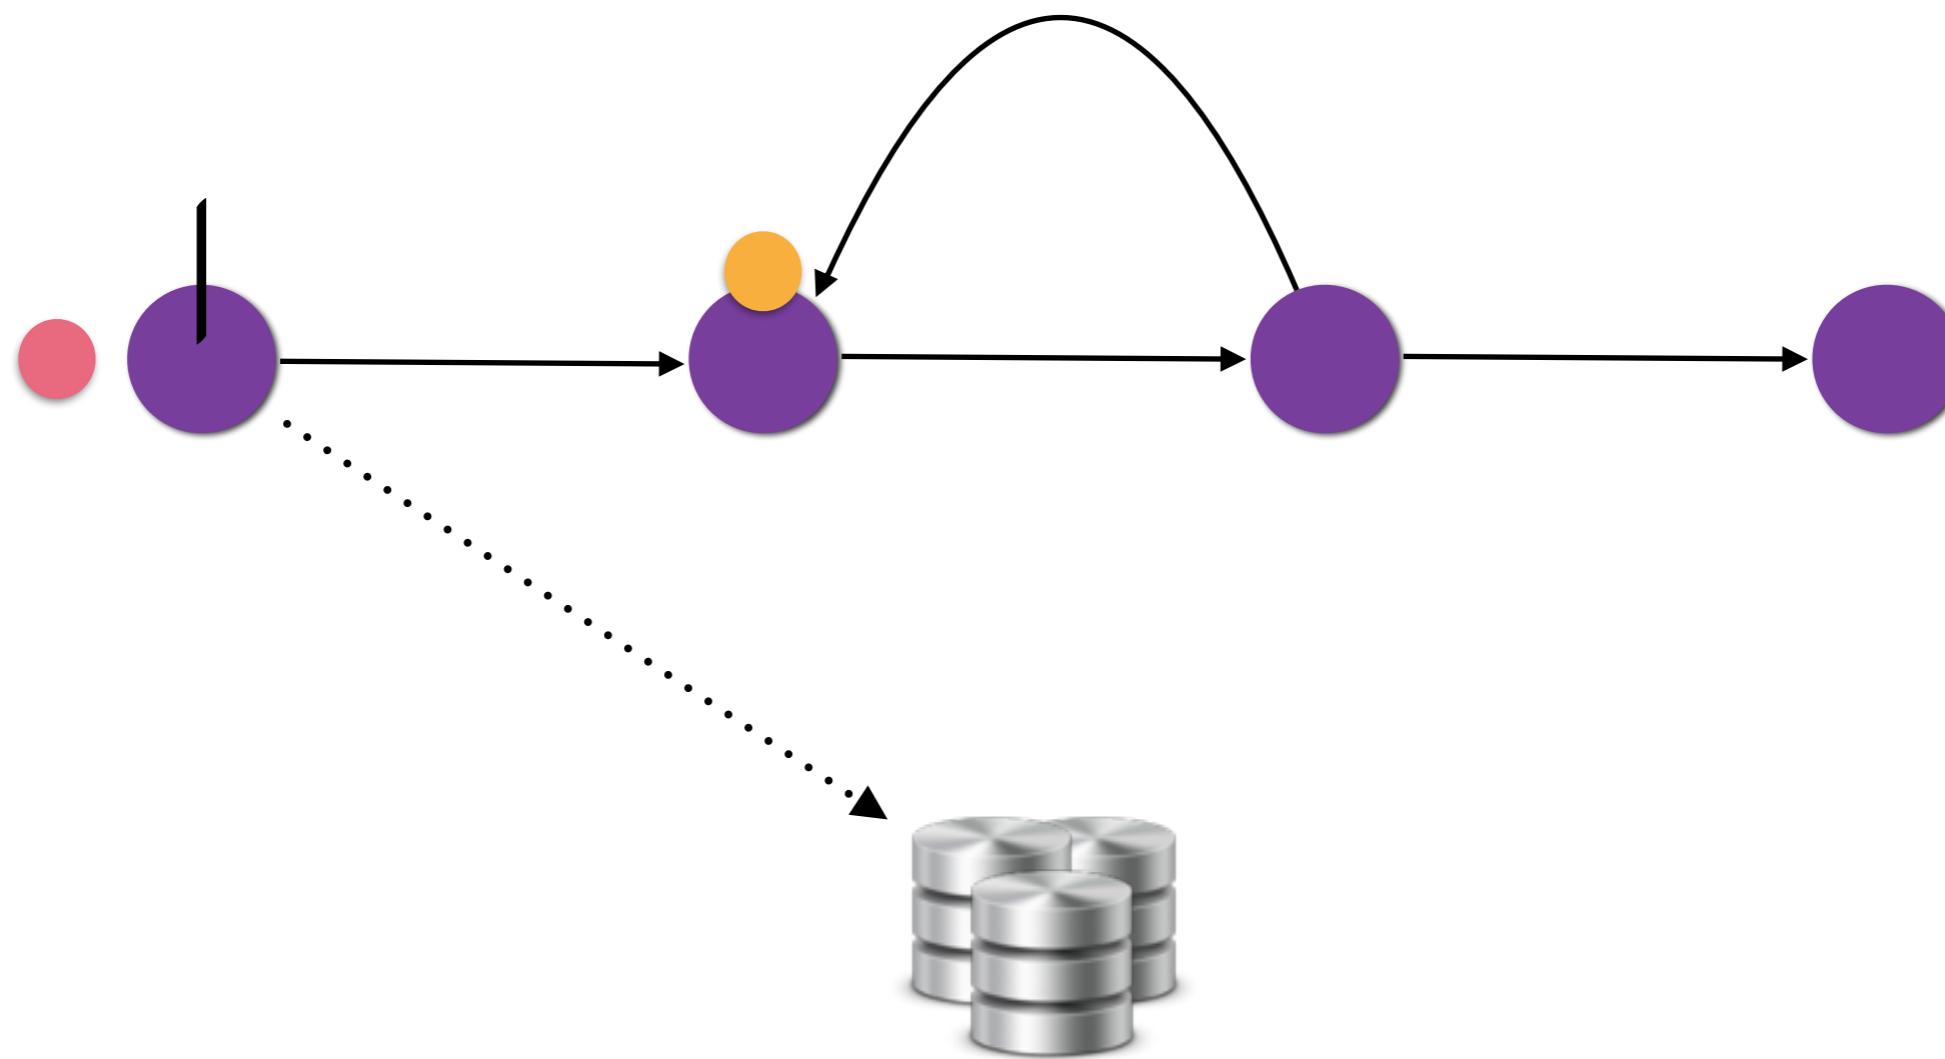
Snapshots on Cyclic Dataflows



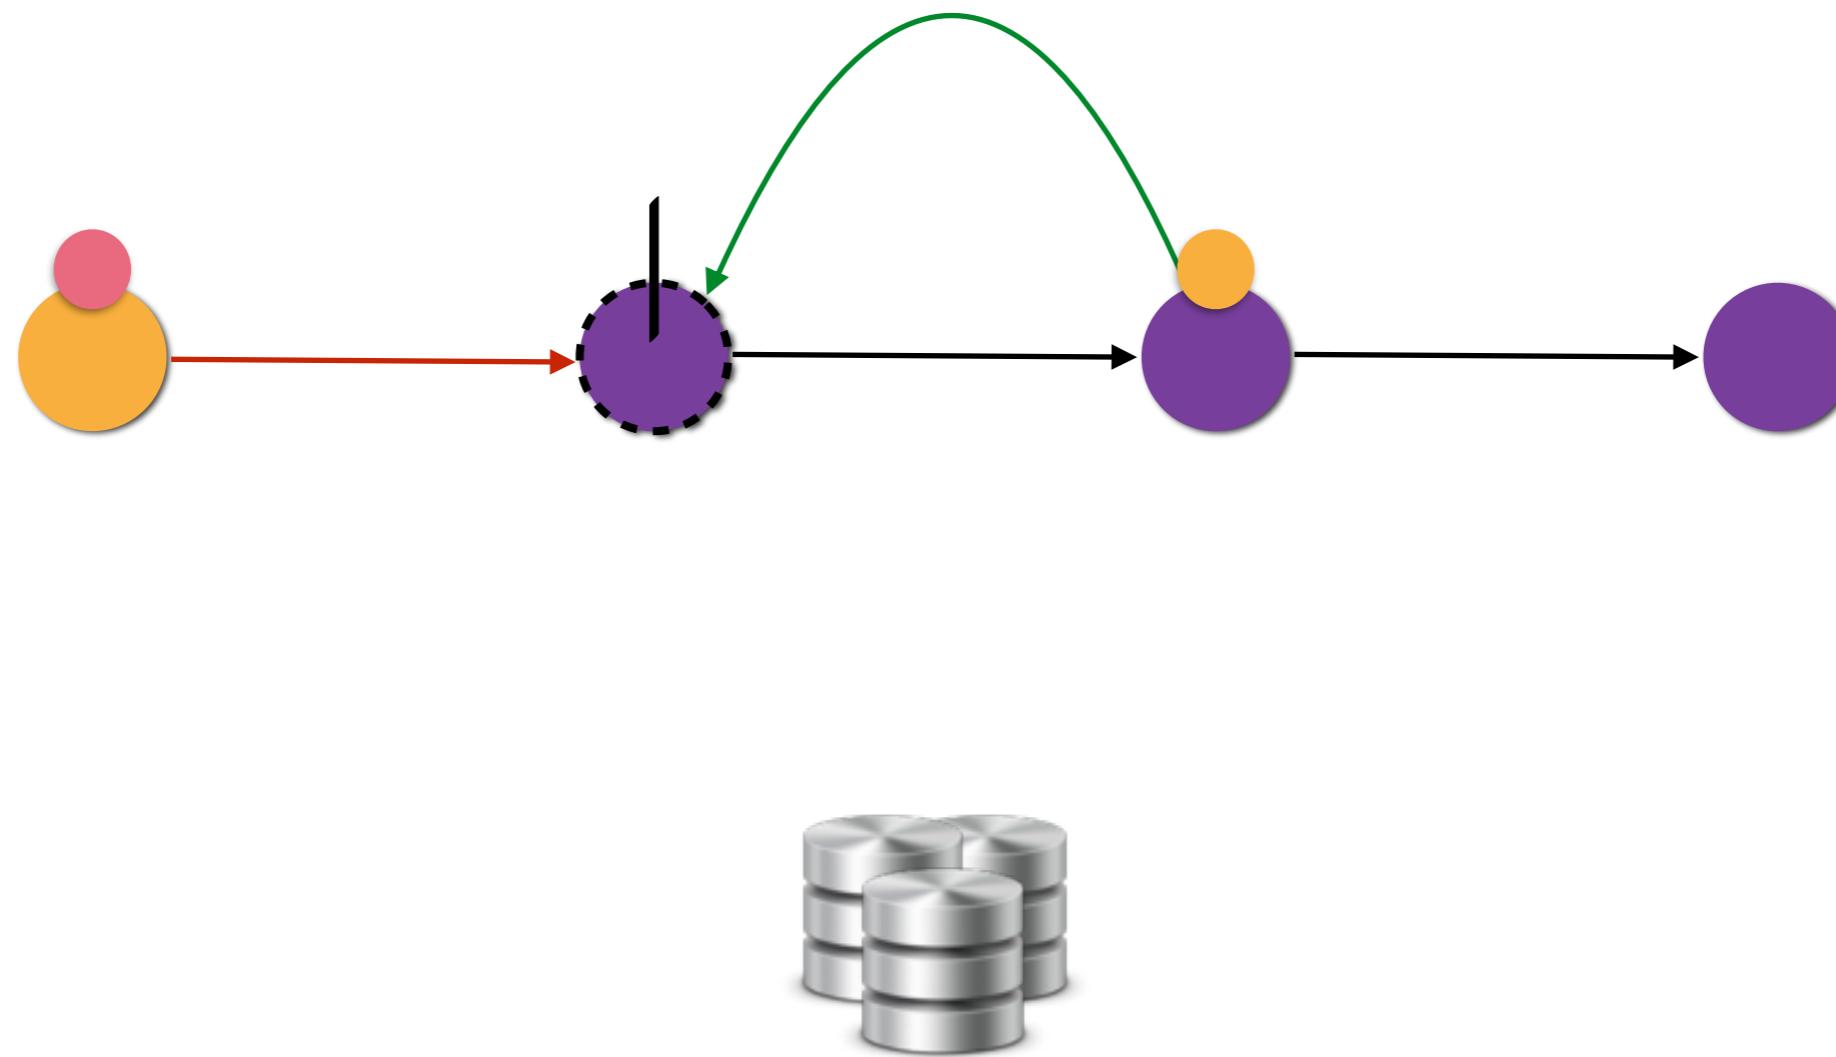
Snapshots on Cyclic Dataflows



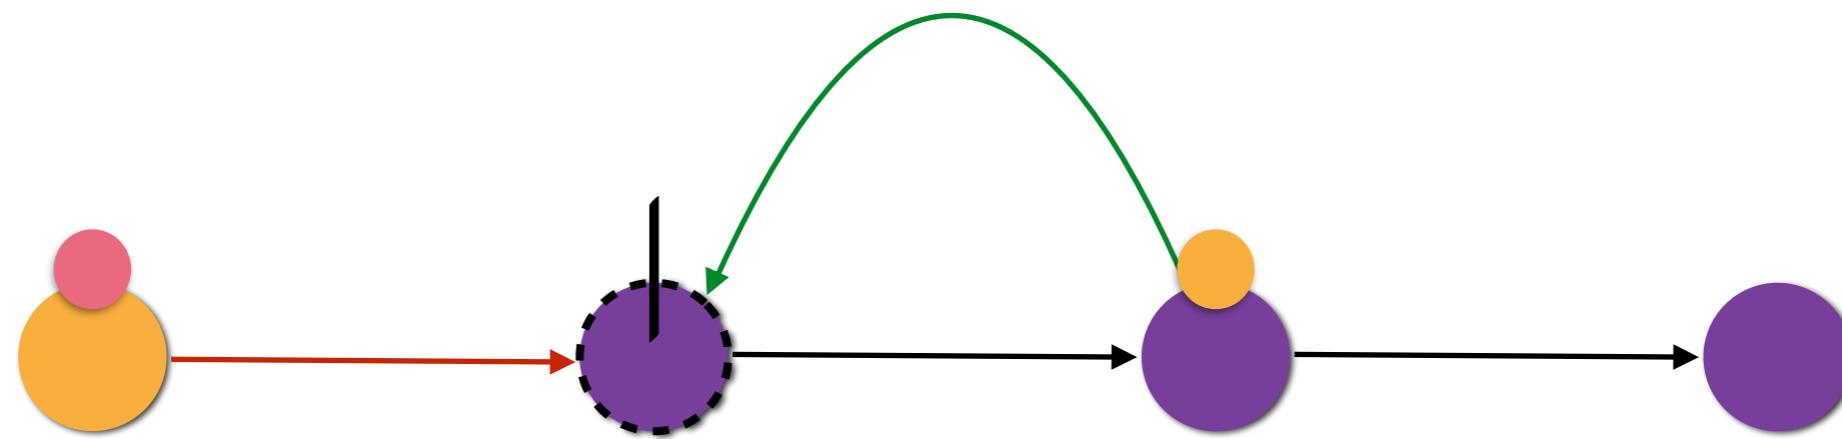
Snapshots on Cyclic Dataflows



Snapshots on Cyclic Dataflows



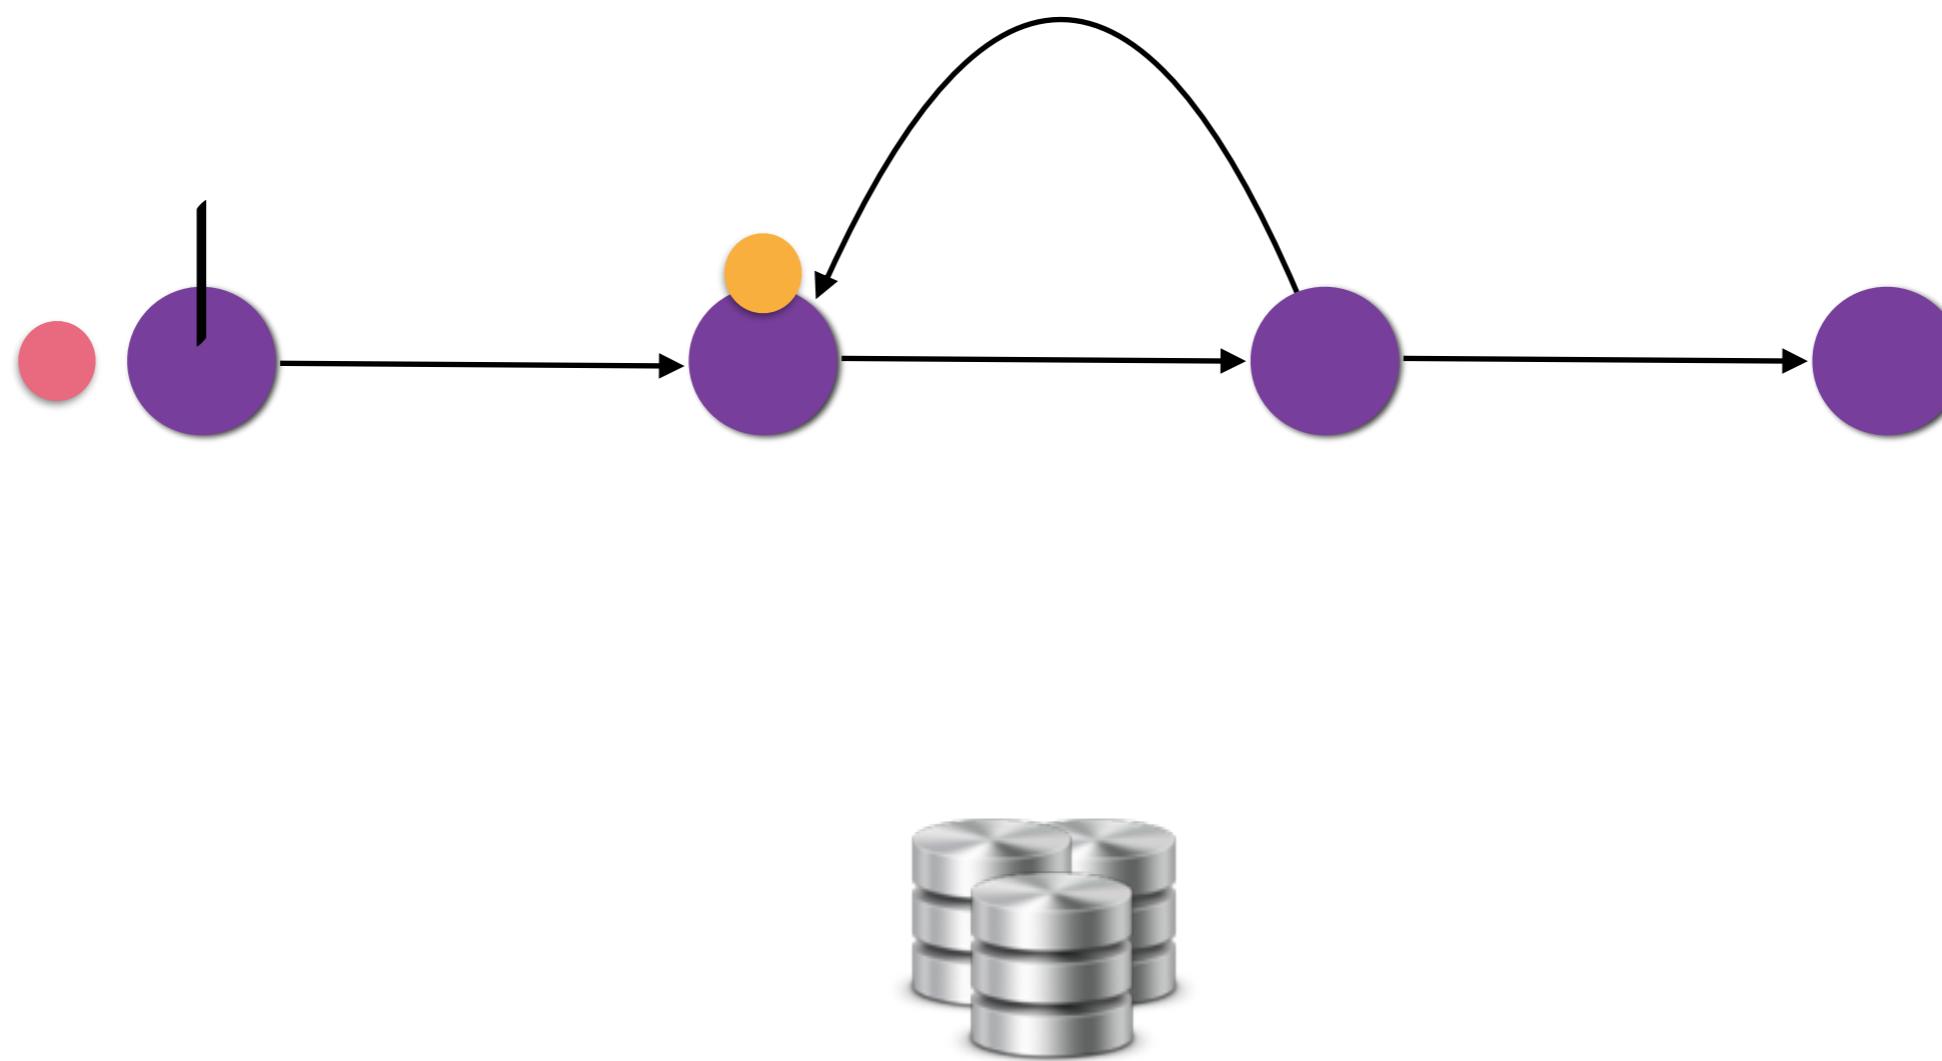
Snapshots on Cyclic Dataflows



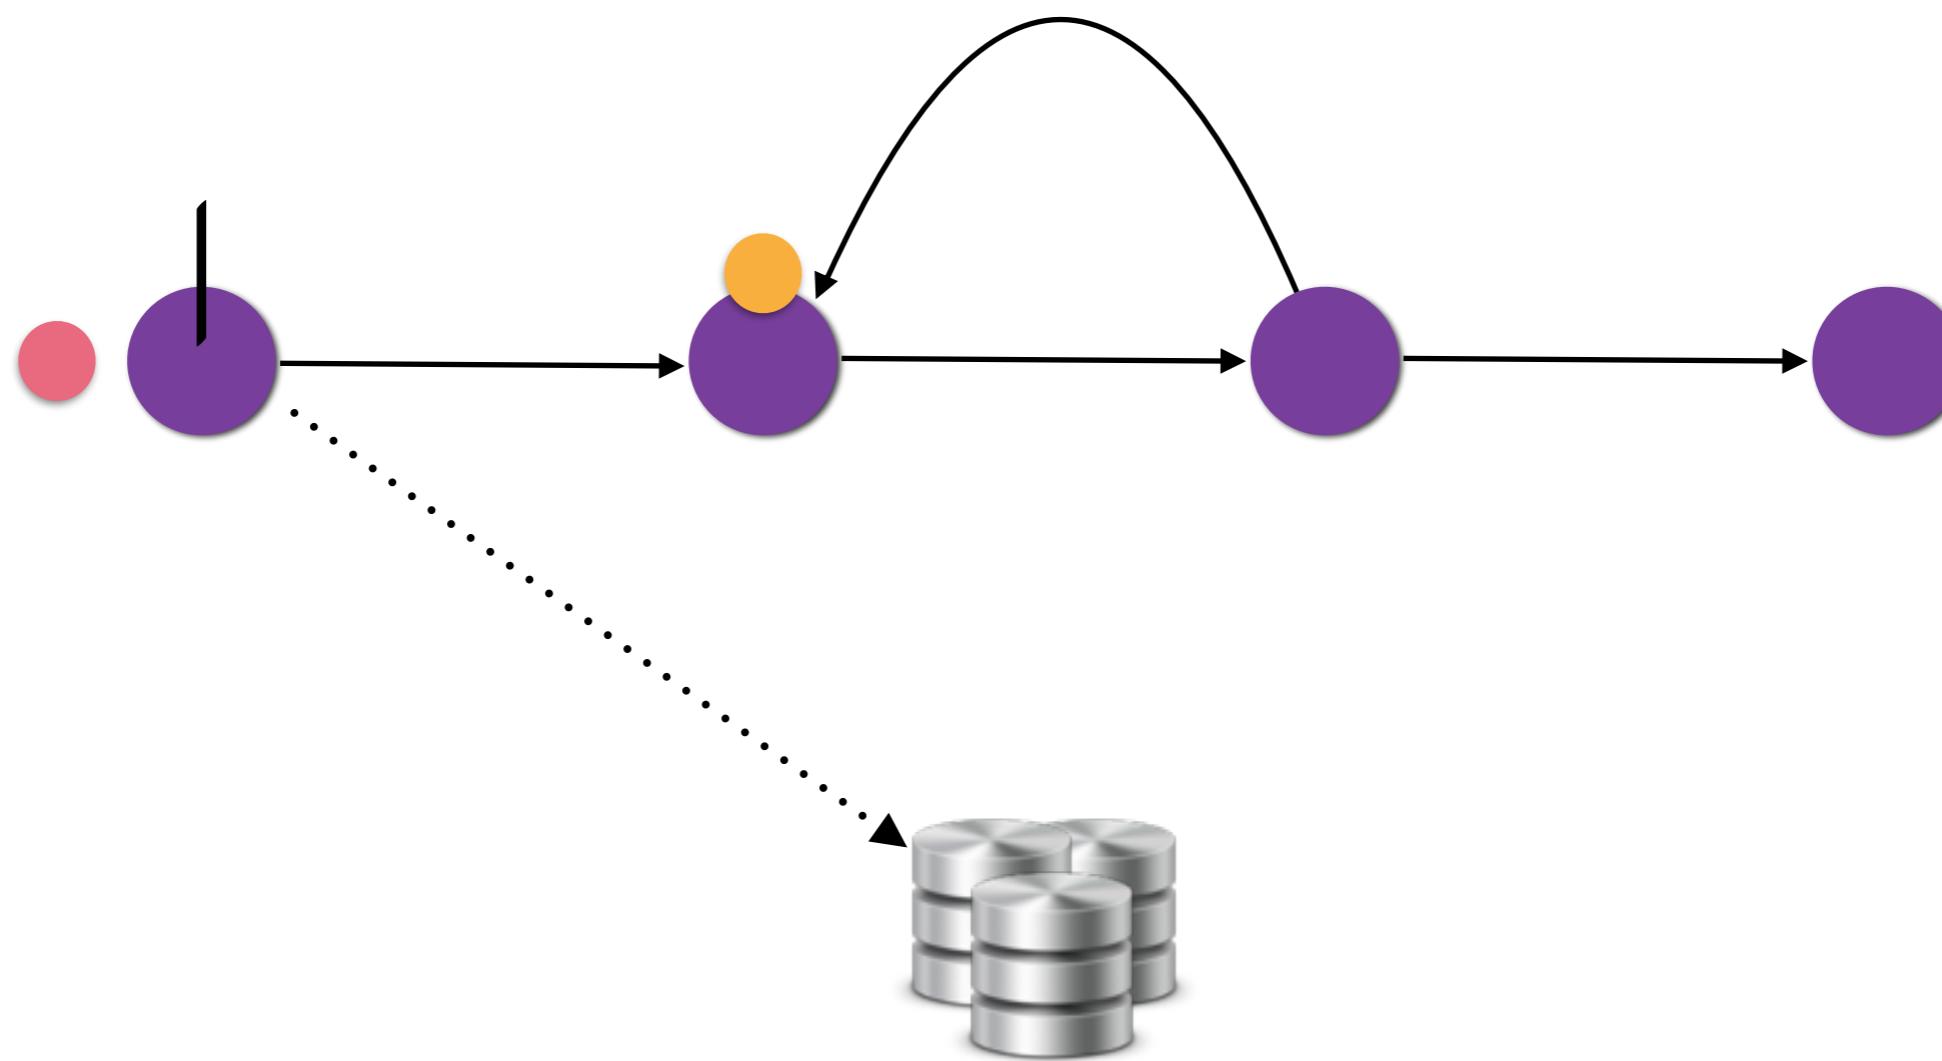
Snapshots on Cyclic Dataflows

- Checkpointing should eventually terminate
- Records in transit within loops should be included in the checkpoint

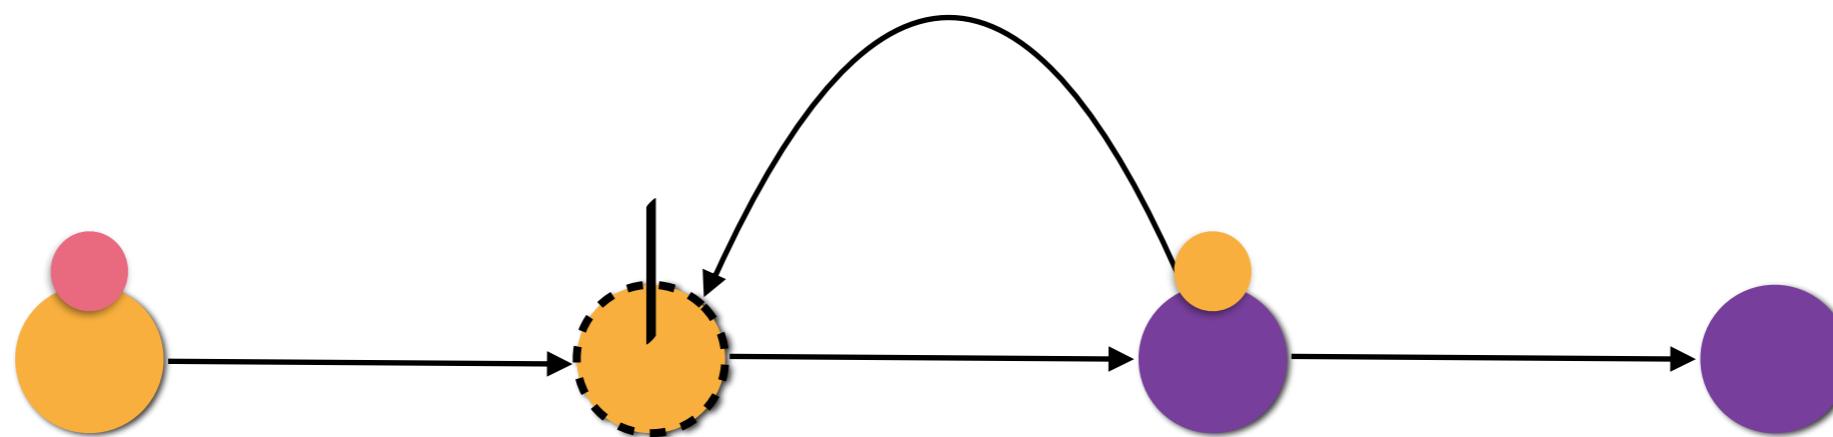
Snapshots on Cyclic Dataflows



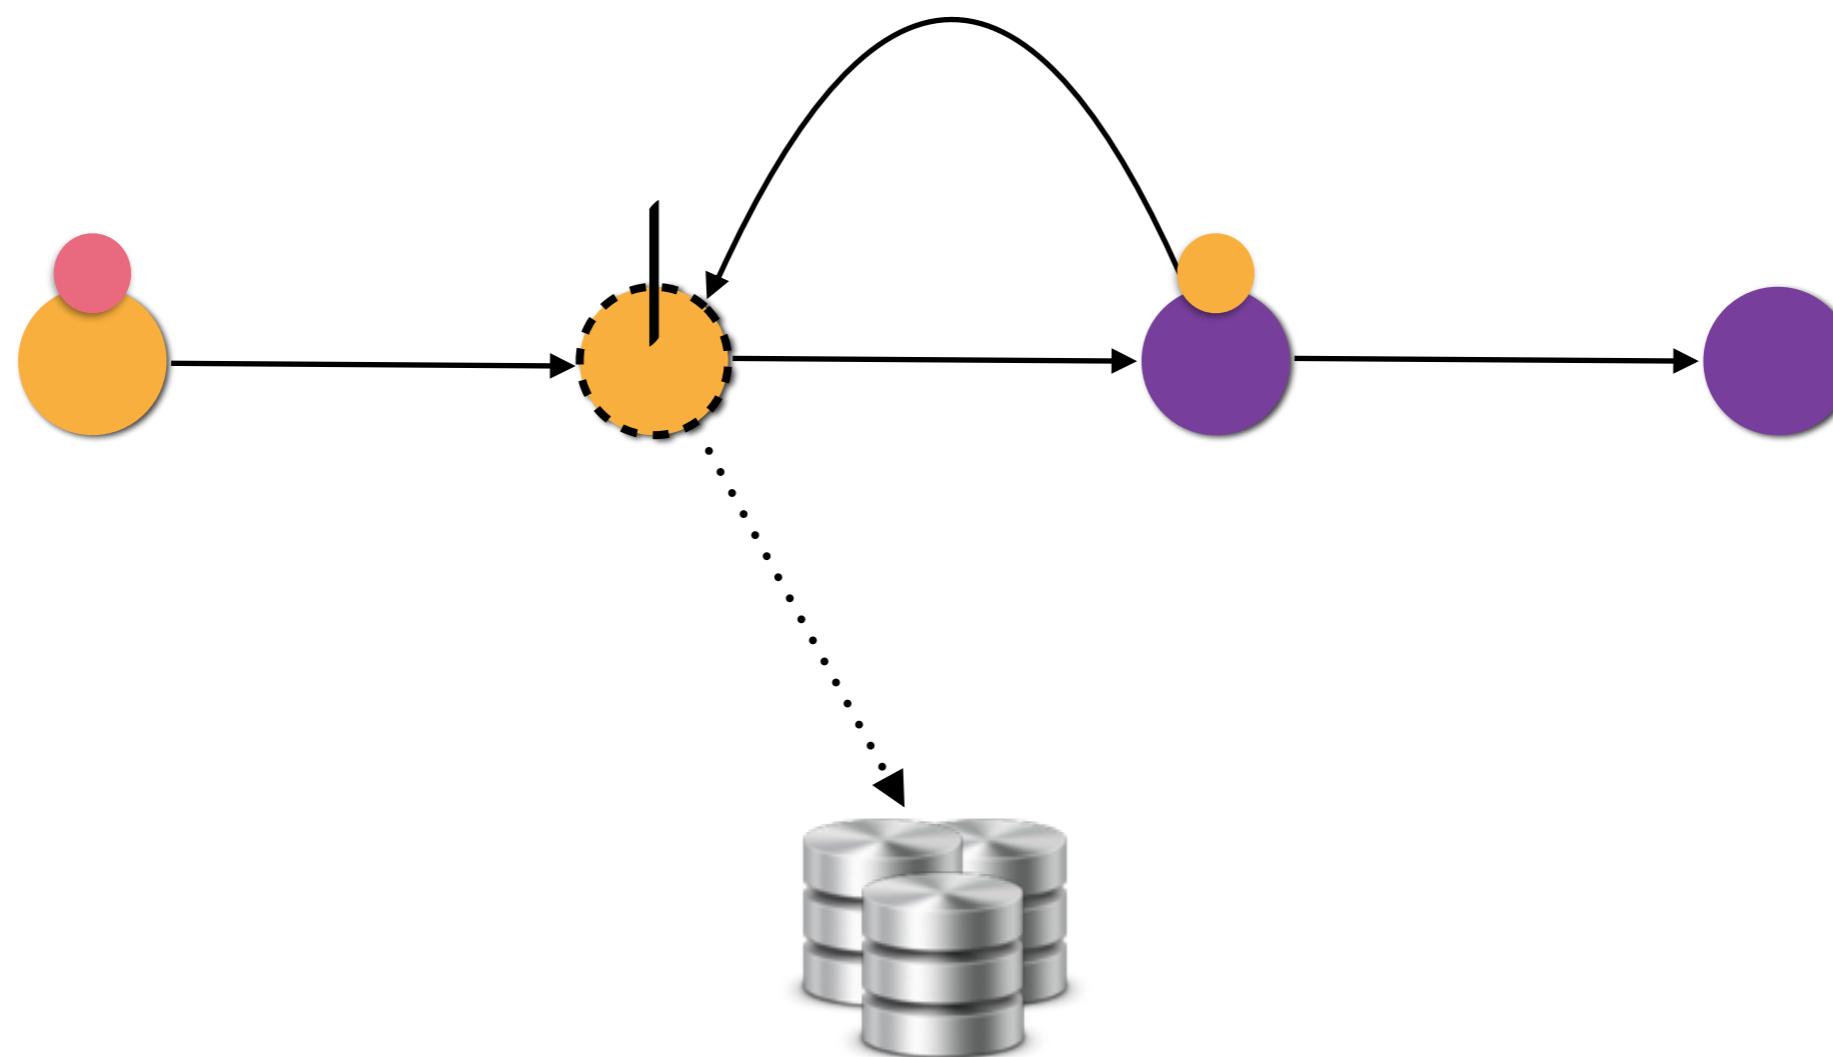
Snapshots on Cyclic Dataflows



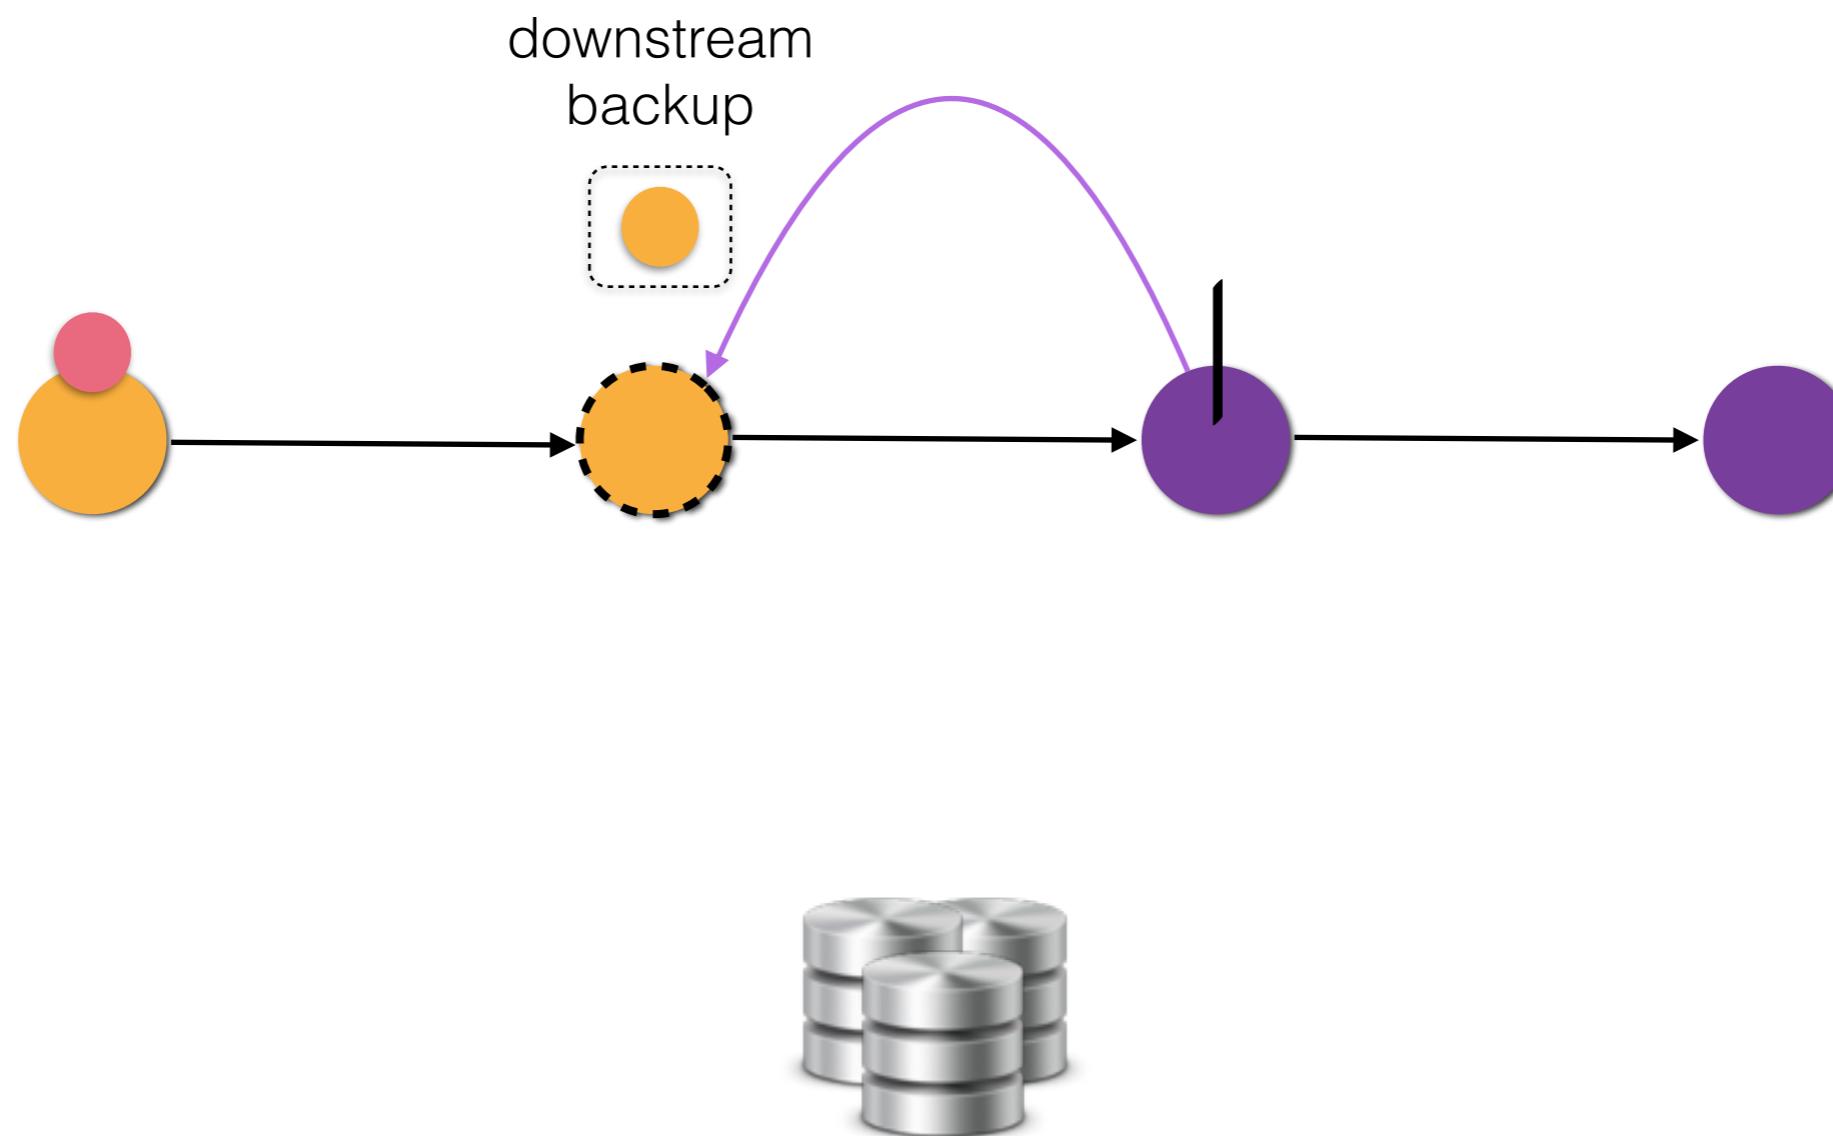
Snapshots on Cyclic Dataflows



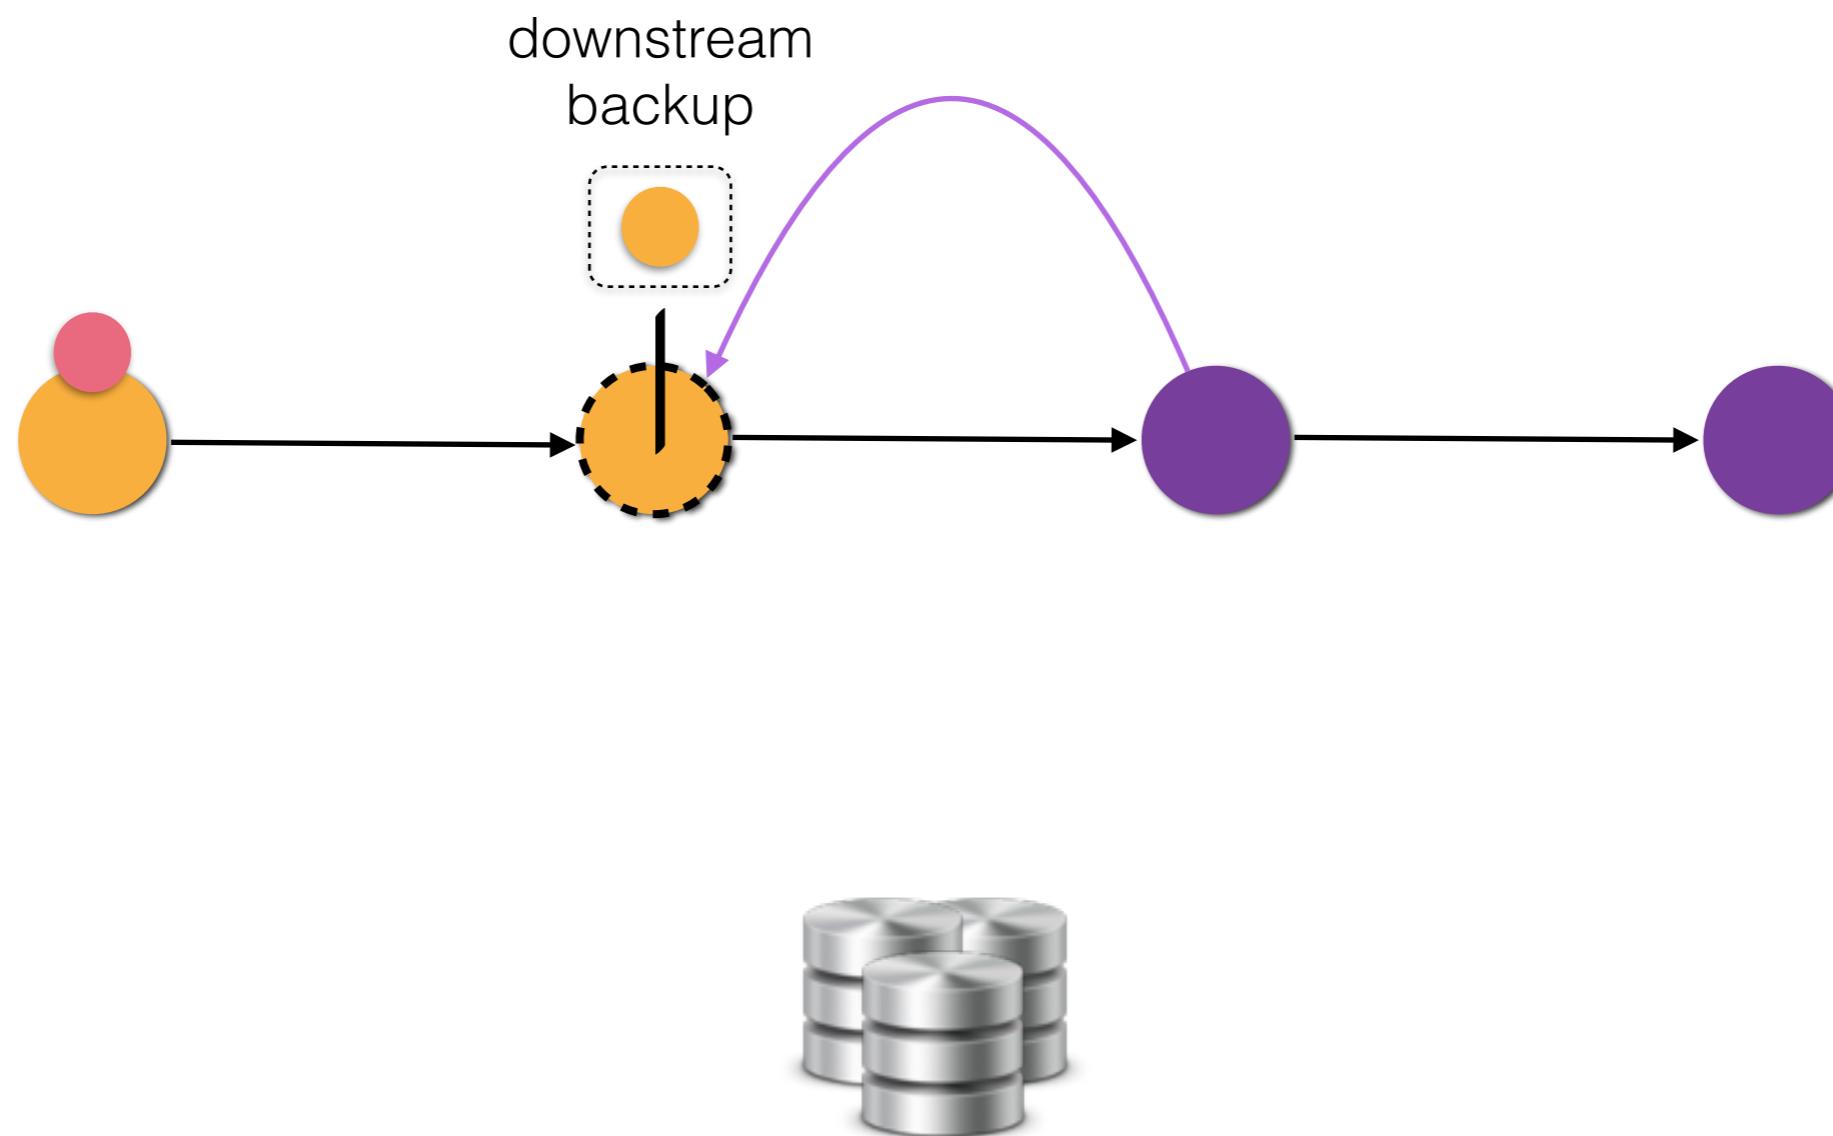
Snapshots on Cyclic Dataflows



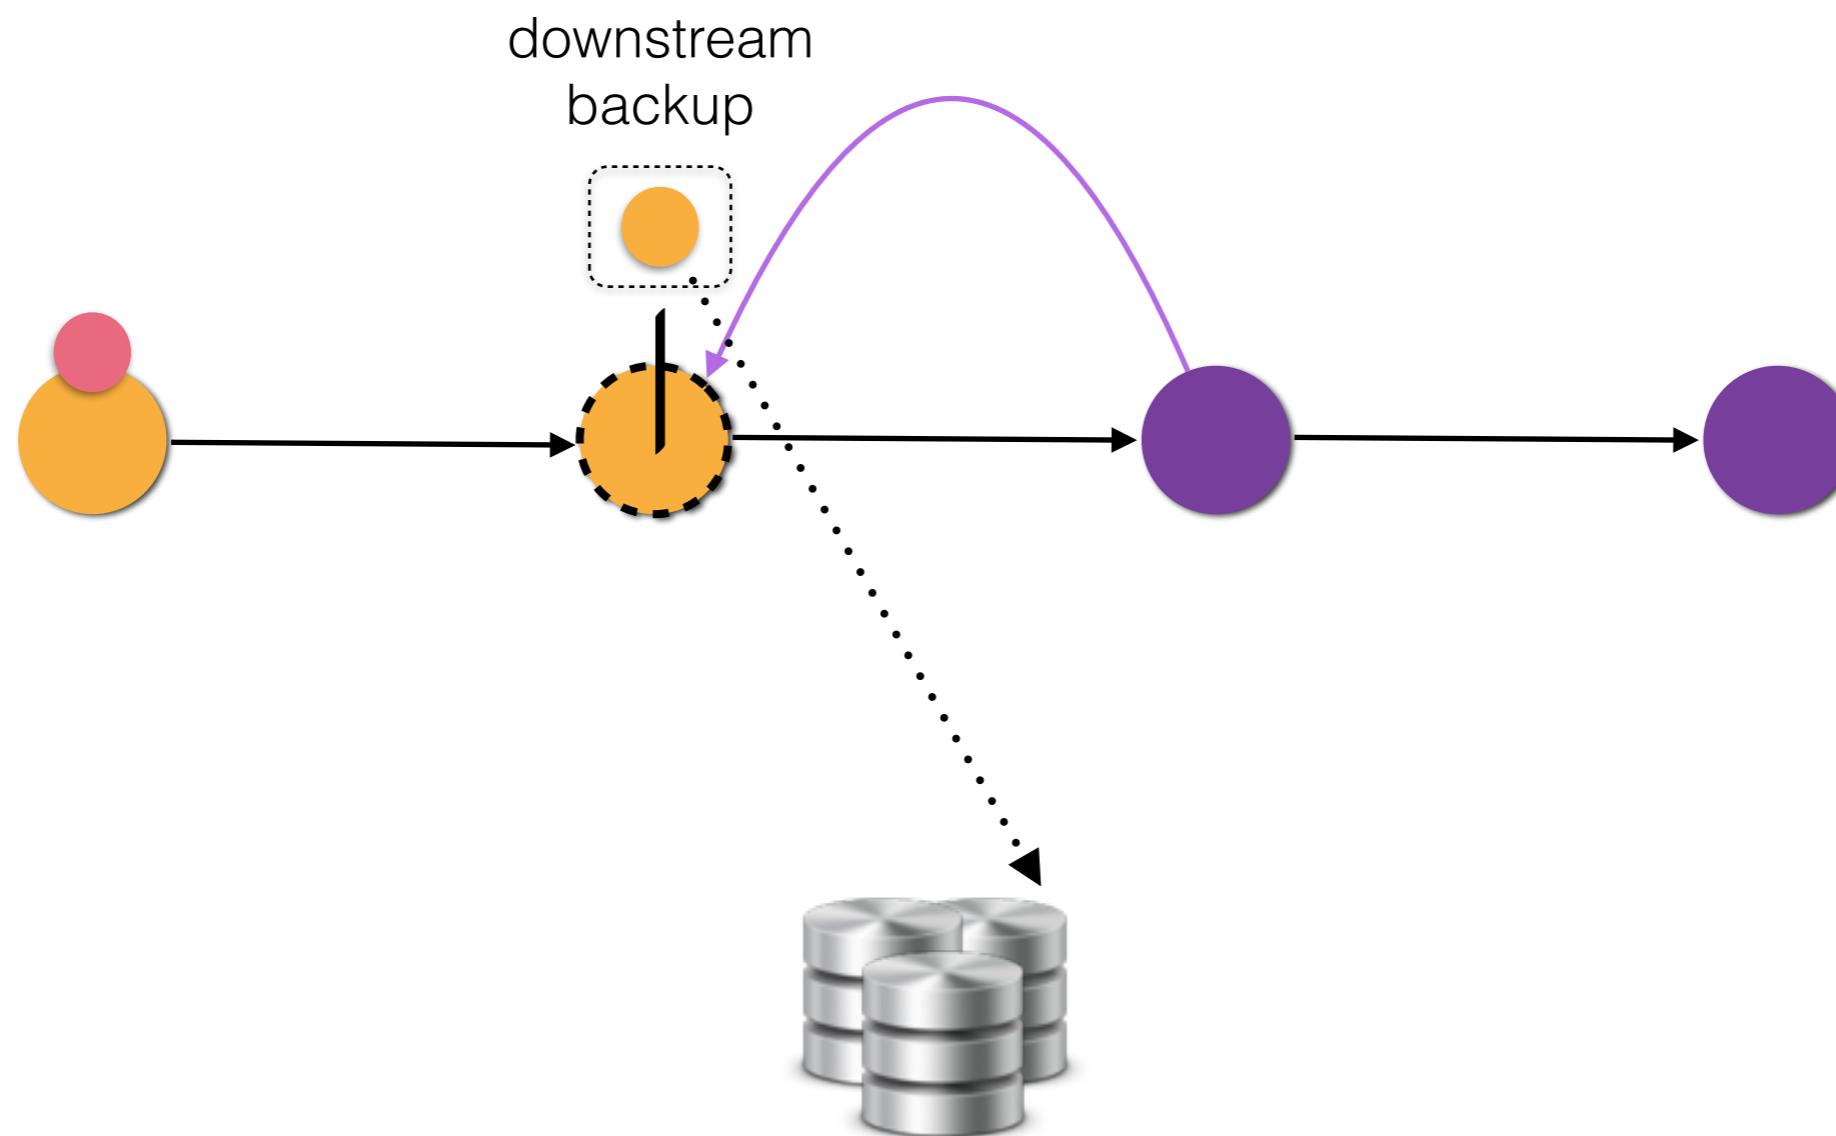
Snapshots on Cyclic Dataflows



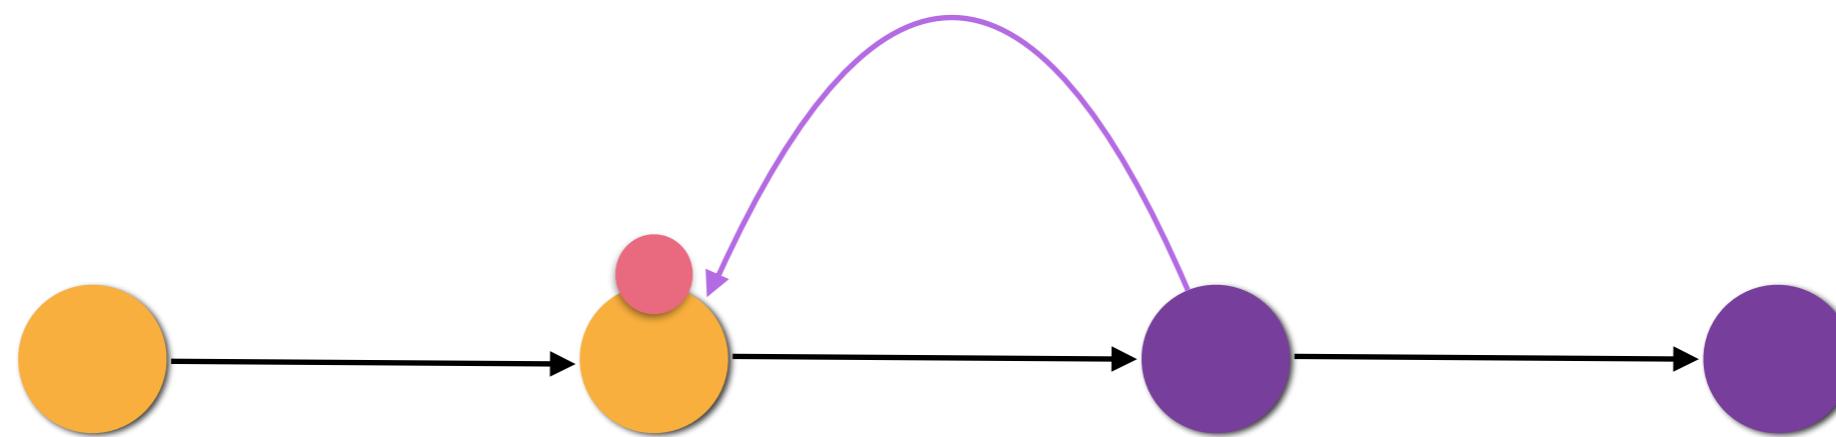
Snapshots on Cyclic Dataflows



Snapshots on Cyclic Dataflows



Snapshots on Cyclic Dataflows



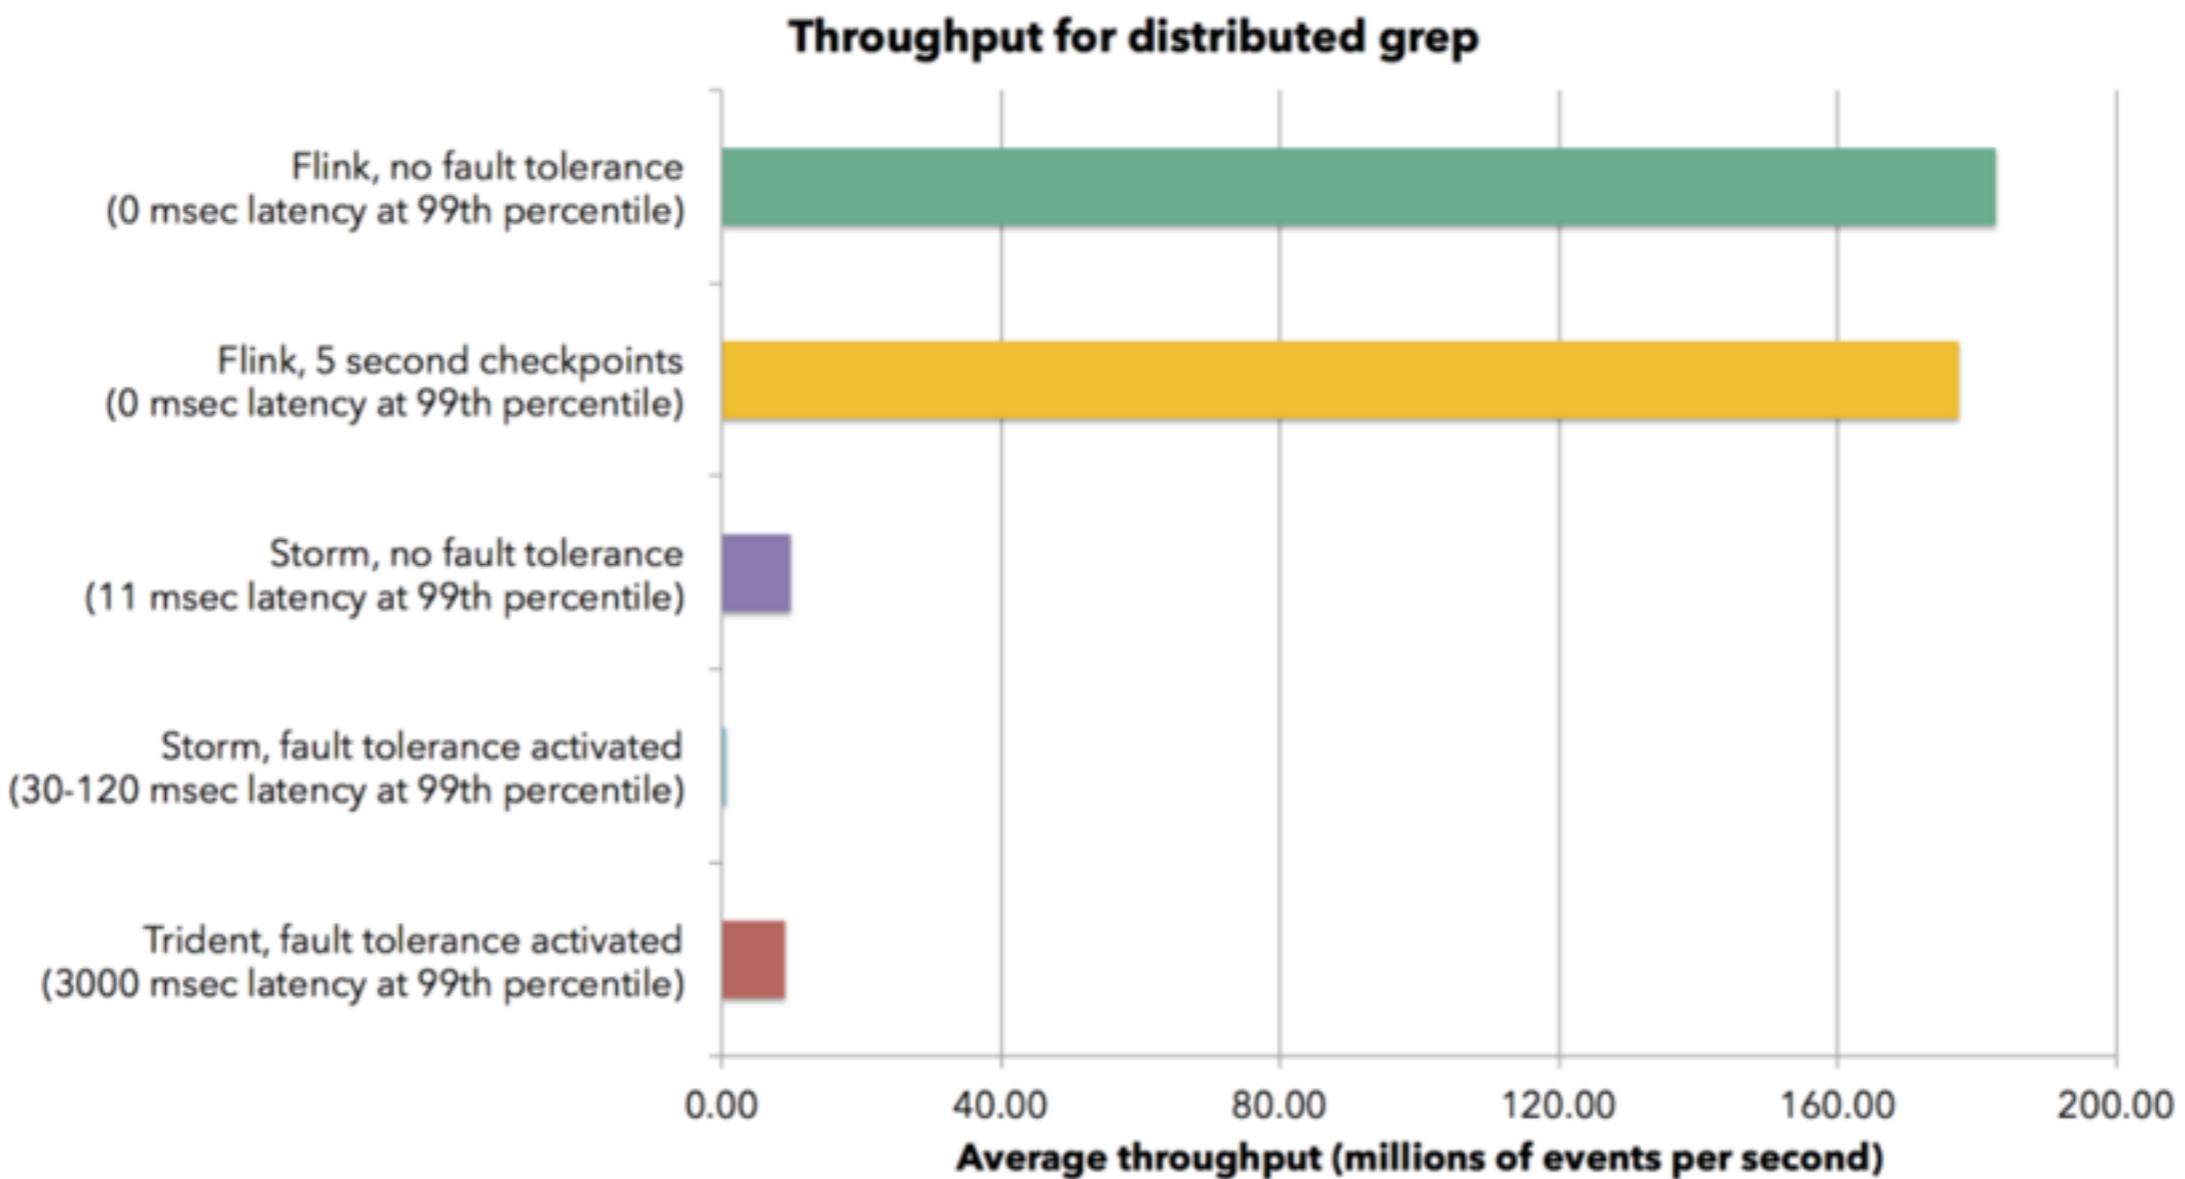
Implementation in Flink

- **Coordinator**/Driver Actor per job in JM
 - sends periodic markers to sources
 - collects acknowledgements with state handles and registers complete snapshots
 - injects references to latest consistent state handles and back-edge records in pending execution graph tasks
- **Tasks**
 - snapshot state transparently in given backend upon receiving a barrier and send ack to coordinator
 - propagate barriers further like normal records

Recovery

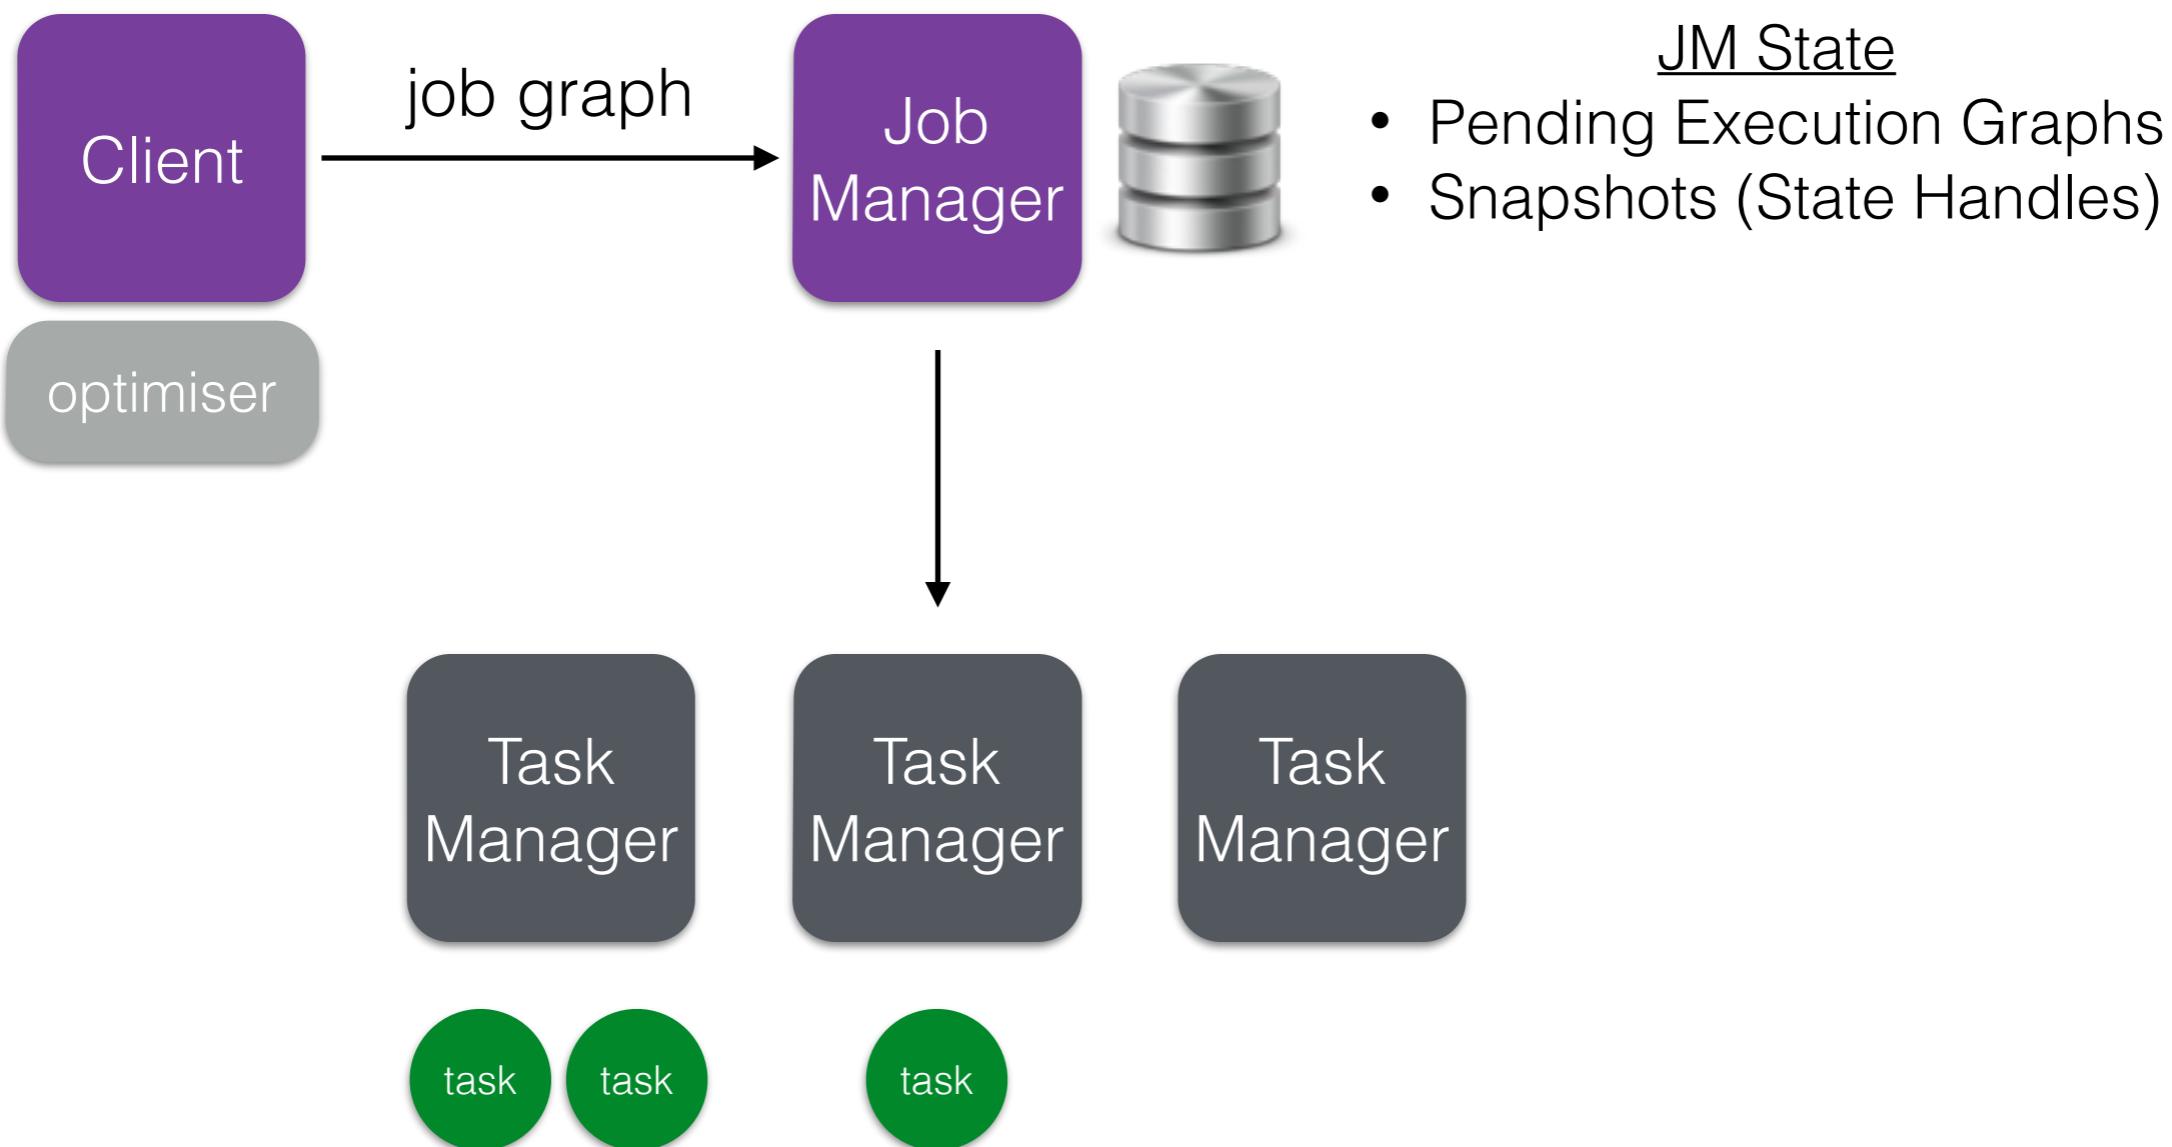
- **Full** rescheduling of the execution graph from the latest checkpoint - simple, no further modifications
- **Partial** rescheduling of the execution graph for all upstream dependencies - trickier, duplicate elimination should be guaranteed

Performance Impact

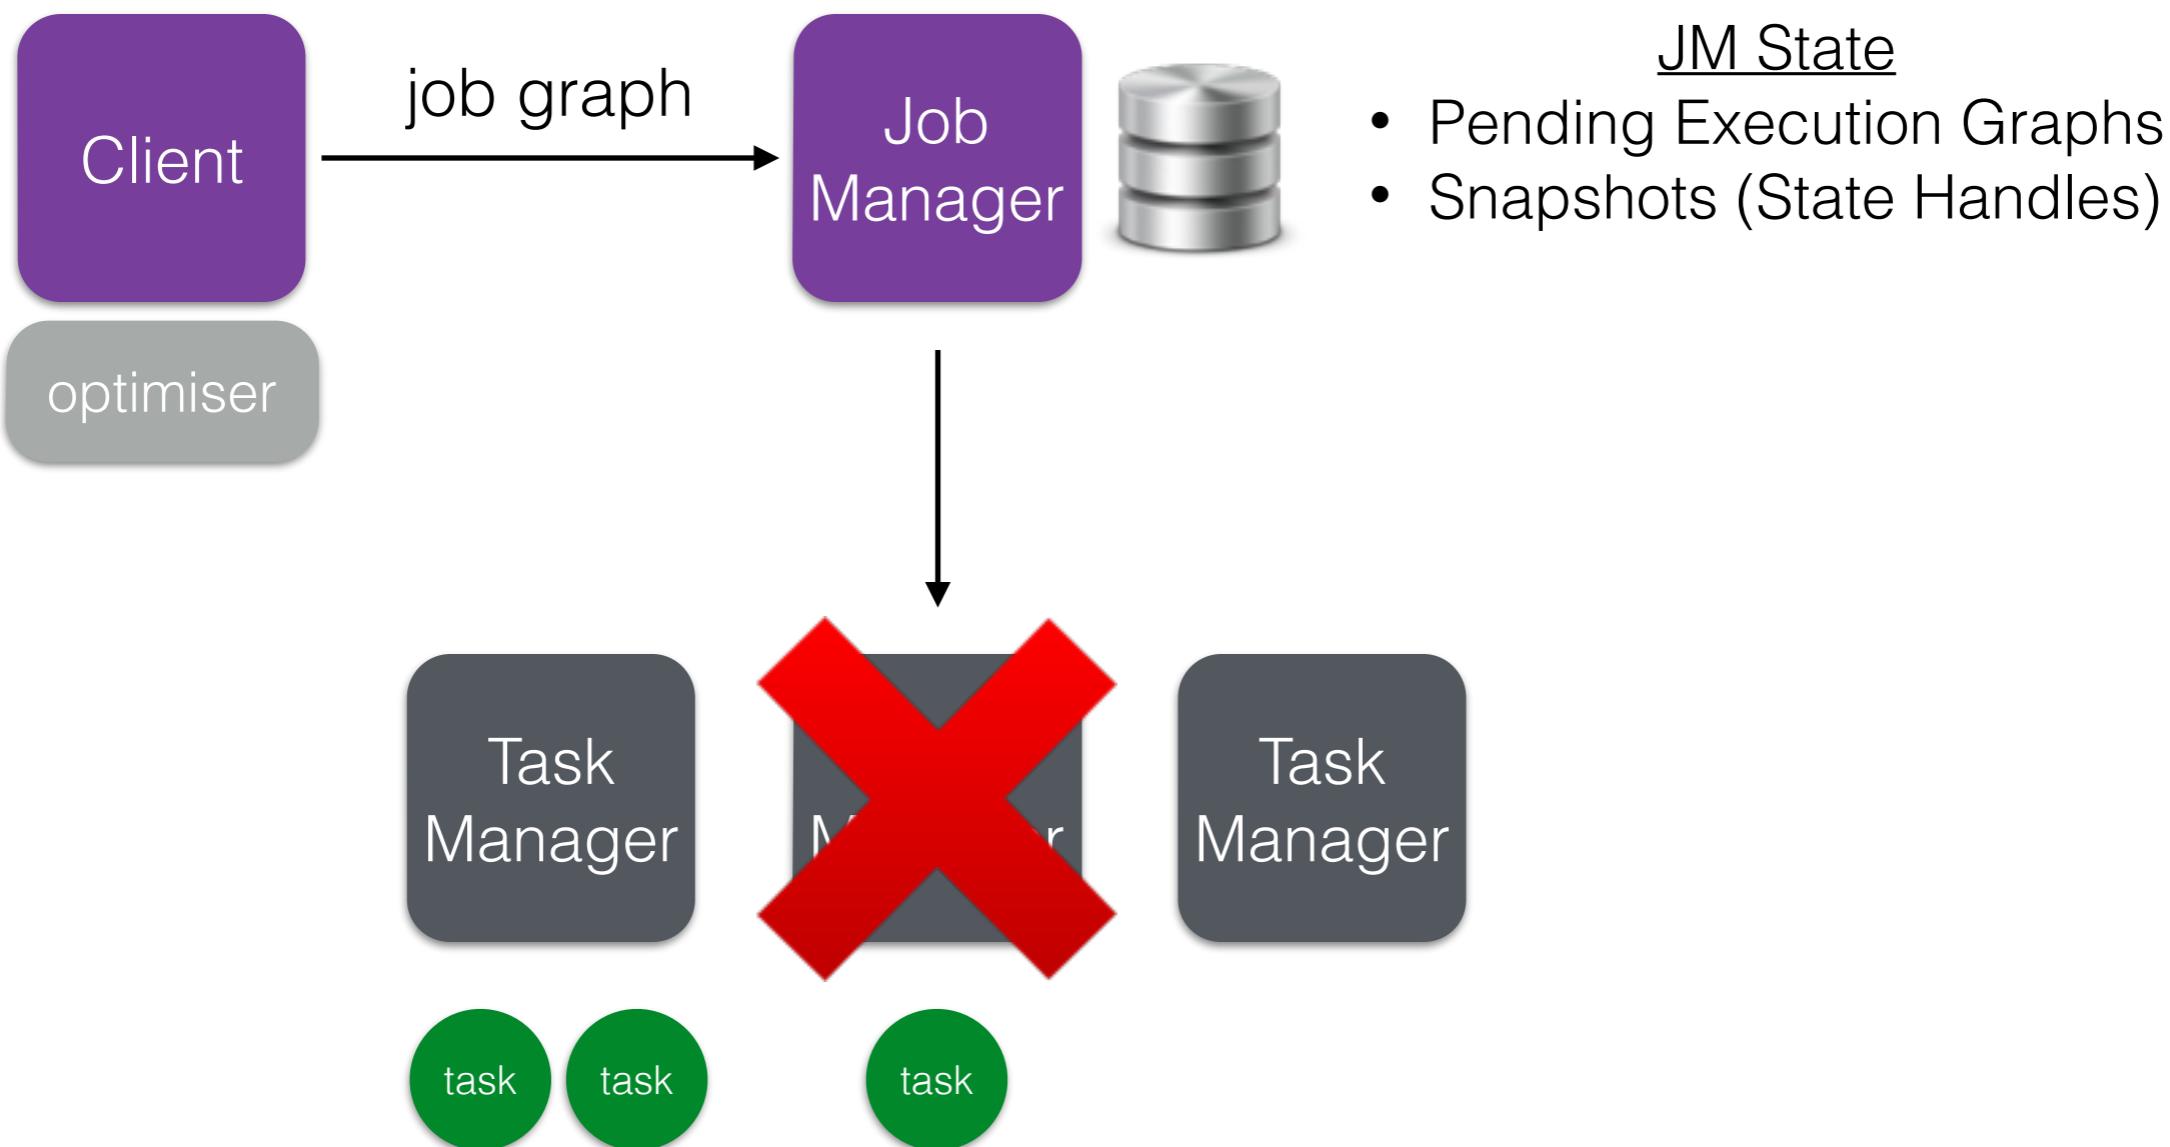


<http://data-artisans.com/high-throughput-low-latency-and-exactly-once-stream-processing-with-apache-flink/>

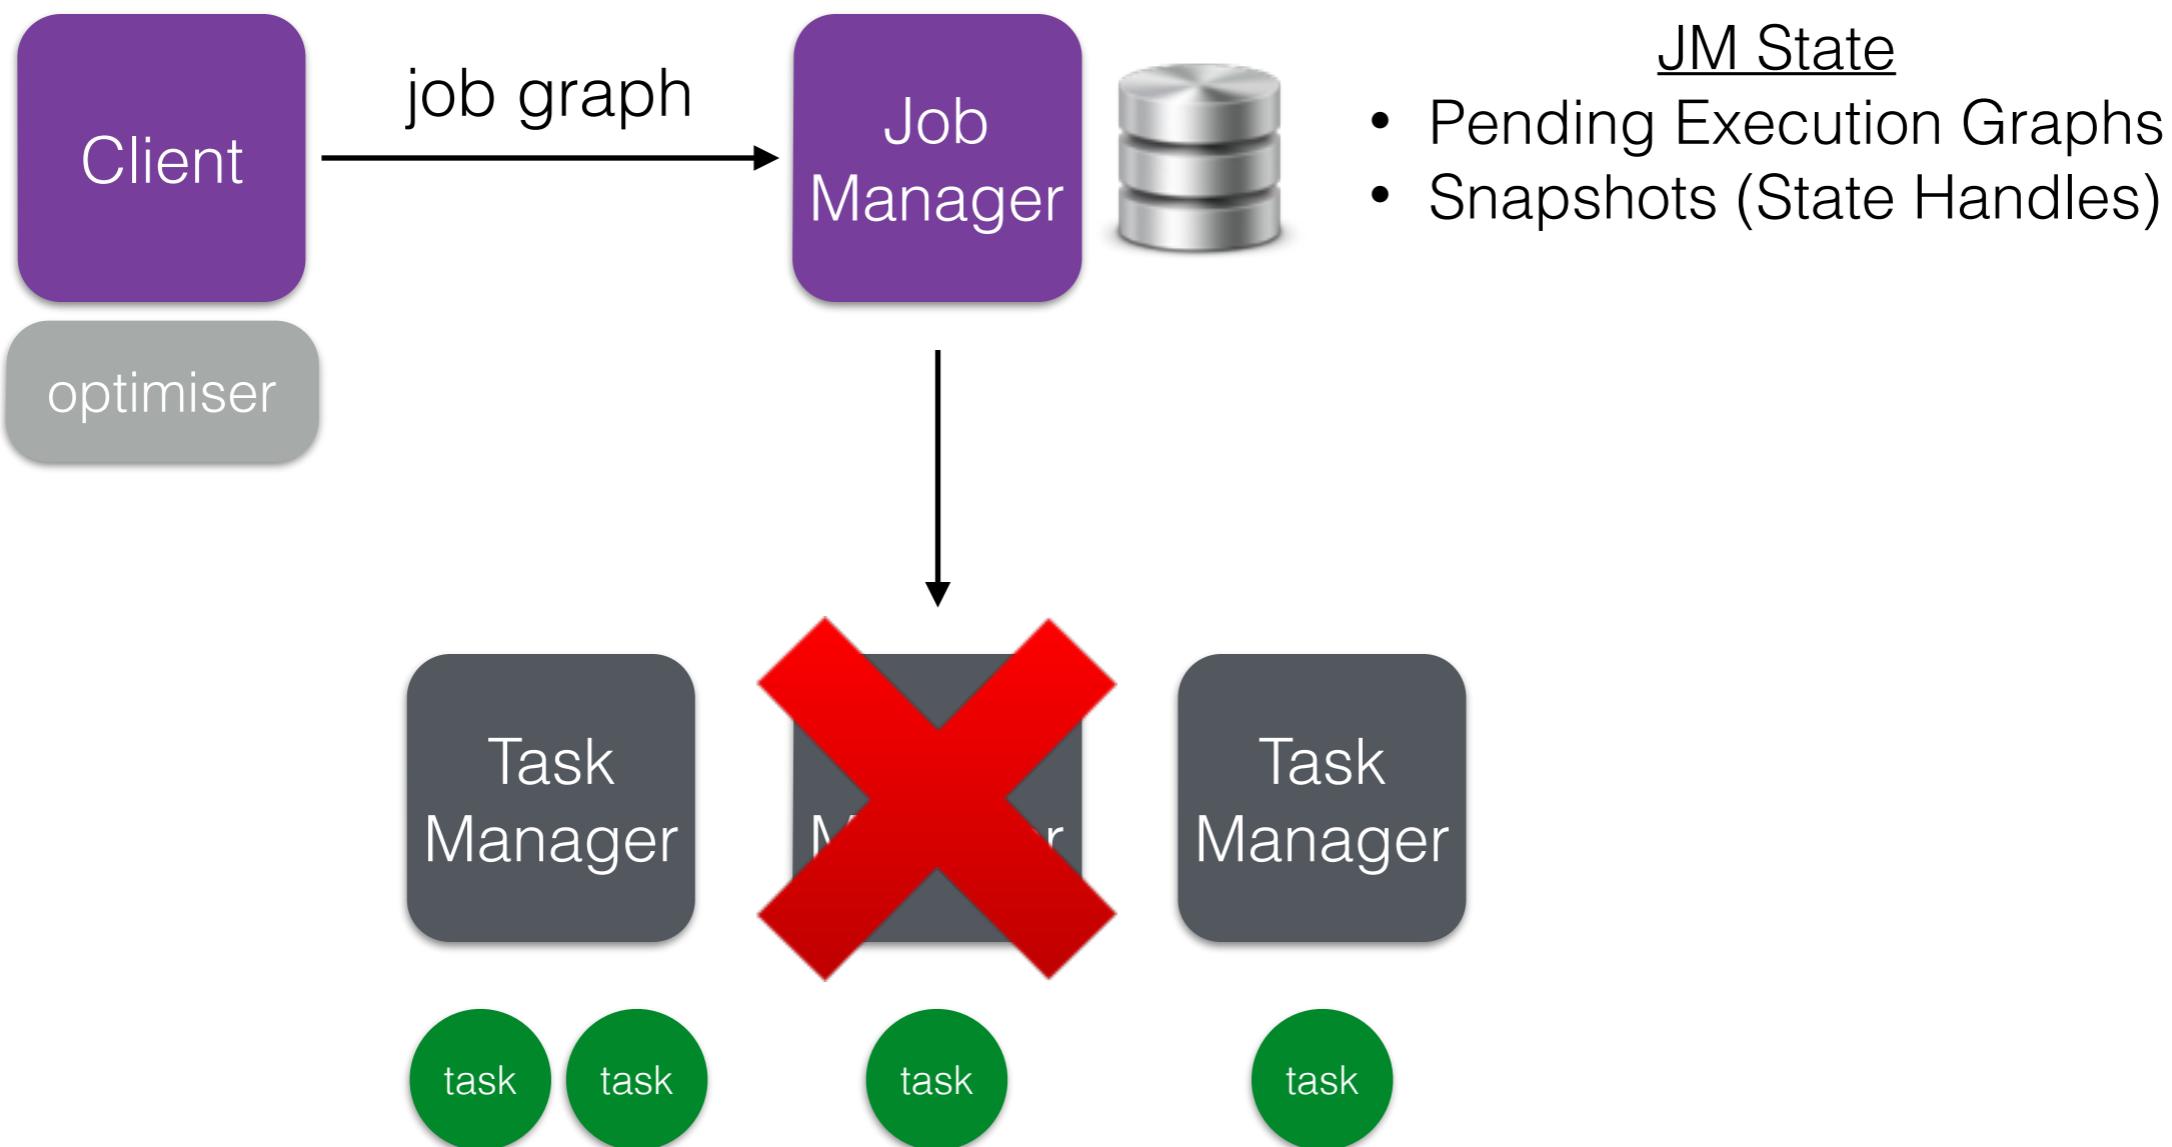
Node Failures



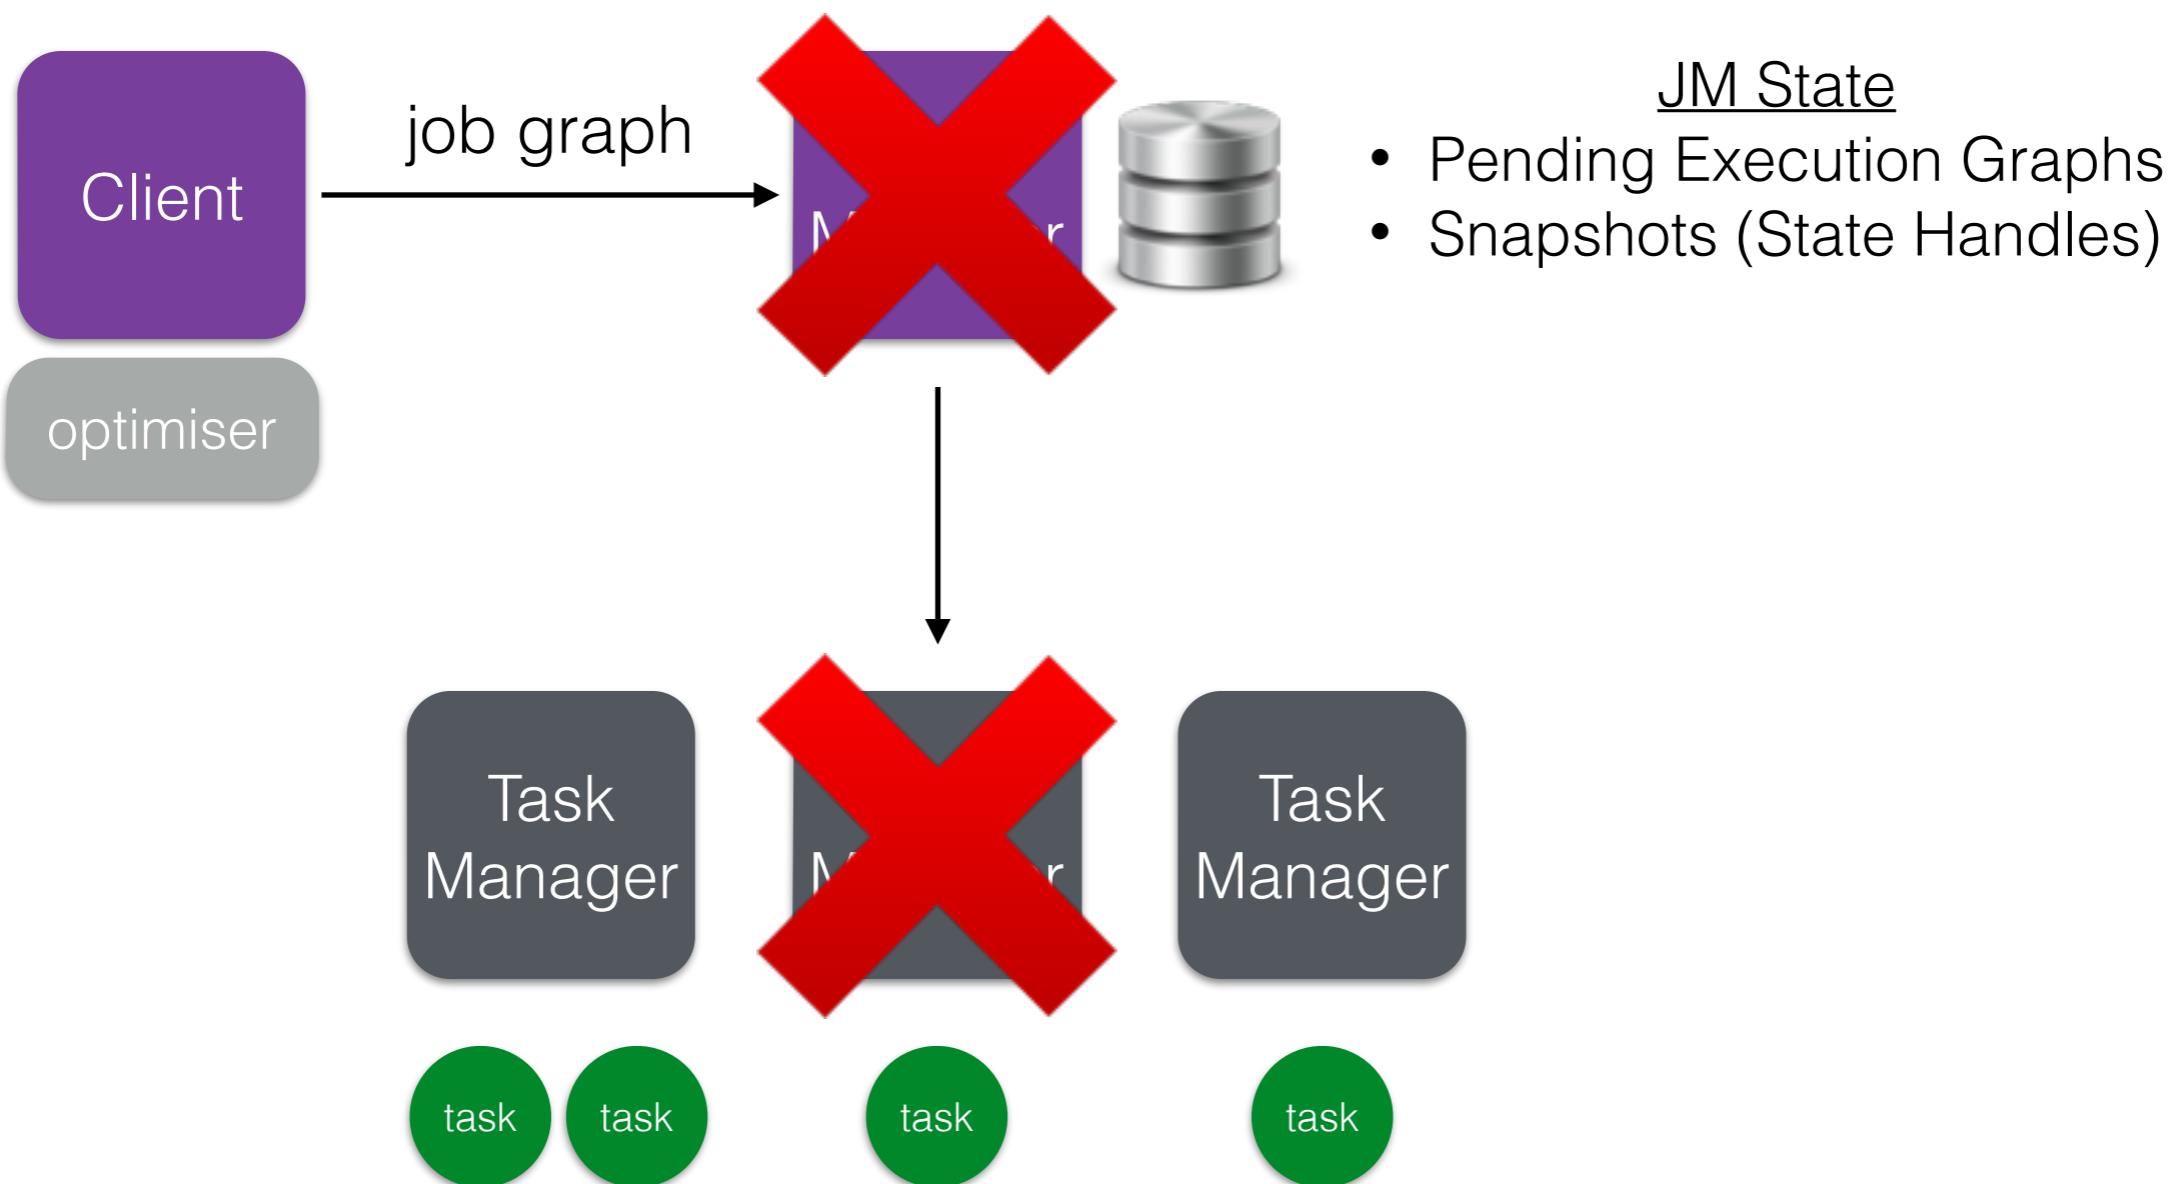
Node Failures



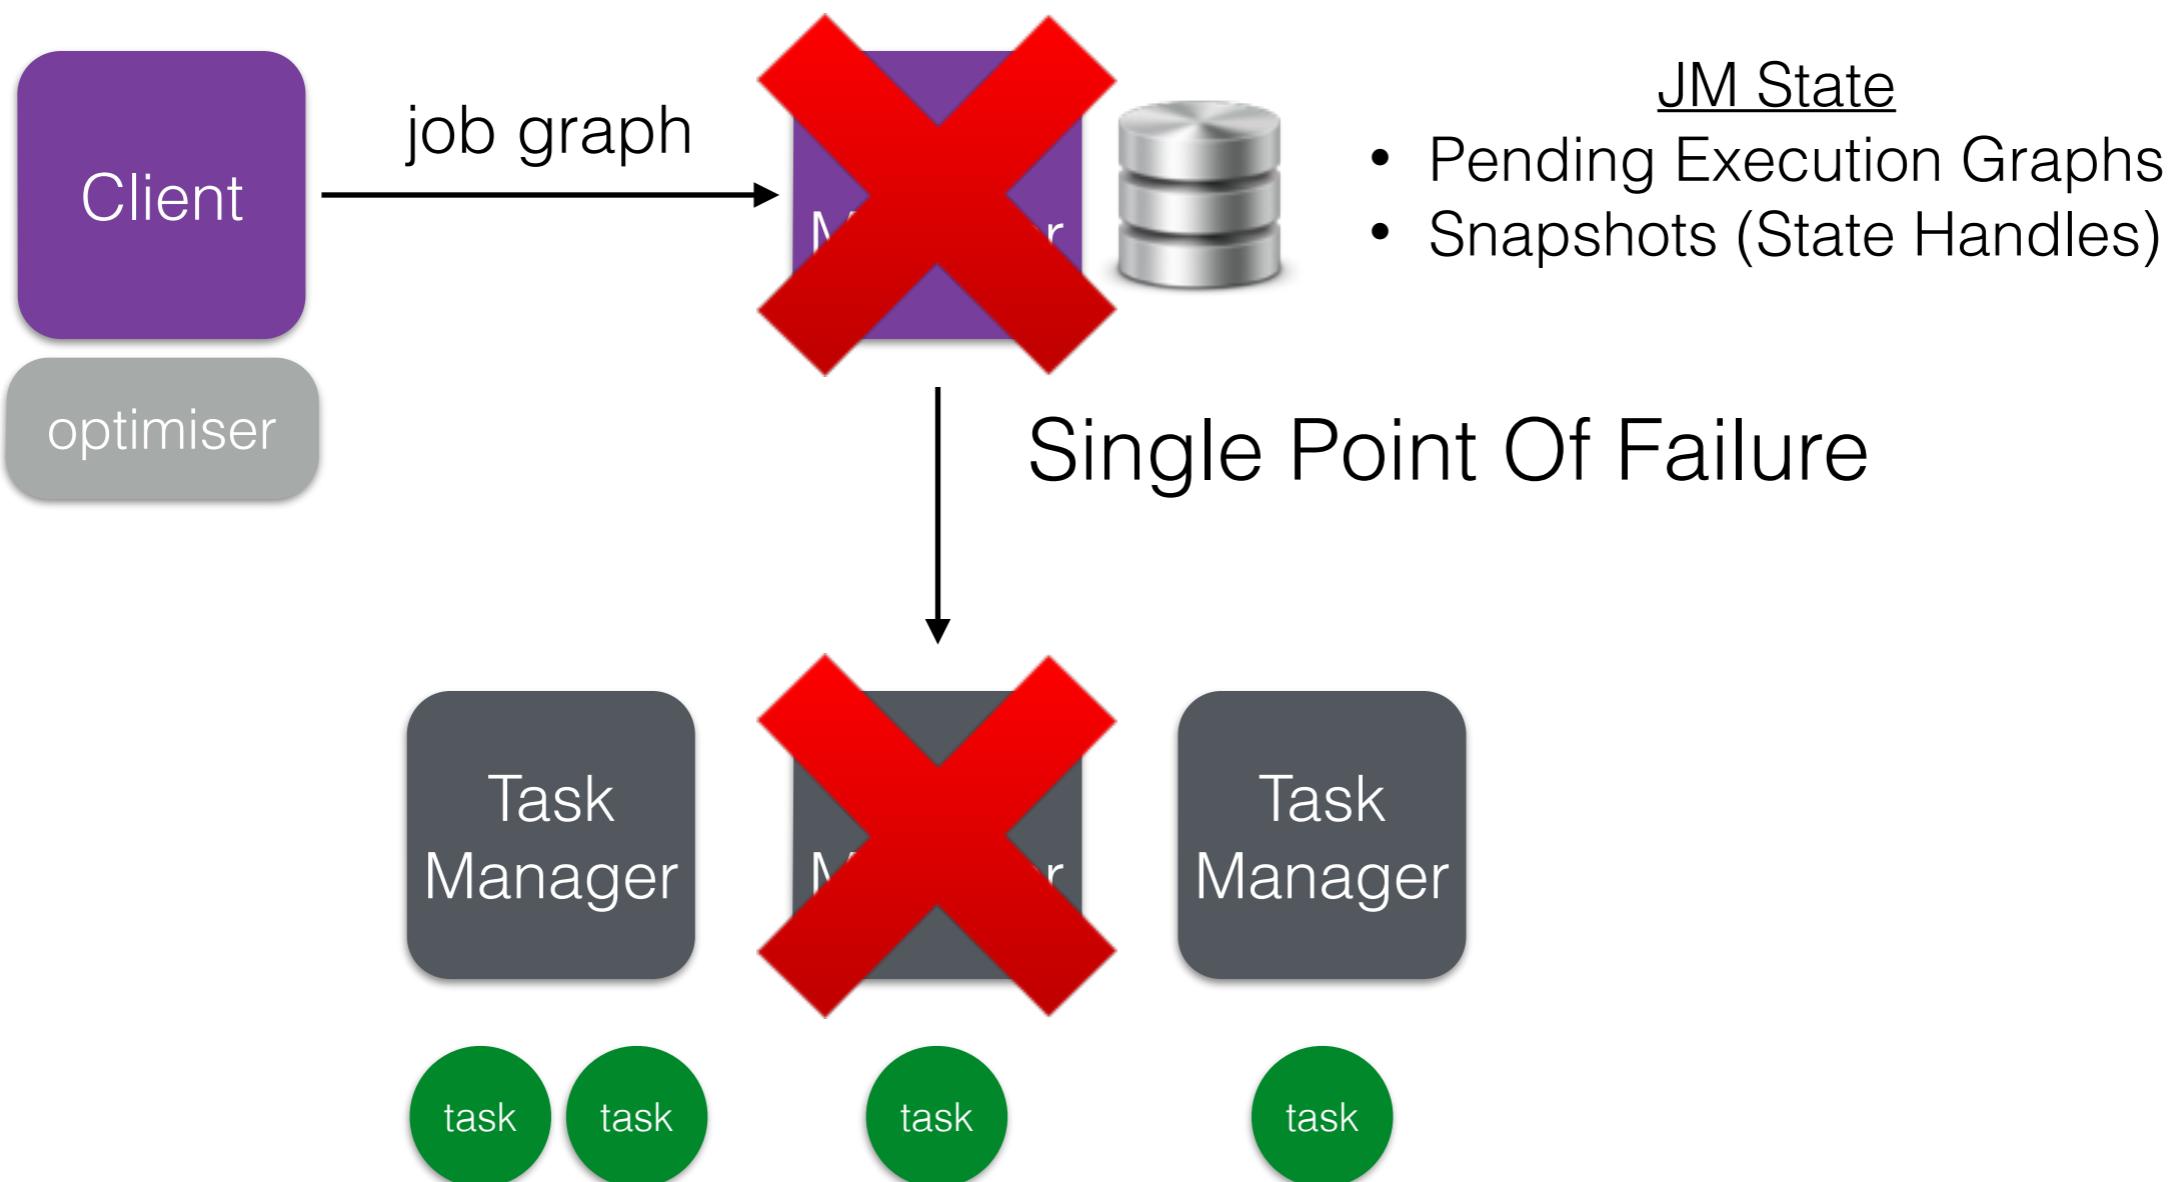
Node Failures



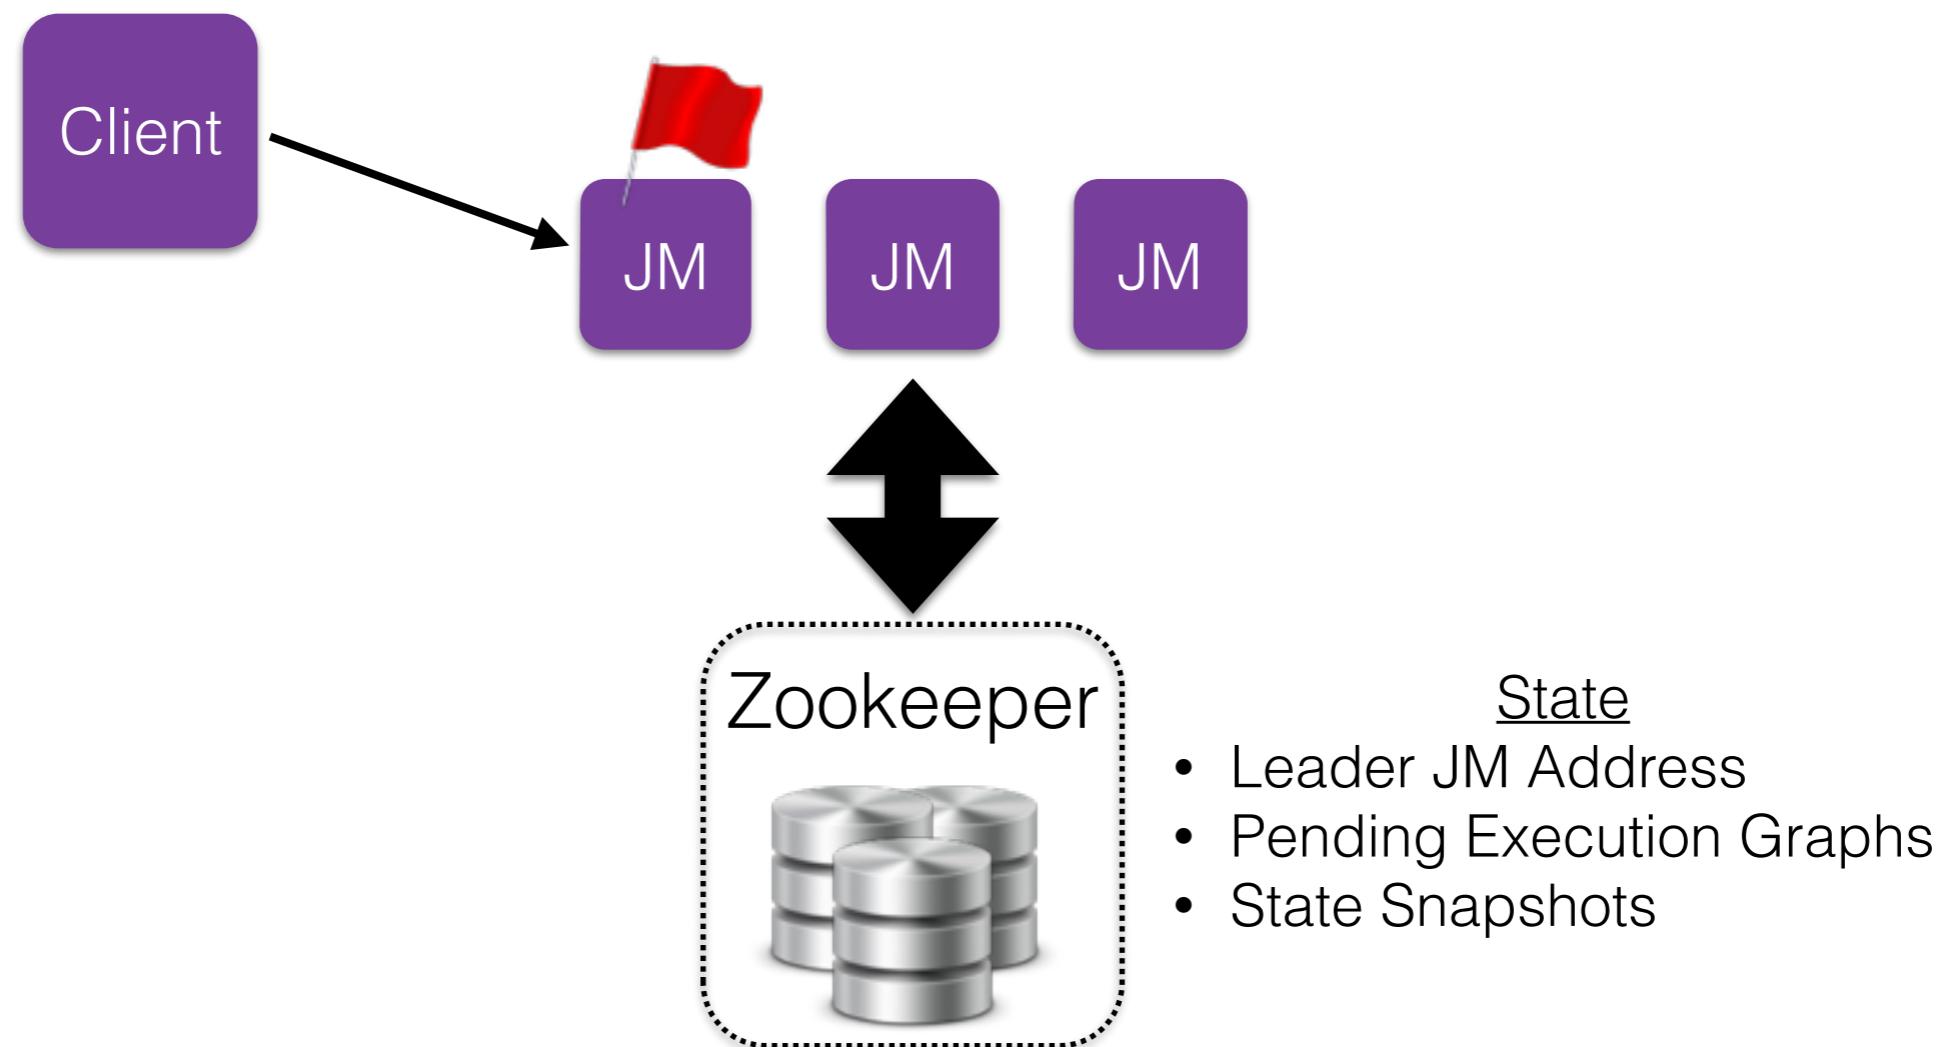
Node Failures



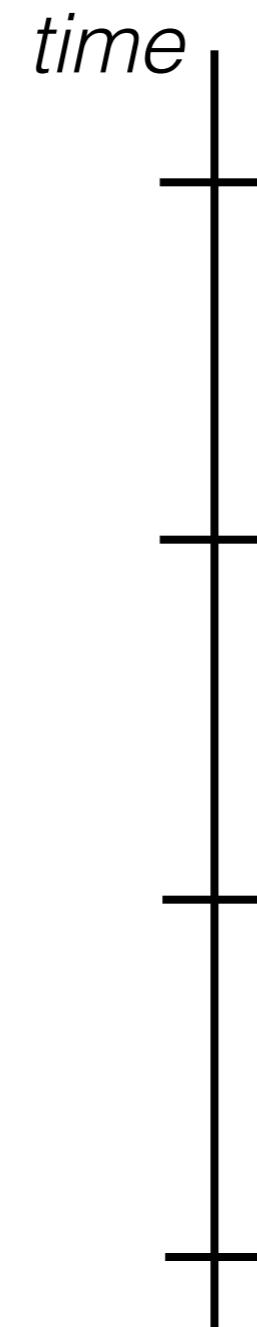
Node Failures



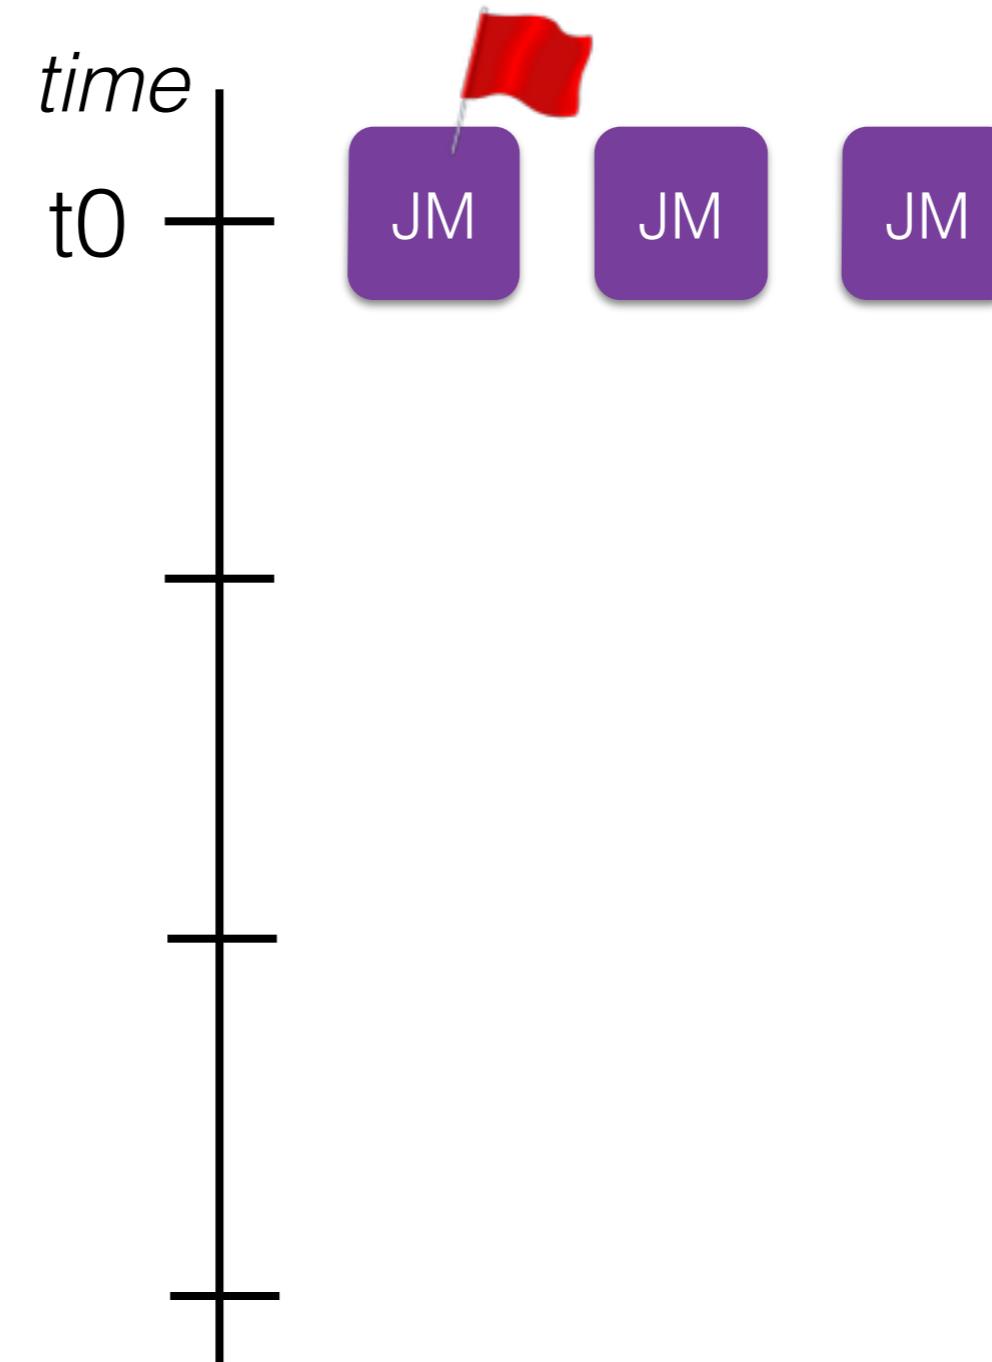
Passive Standby JM



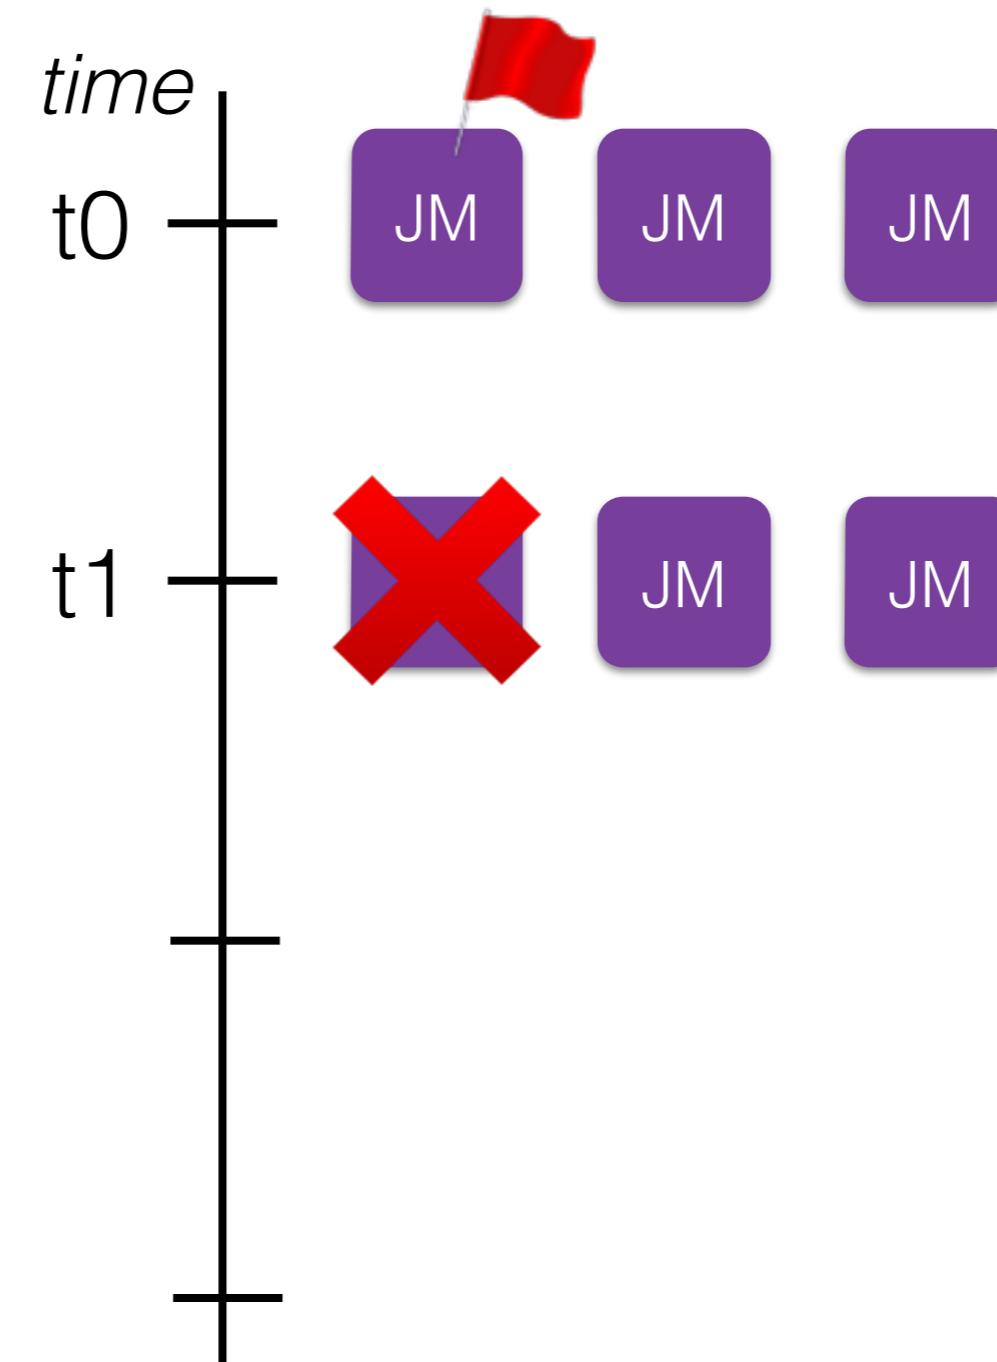
Eventual Leader Election



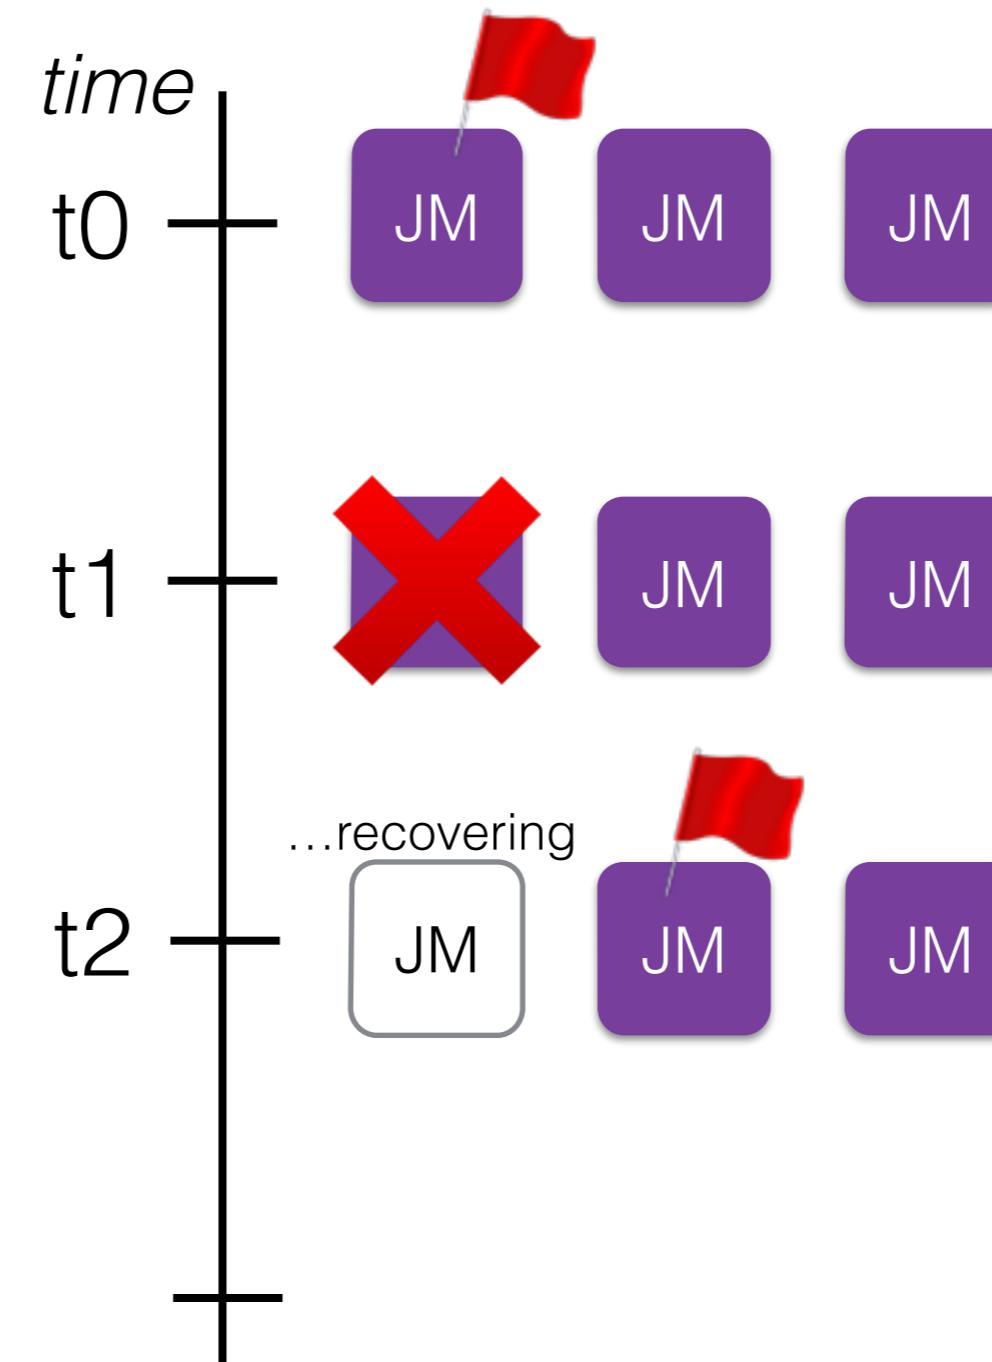
Eventual Leader Election



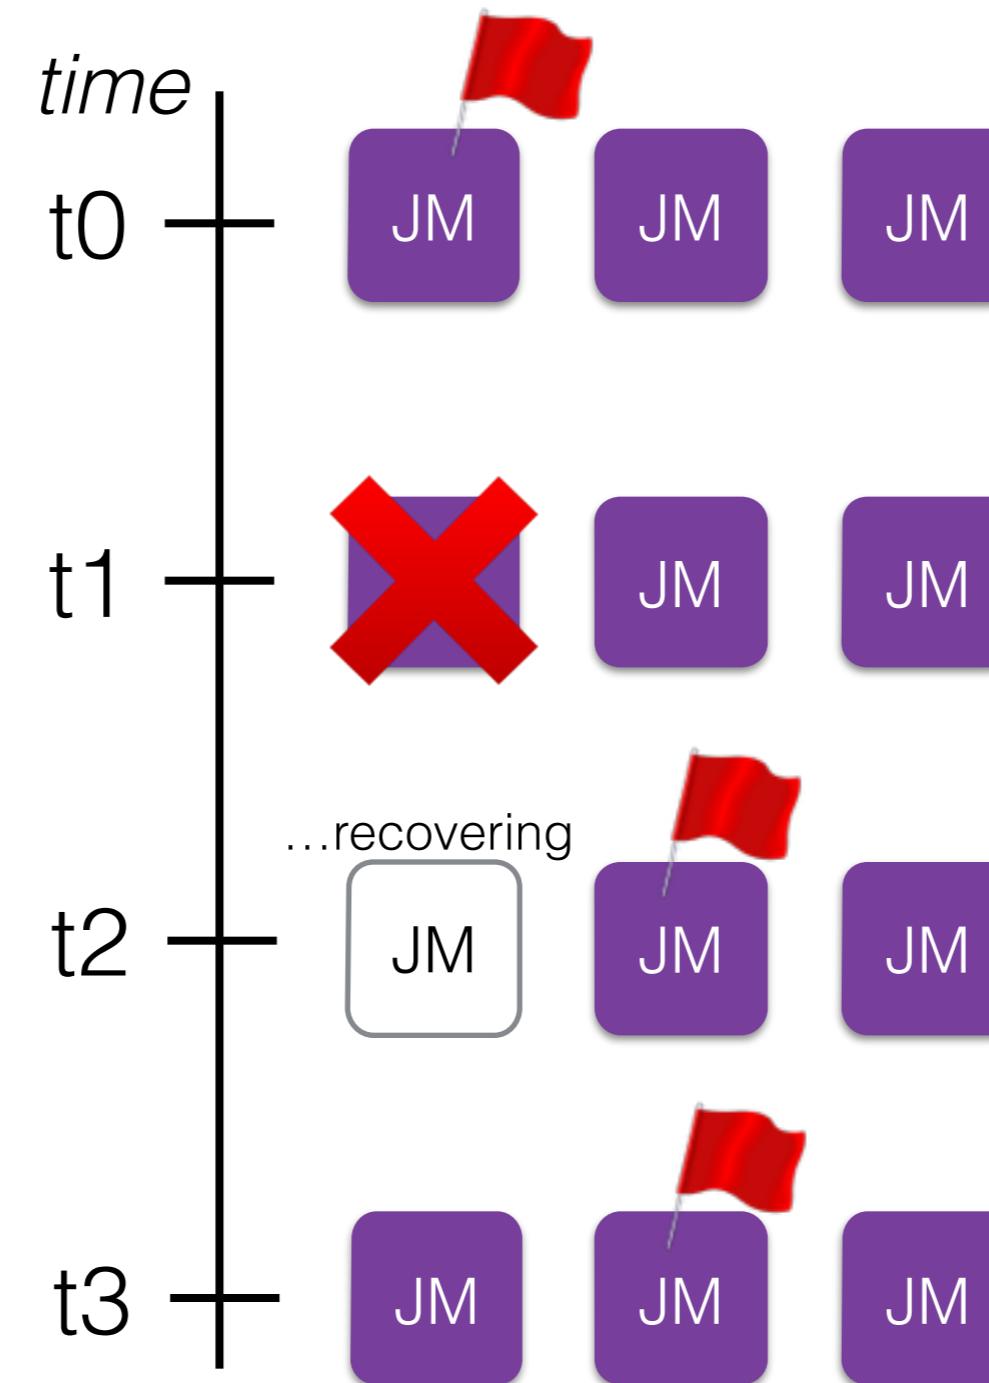
Eventual Leader Election



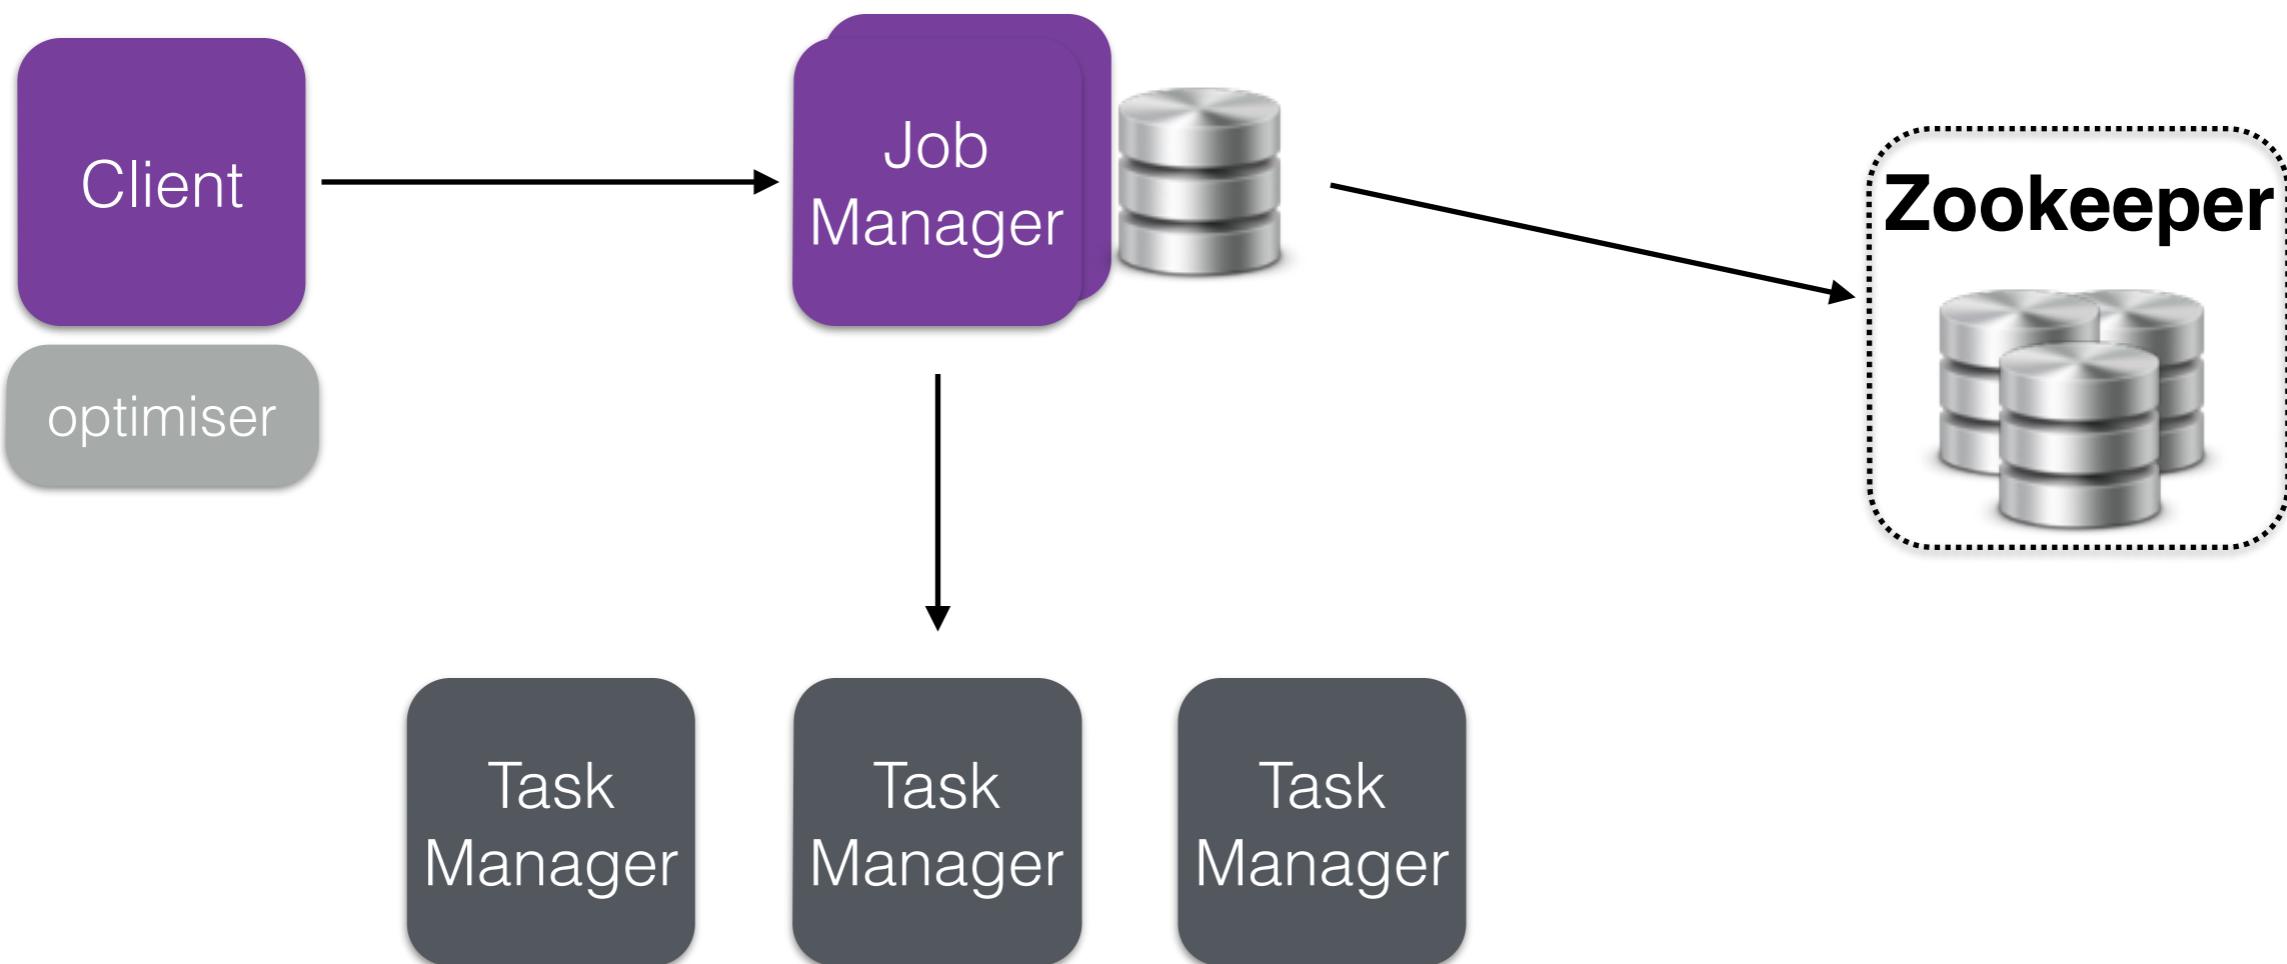
Eventual Leader Election



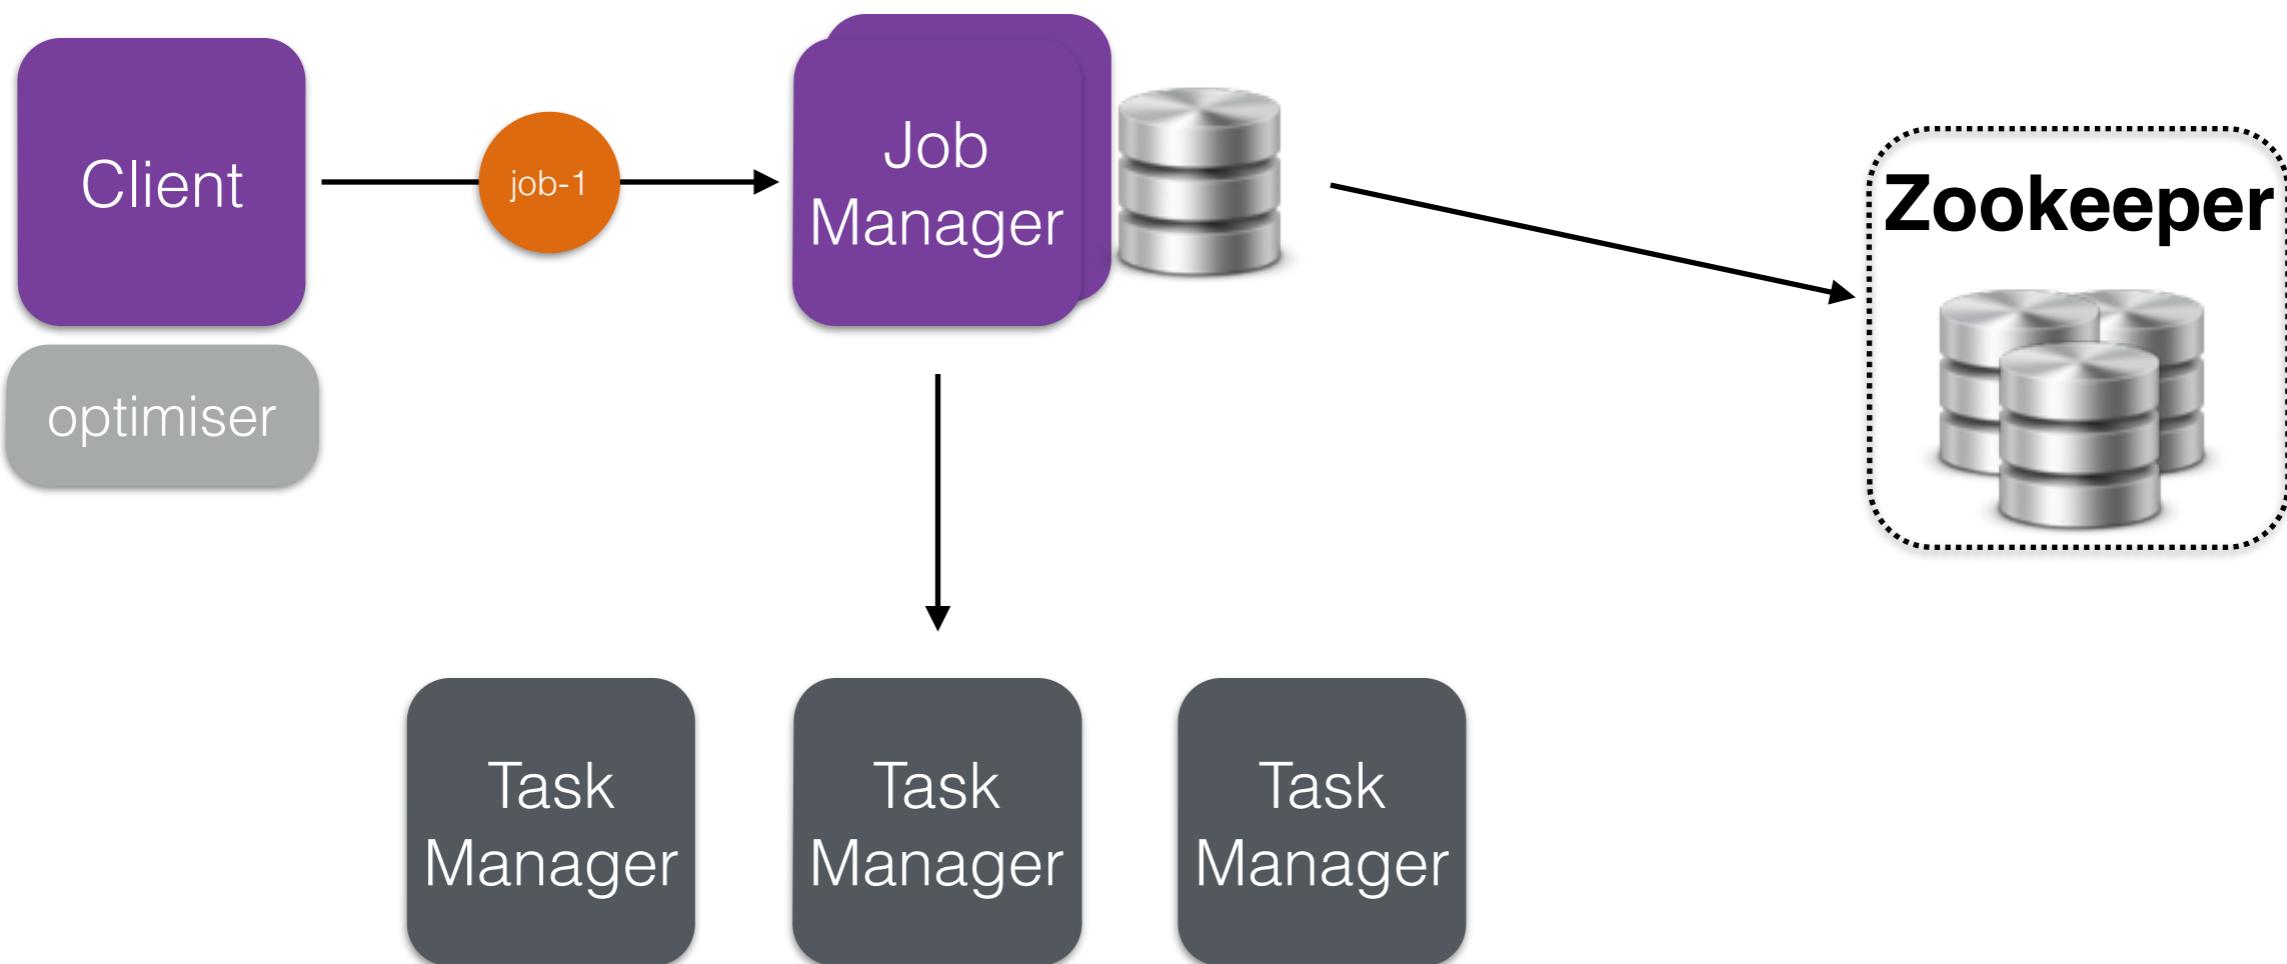
Eventual Leader Election



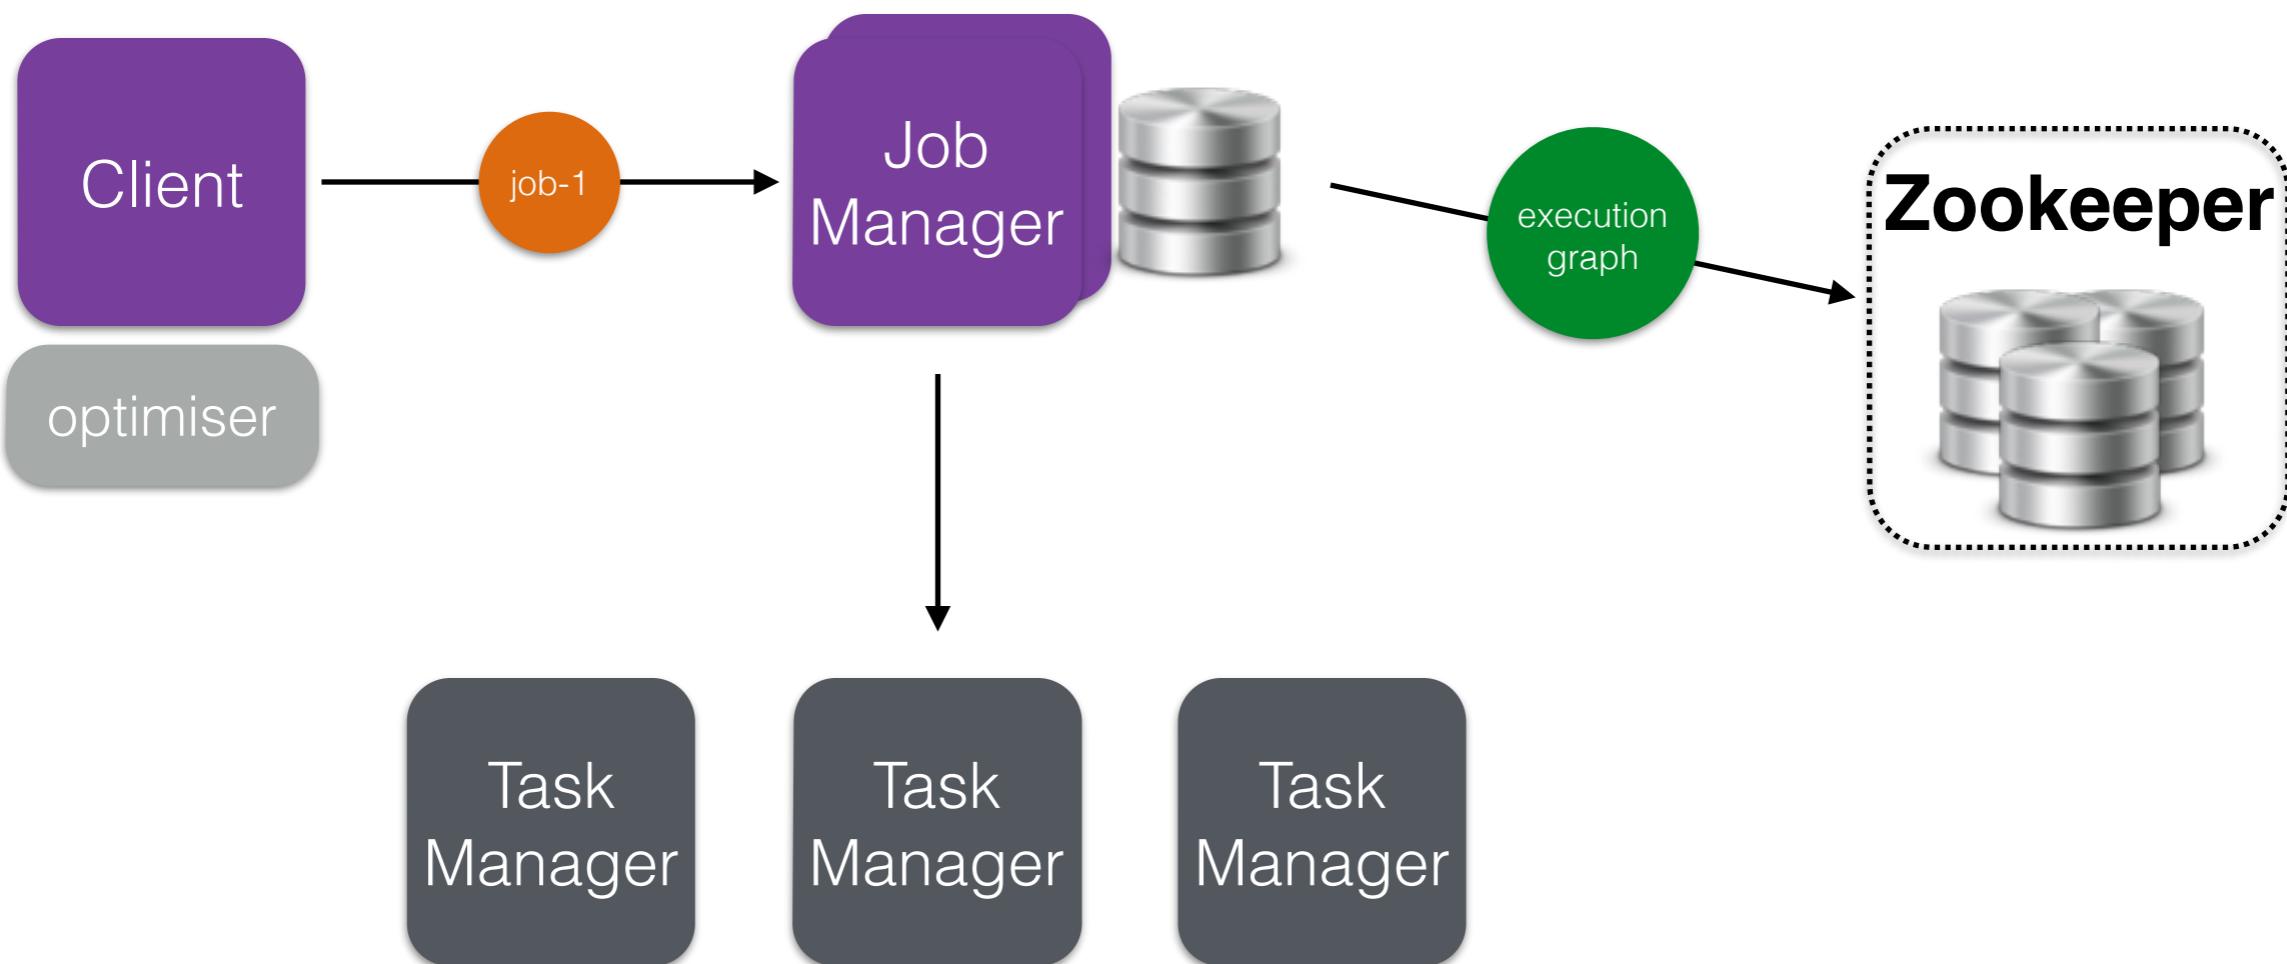
Scheduling Jobs



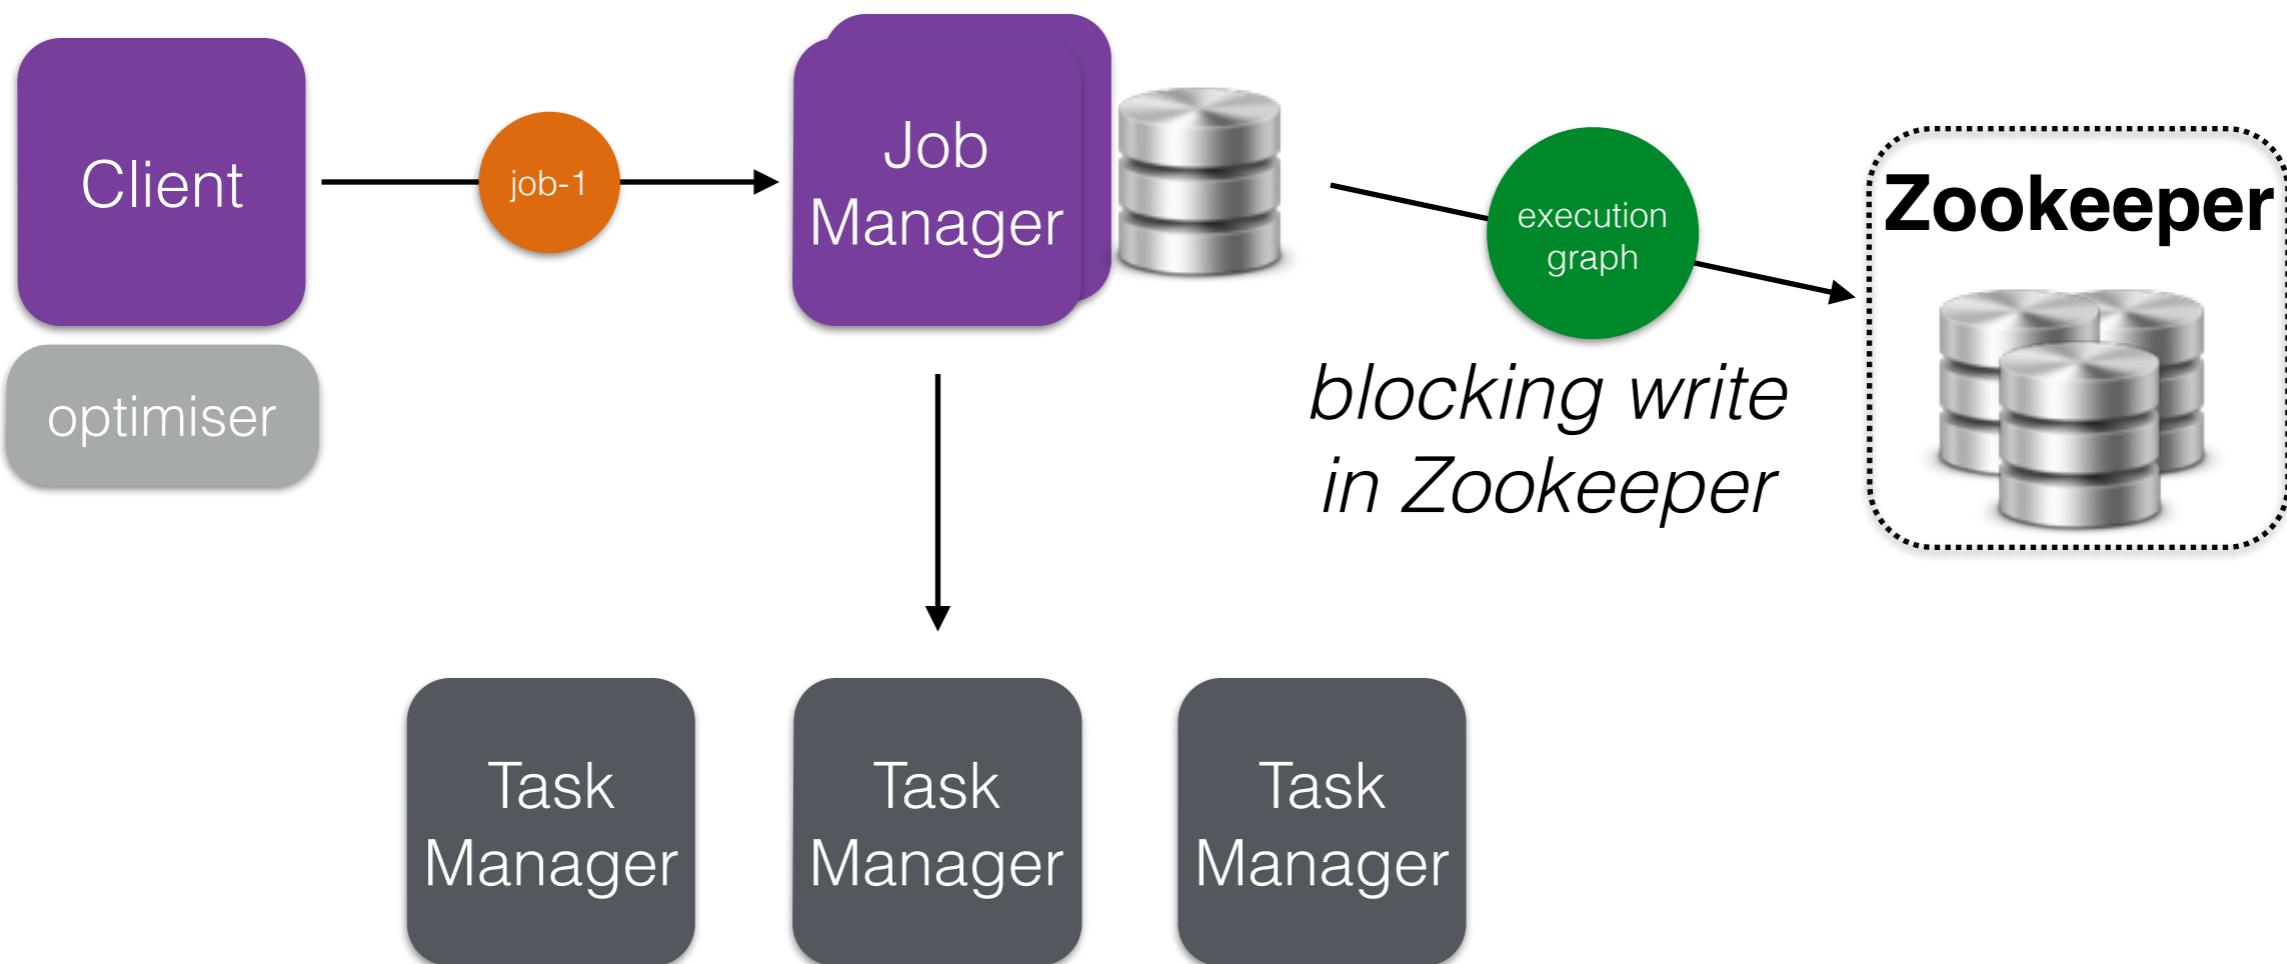
Scheduling Jobs



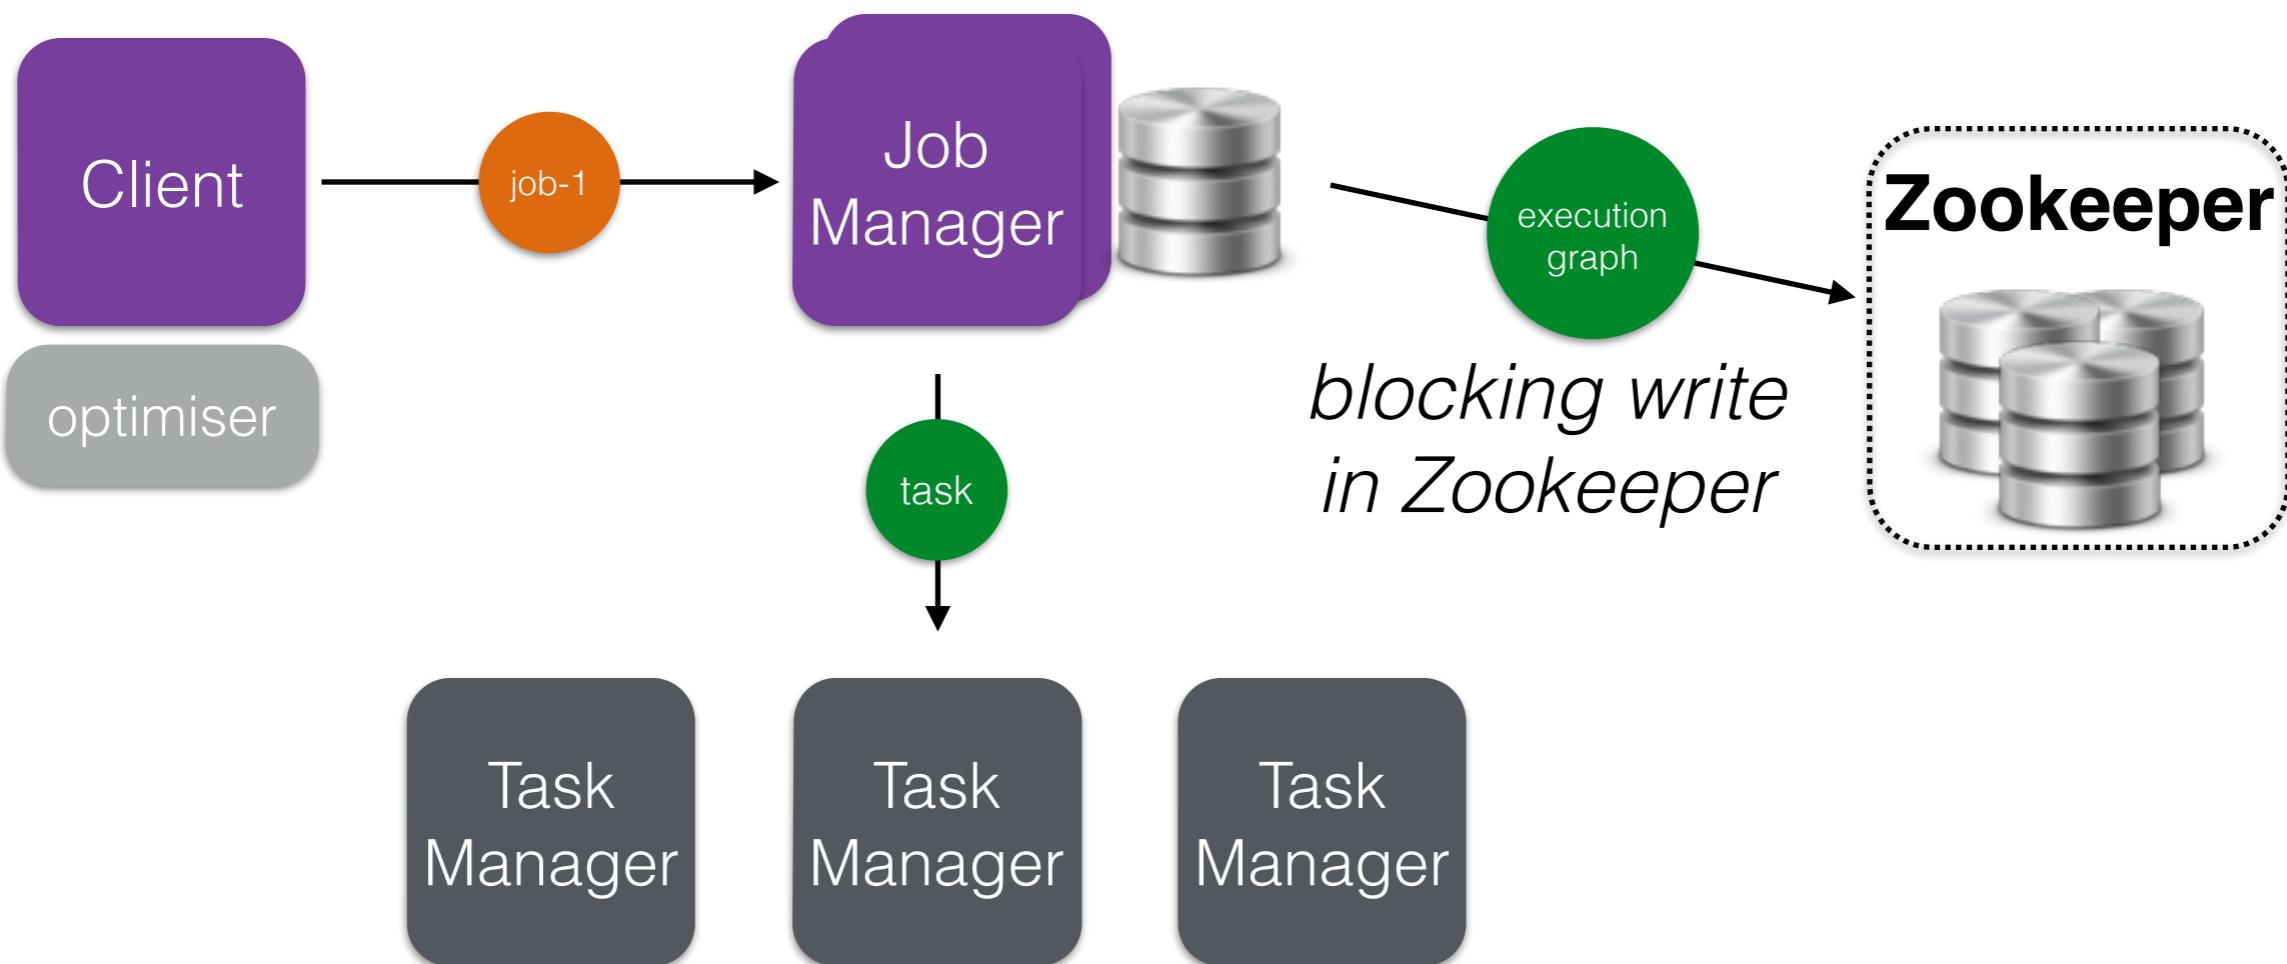
Scheduling Jobs



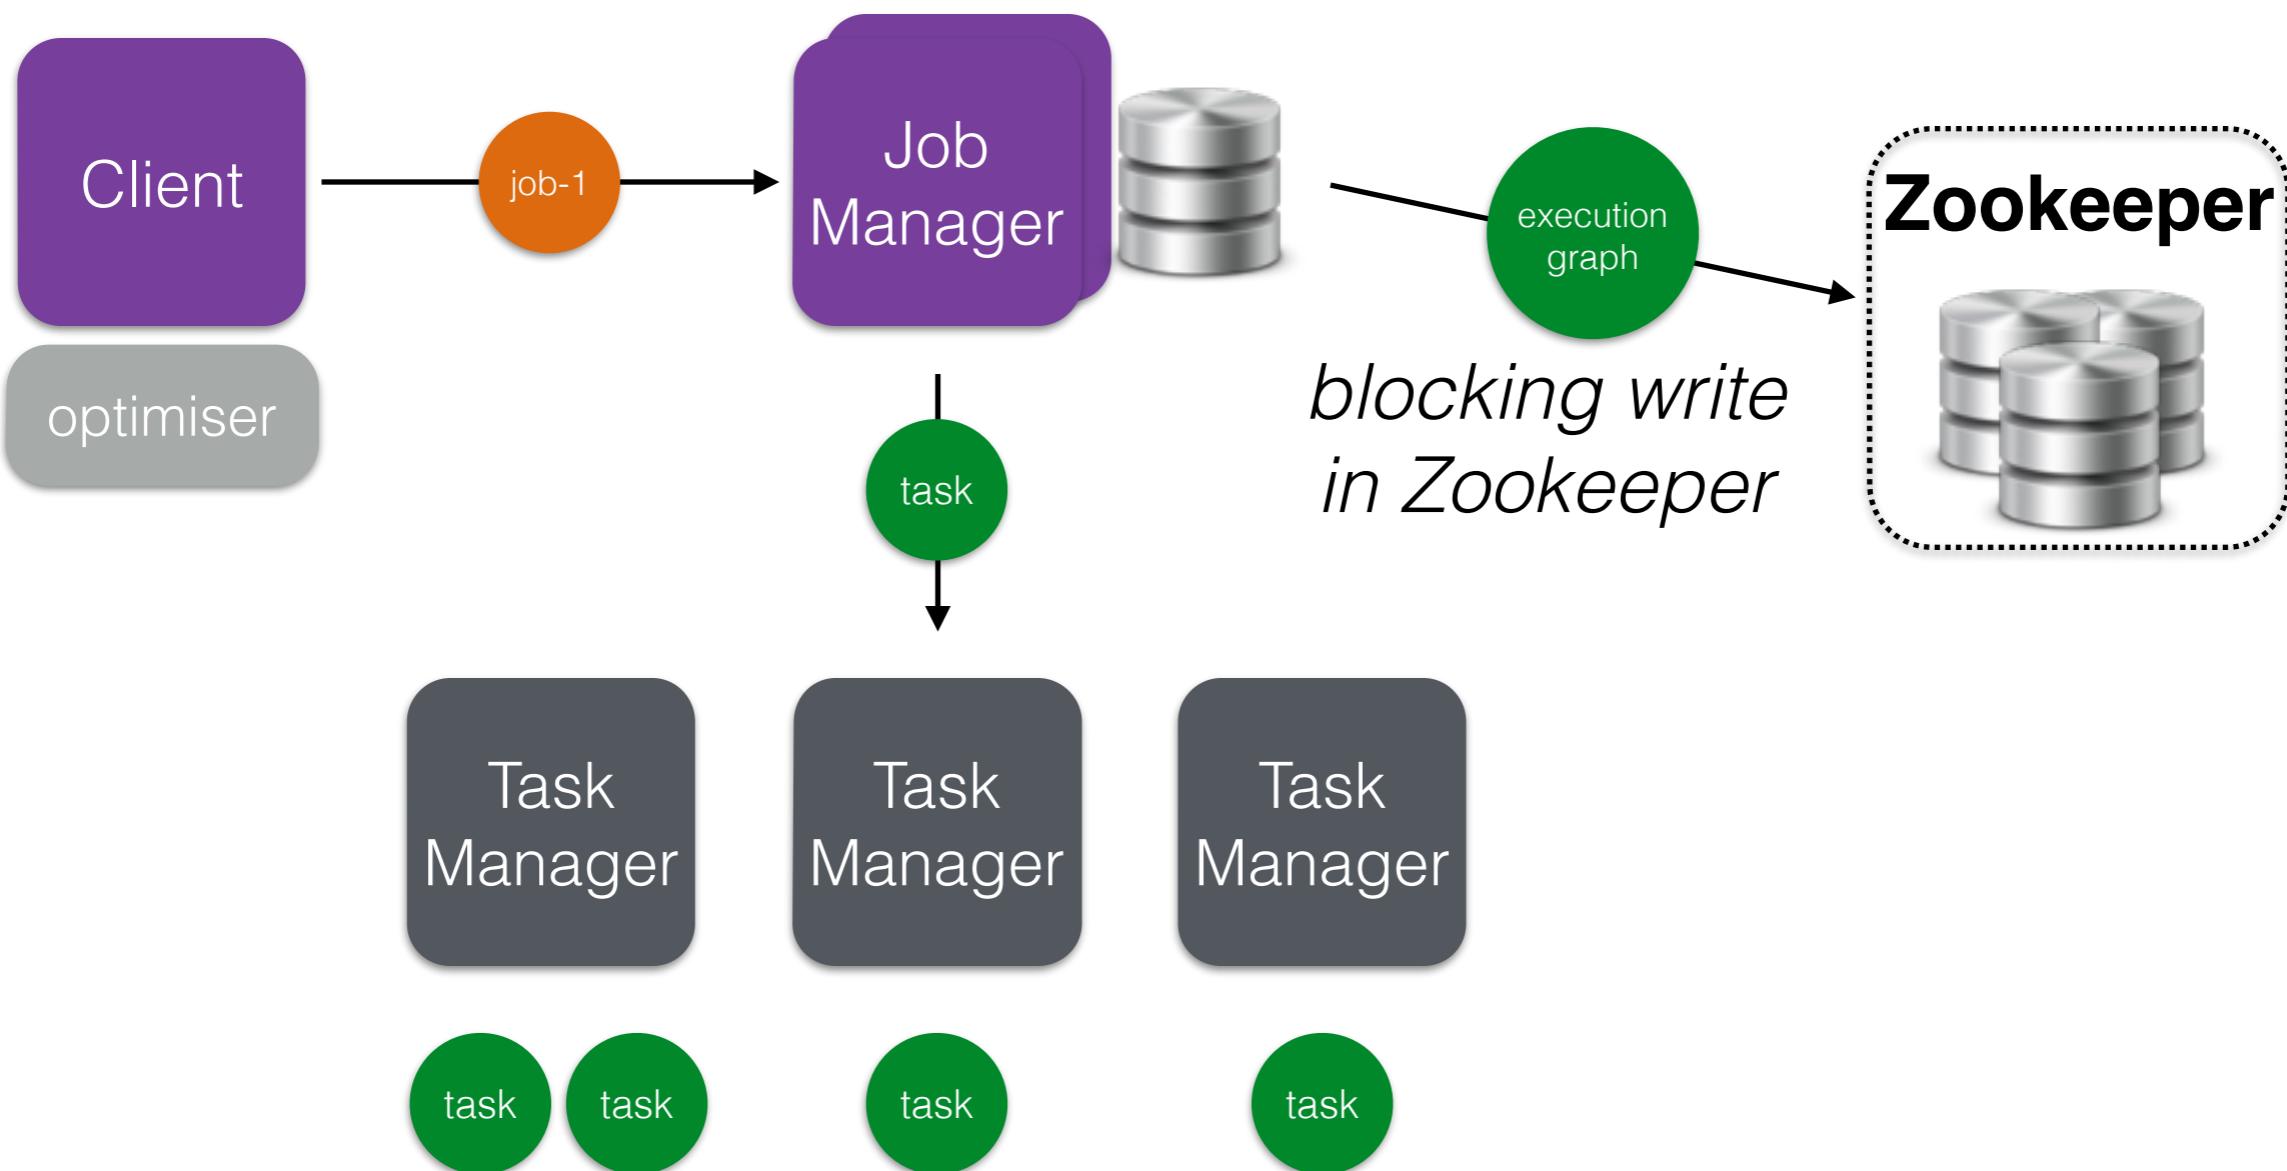
Scheduling Jobs



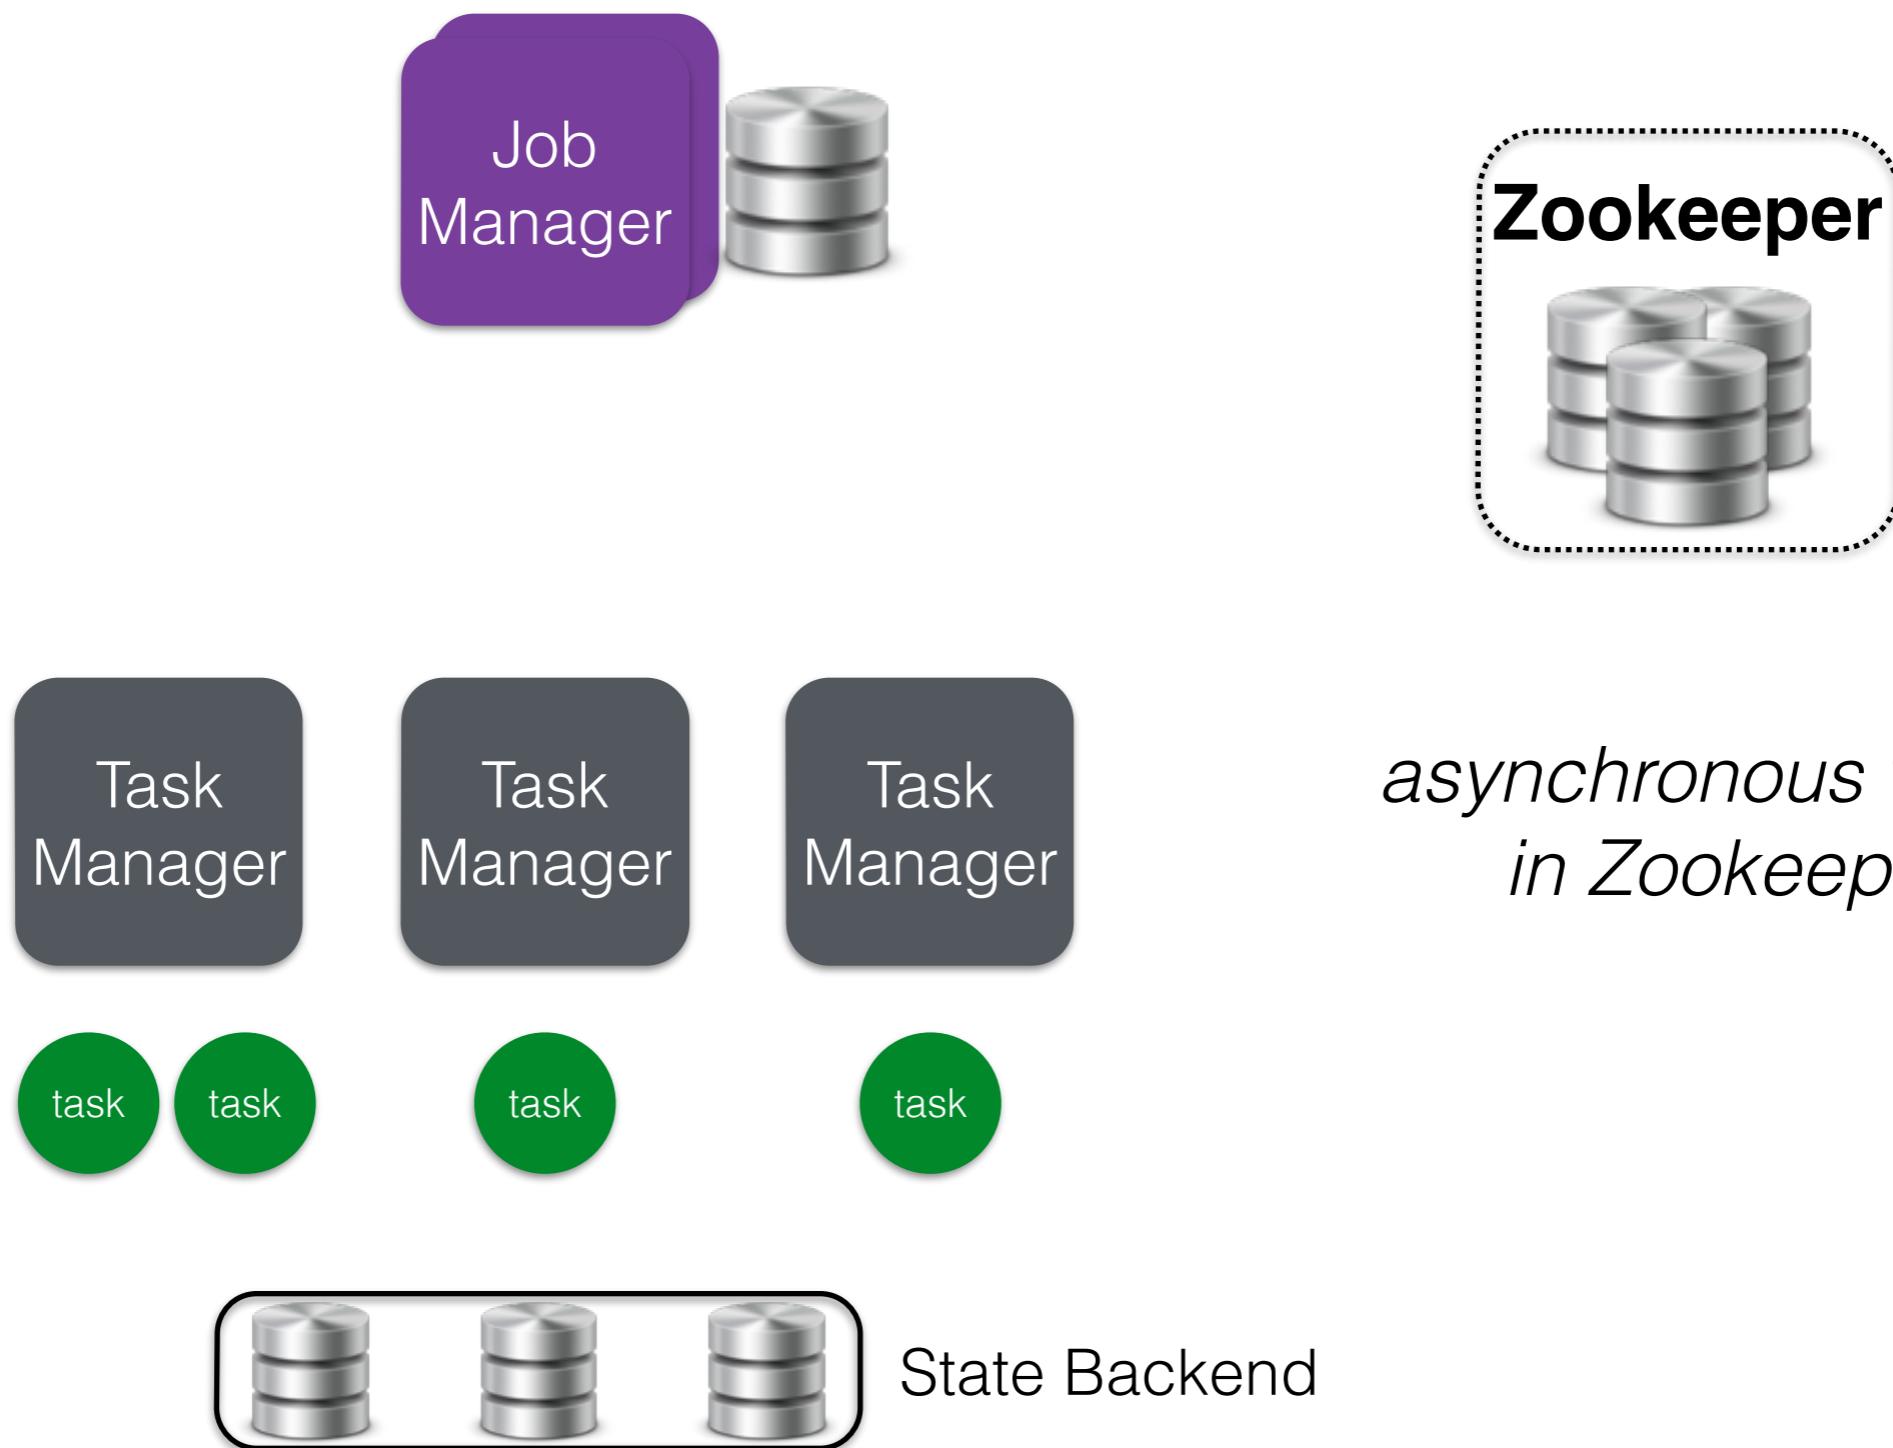
Scheduling Jobs



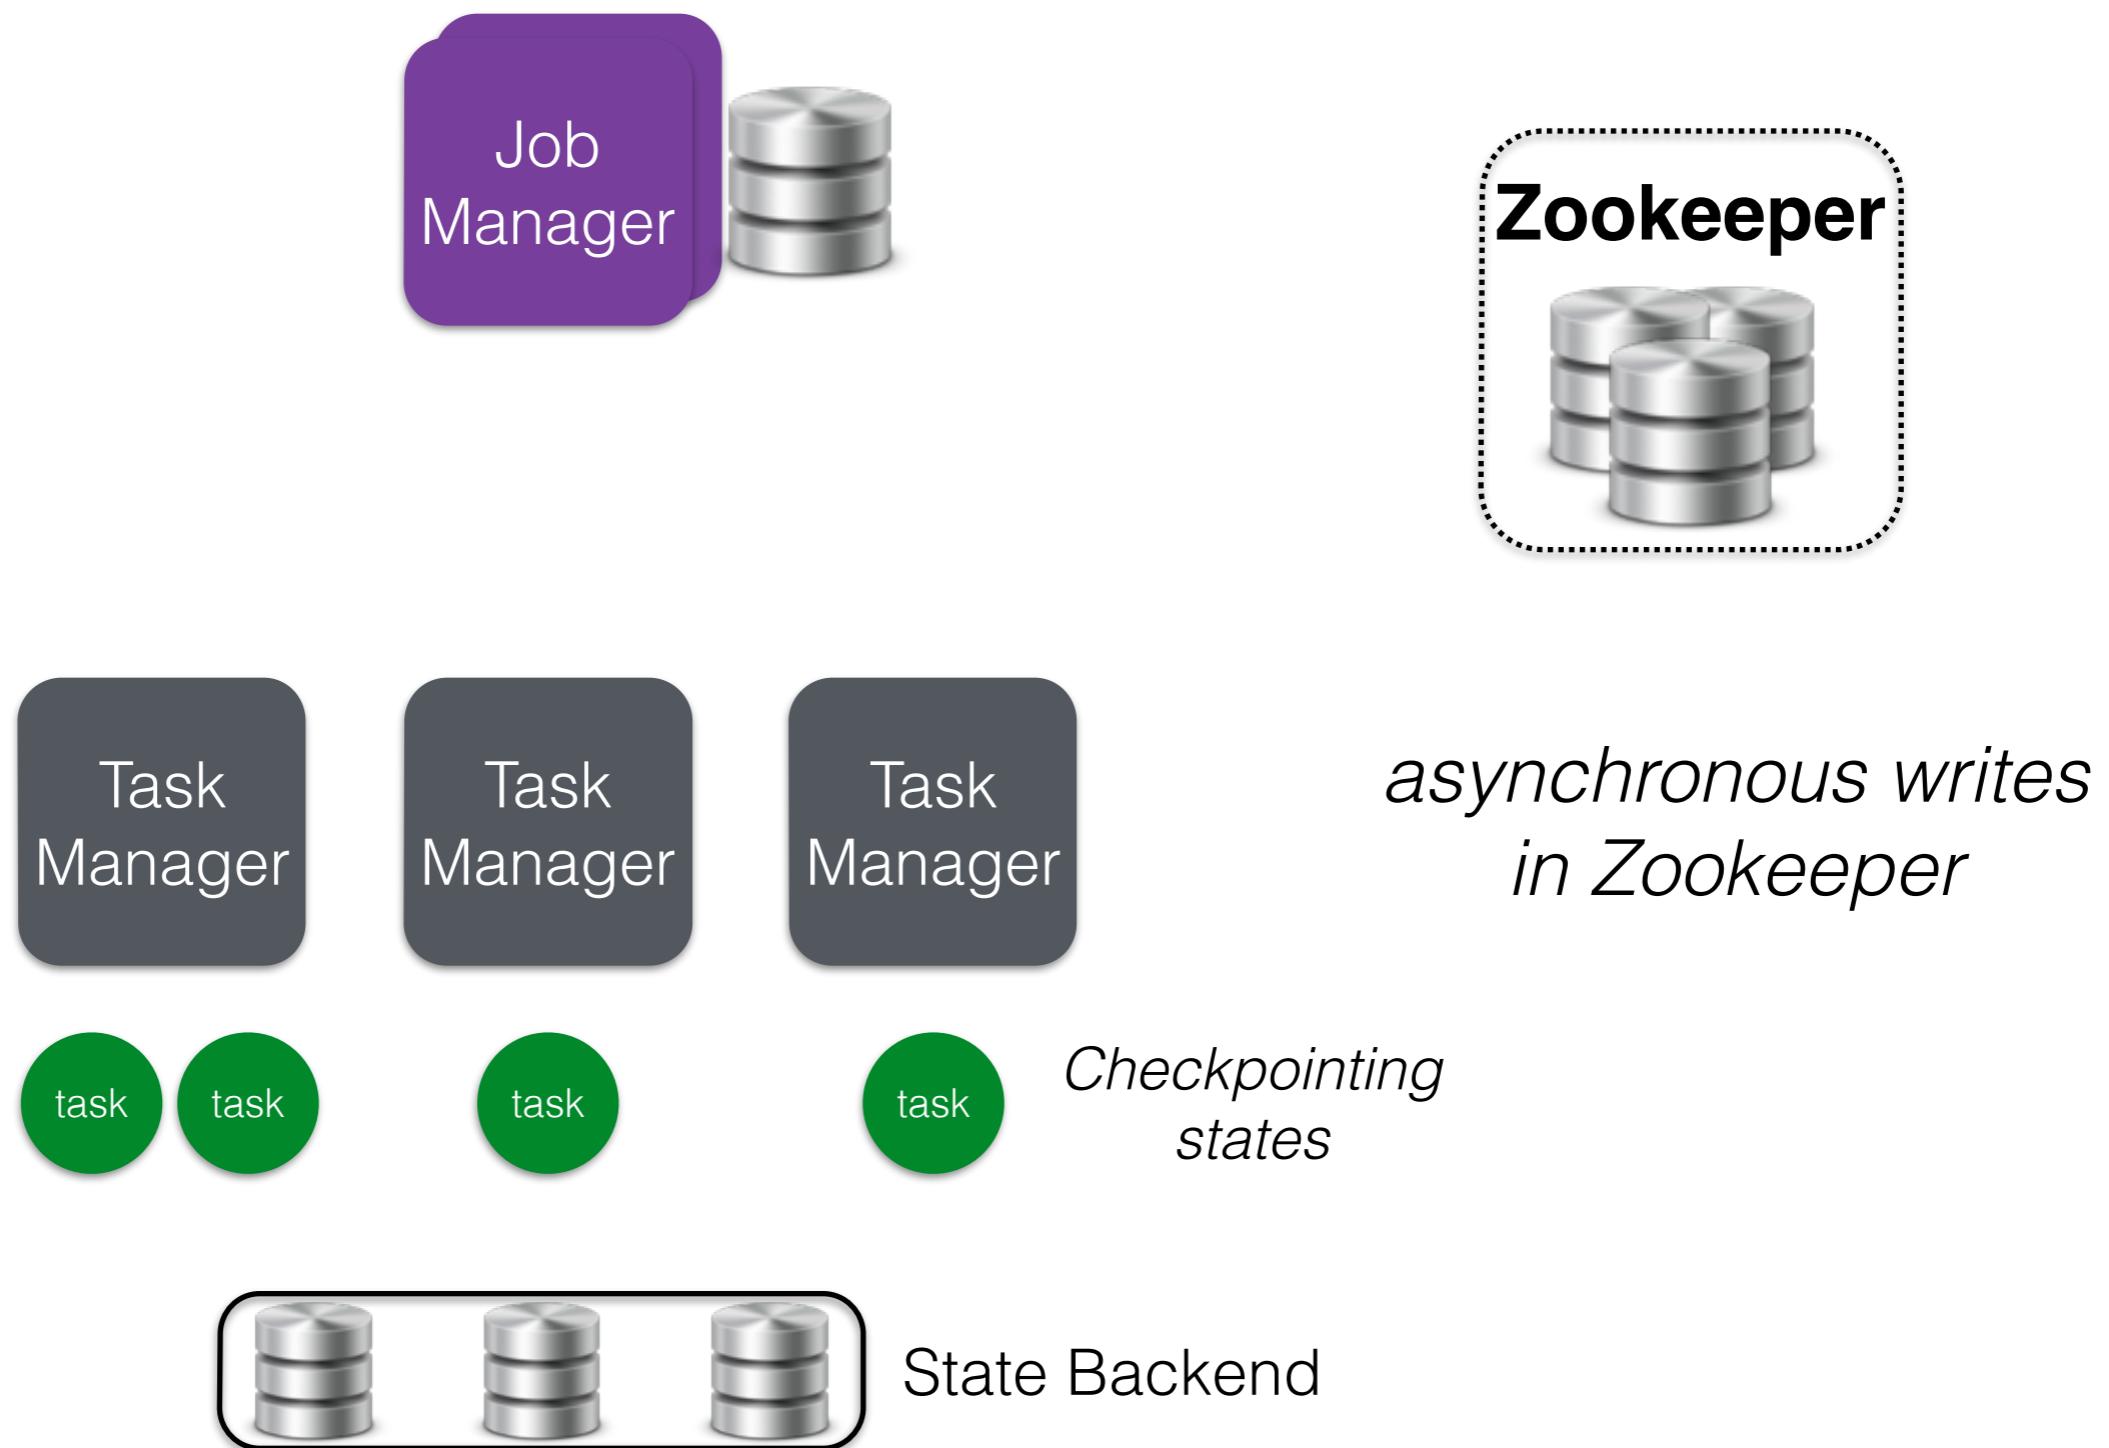
Scheduling Jobs



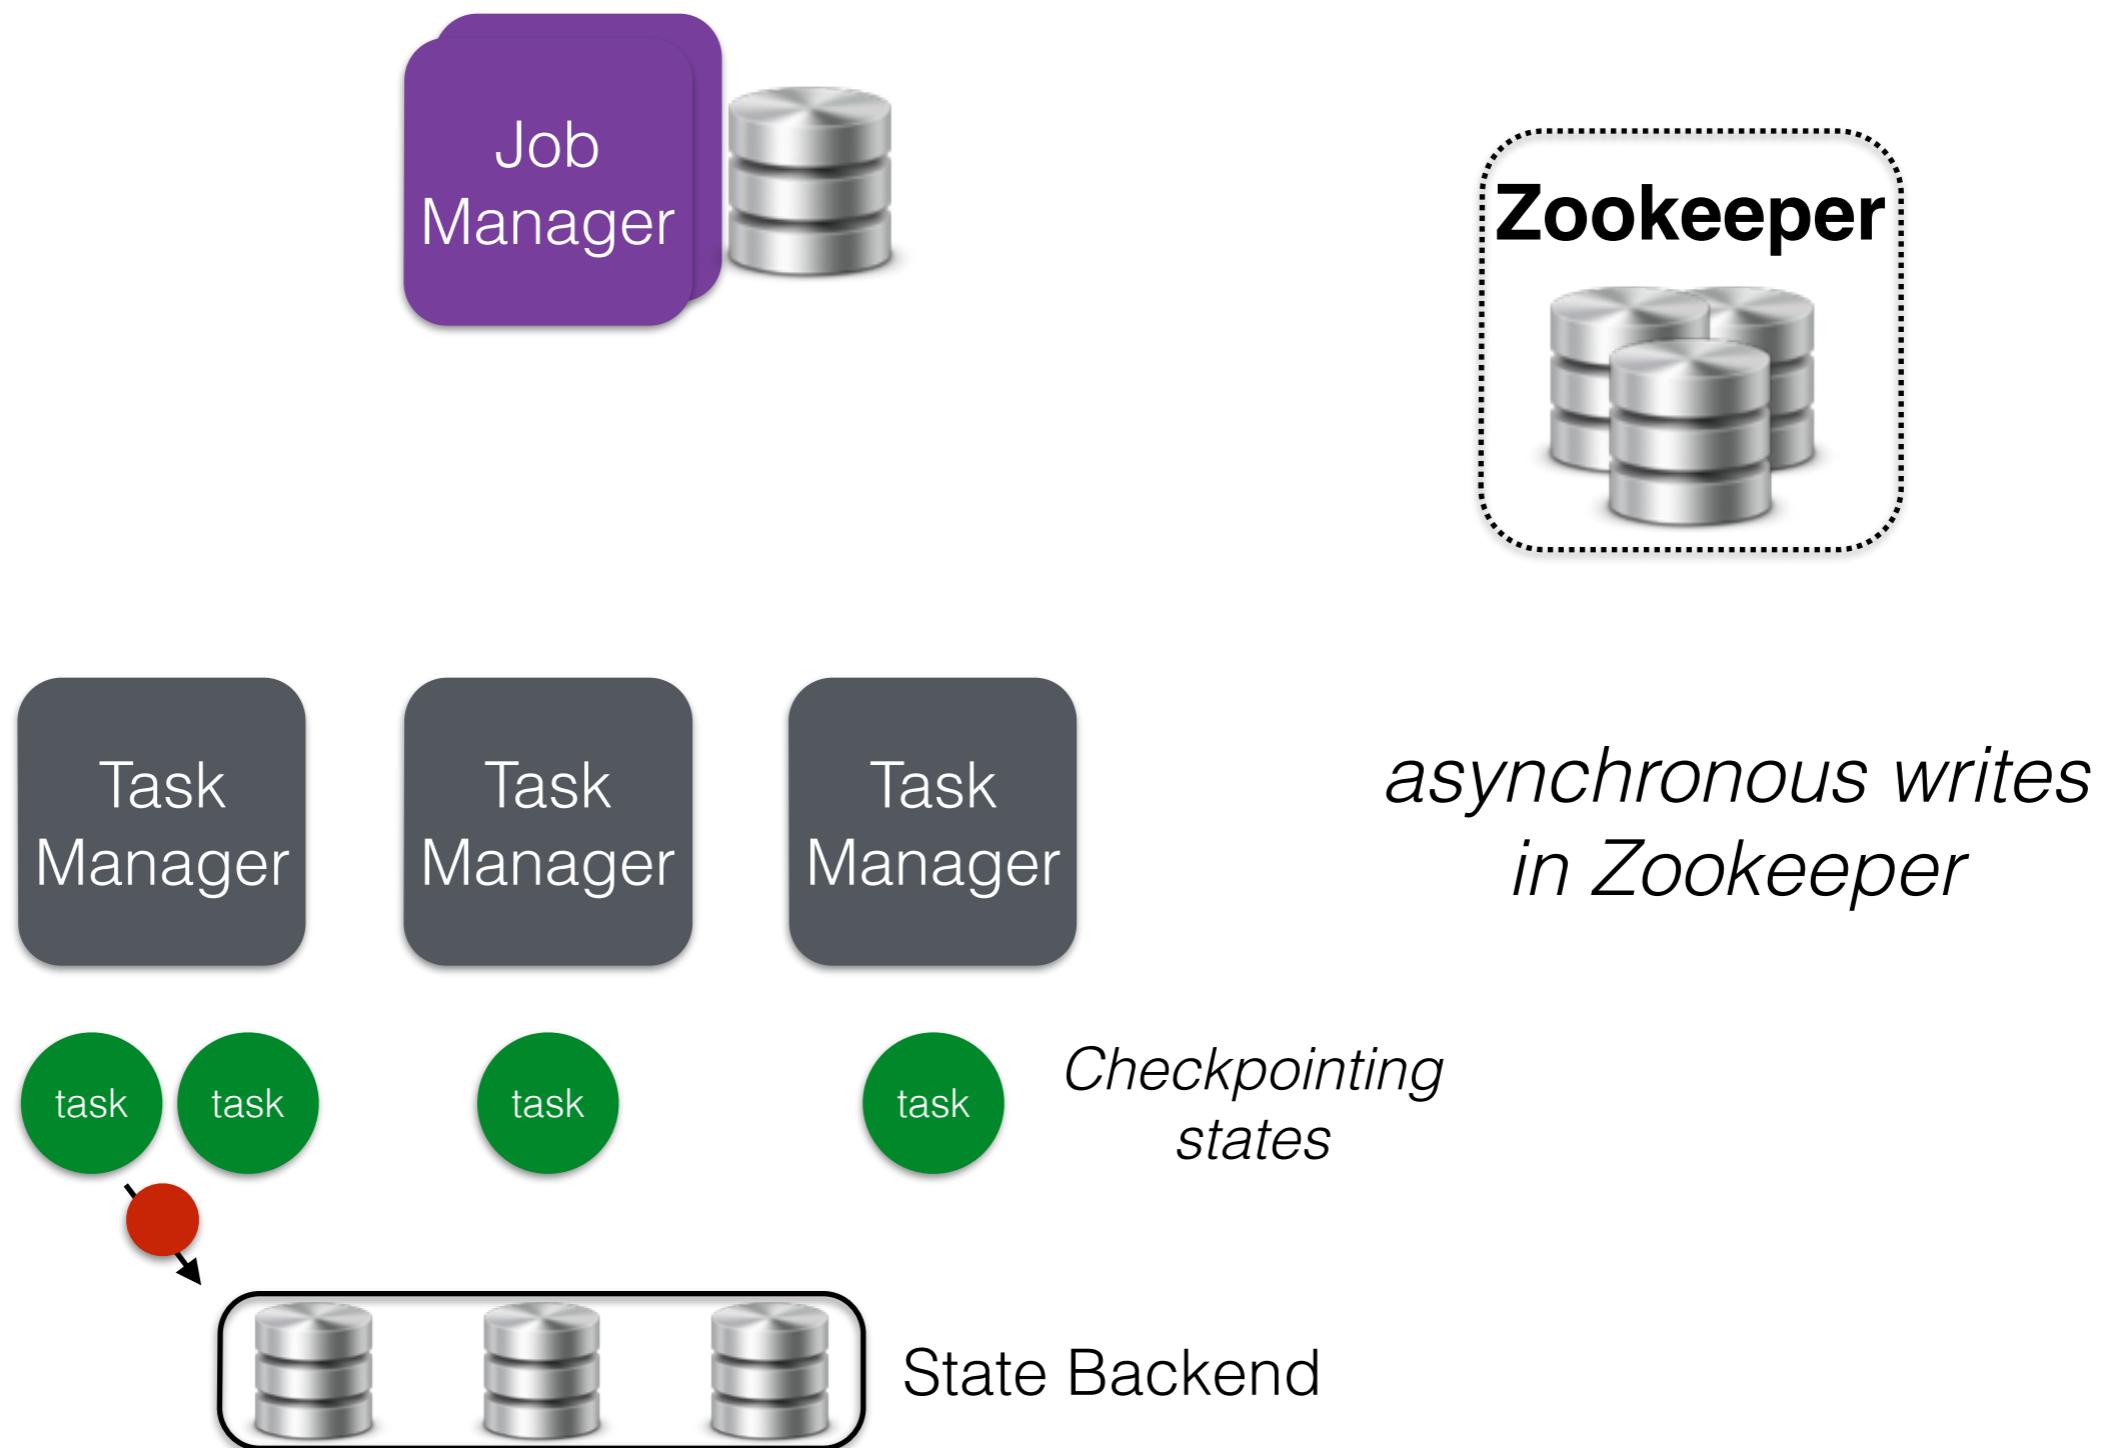
Logging Snapshots



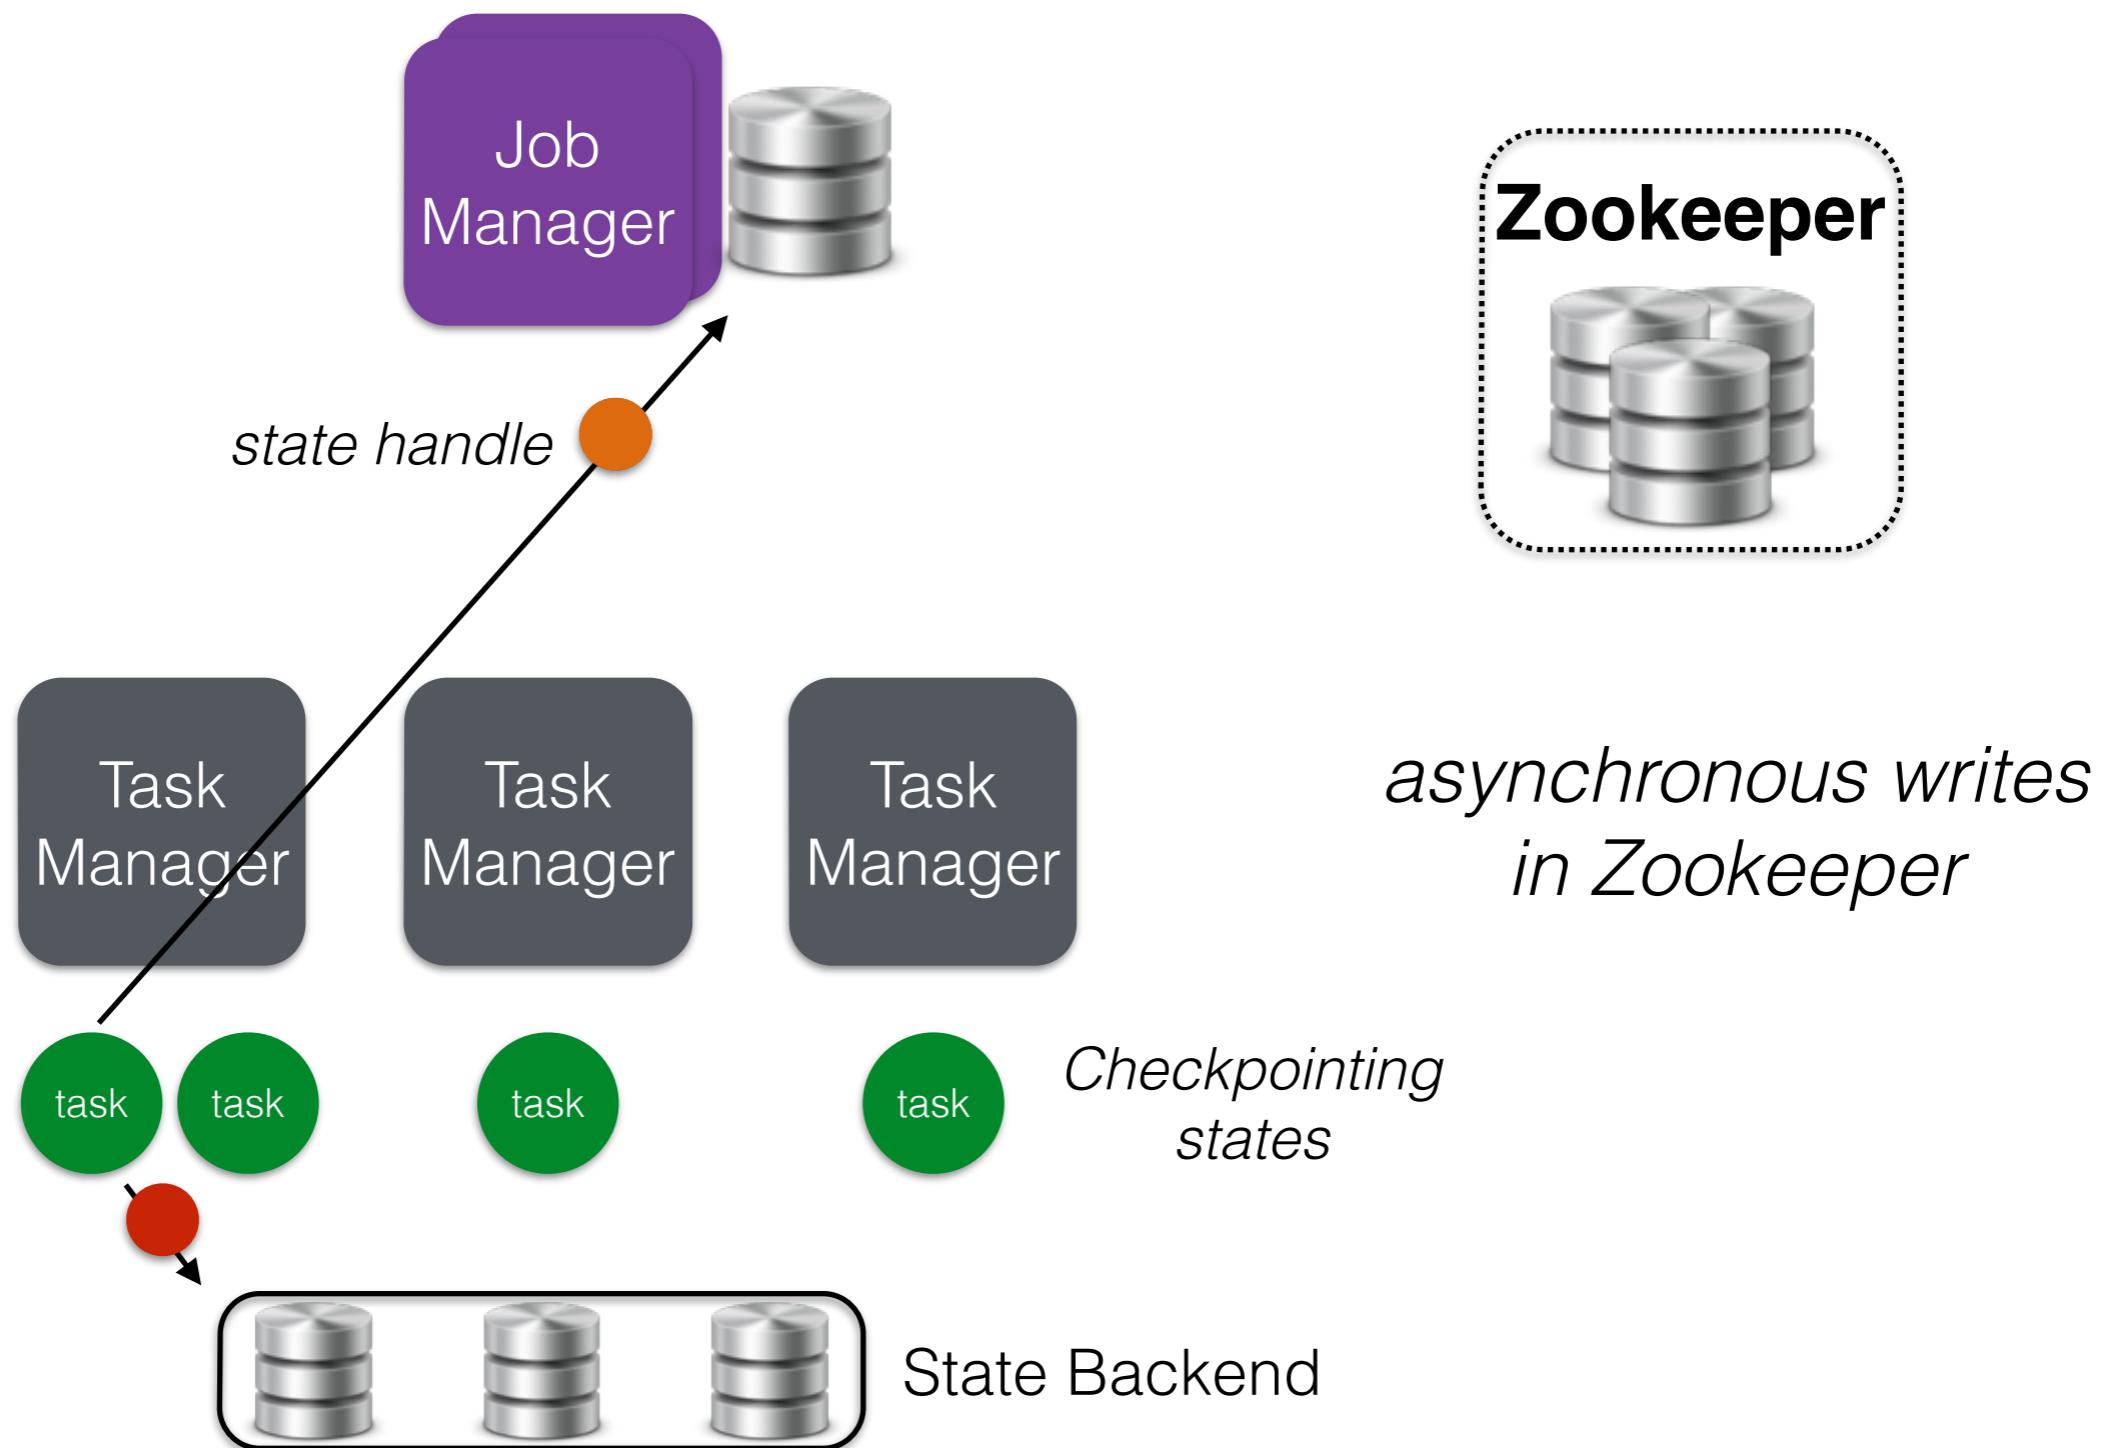
Logging Snapshots



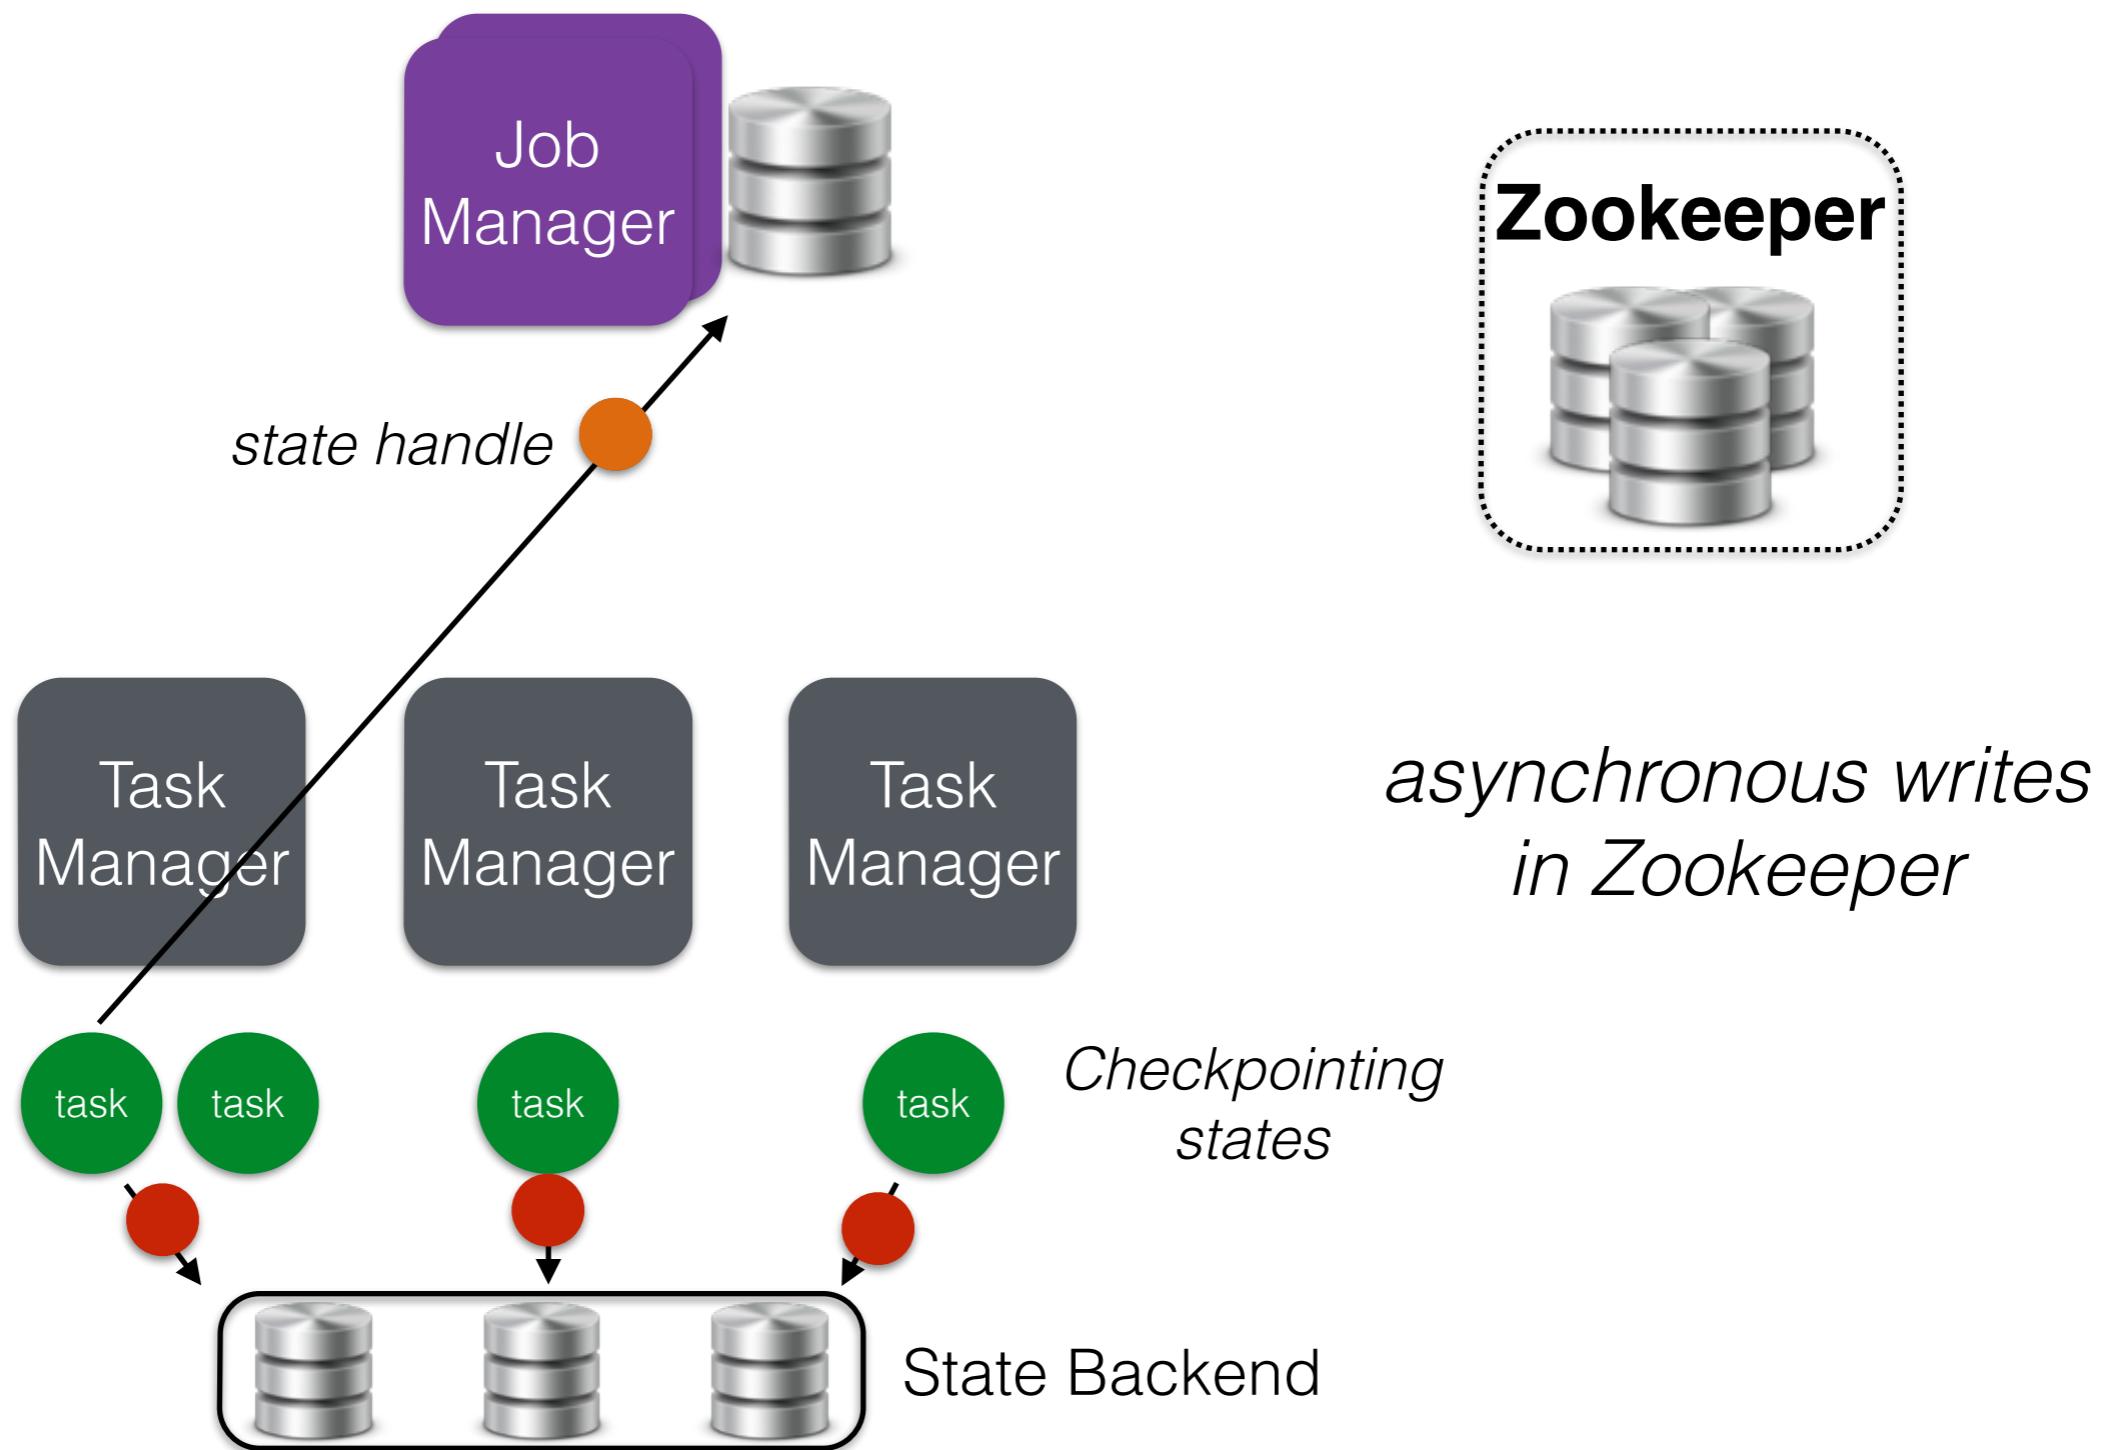
Logging Snapshots



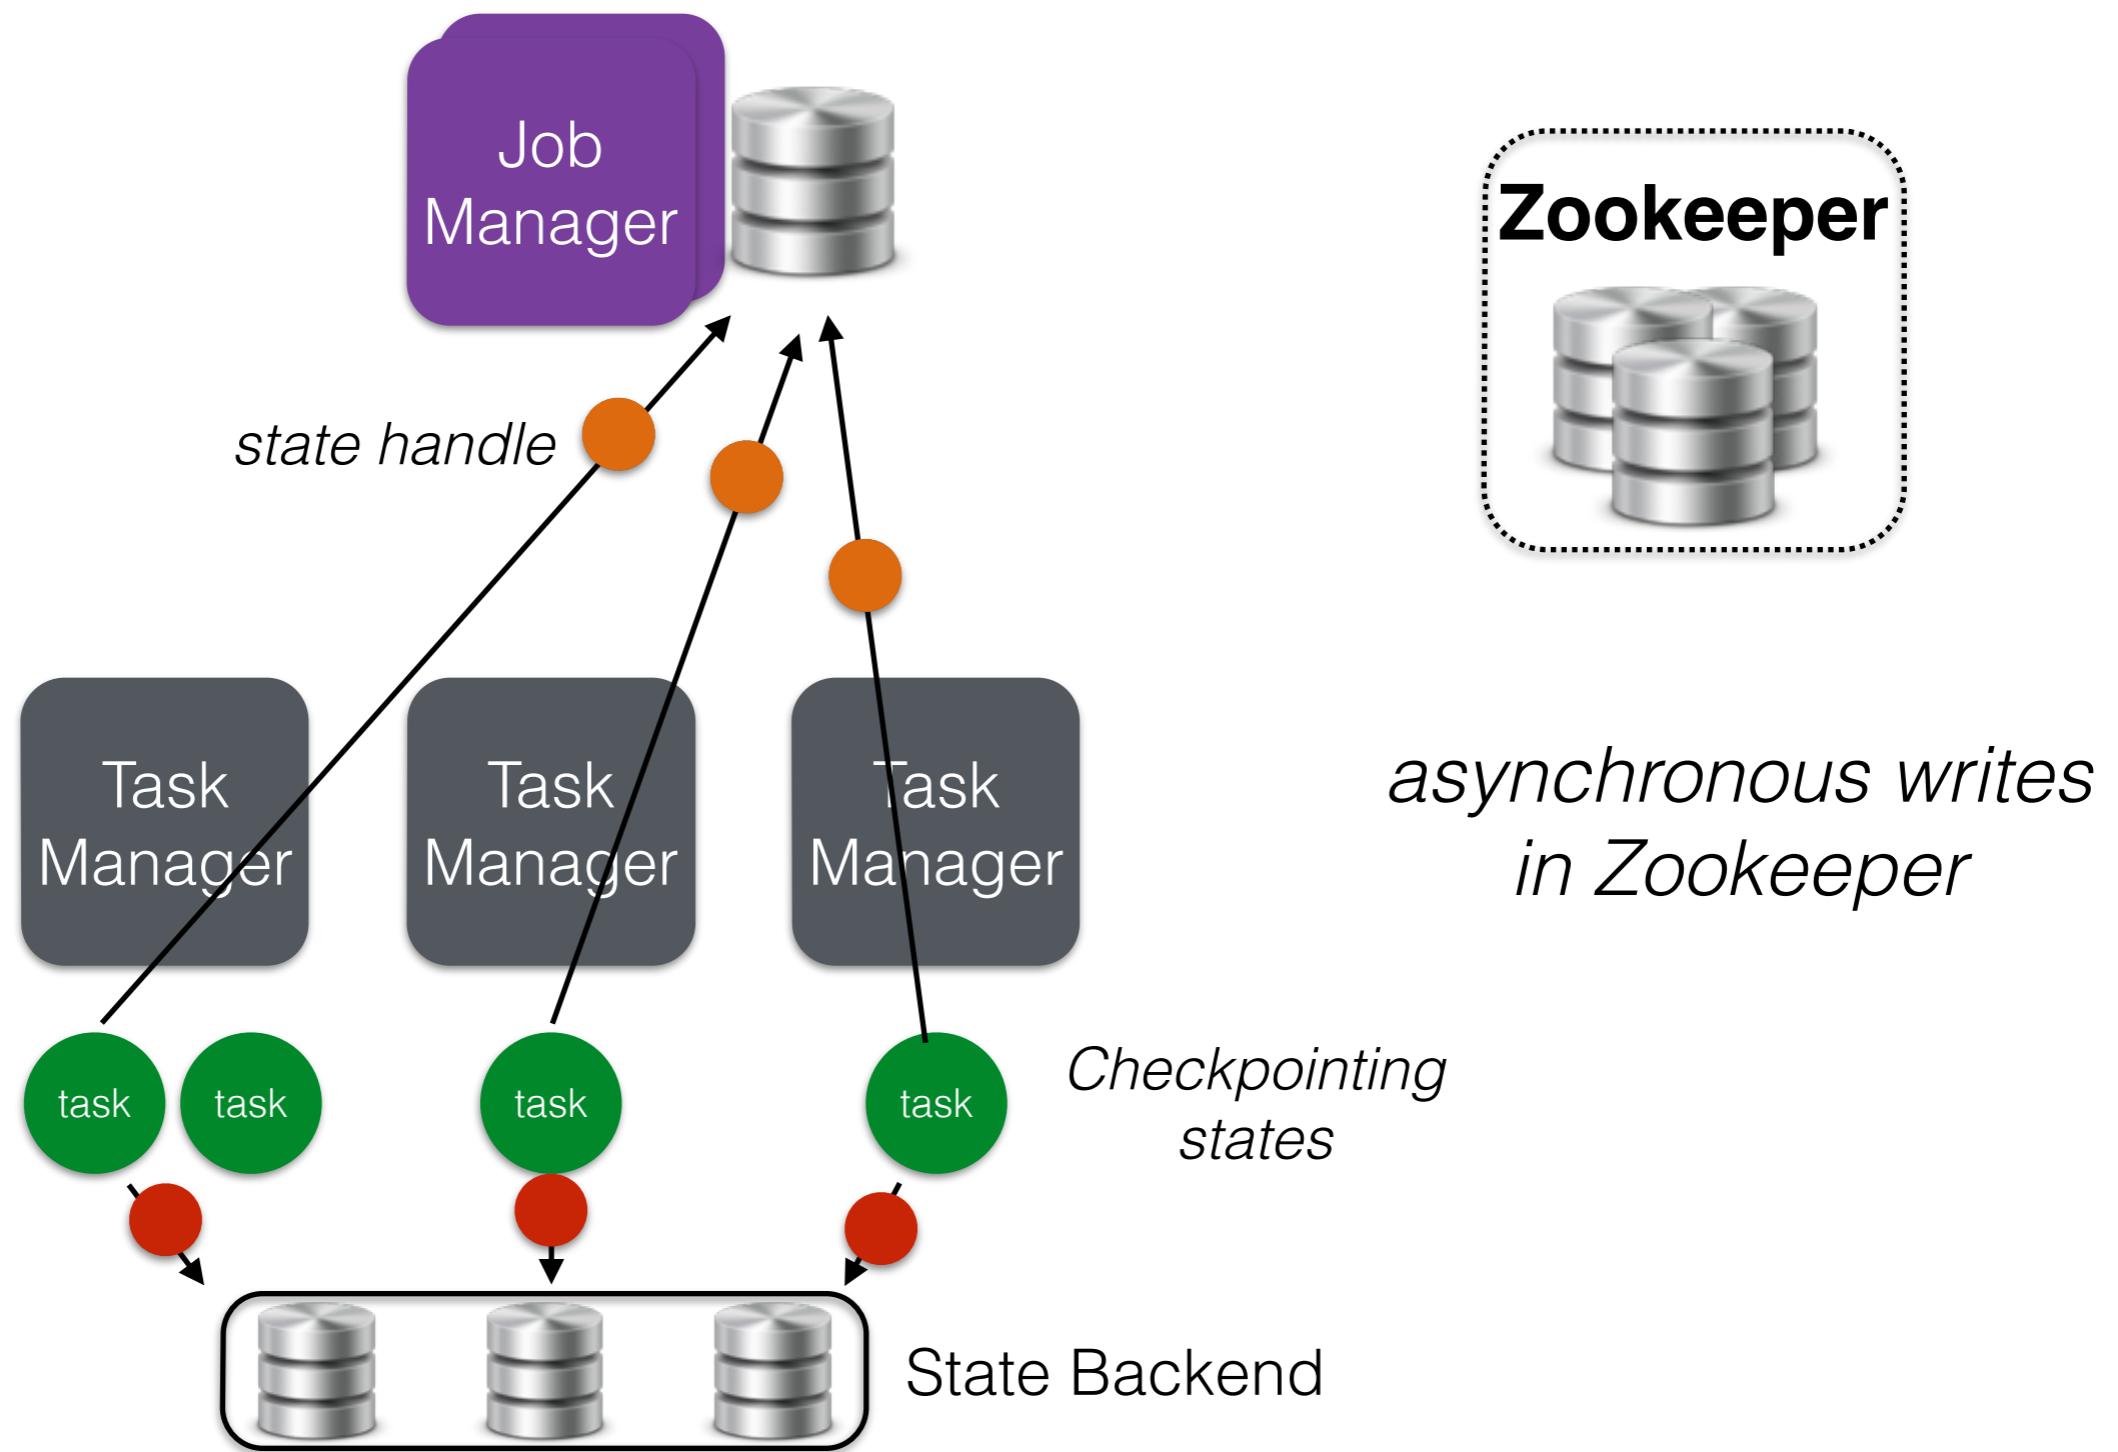
Logging Snapshots



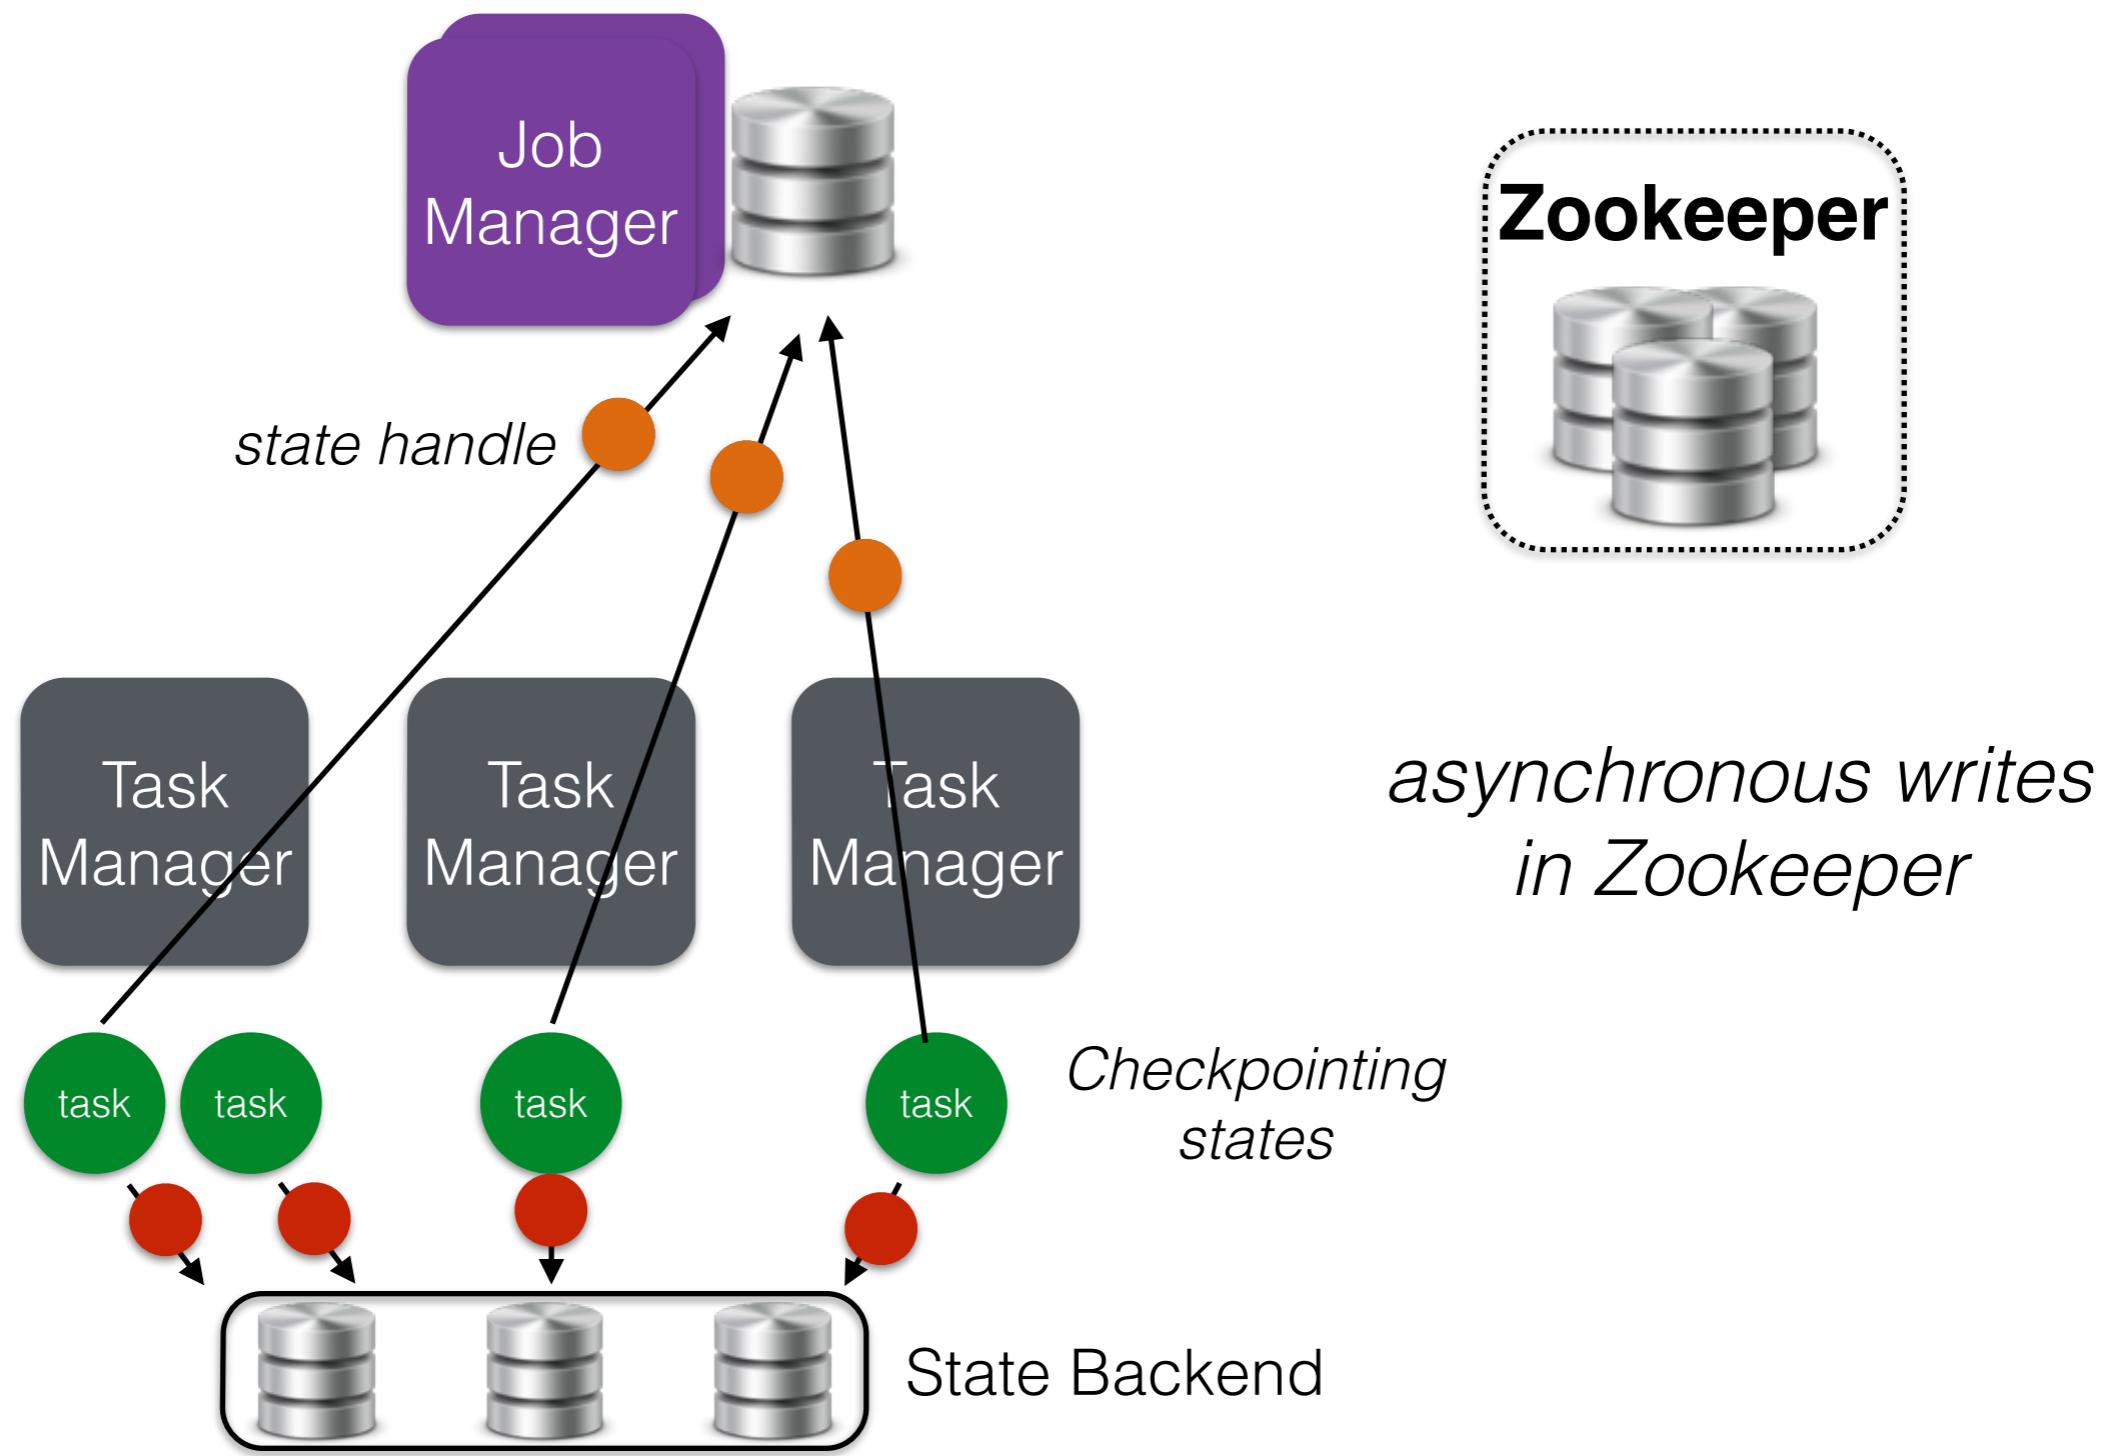
Logging Snapshots



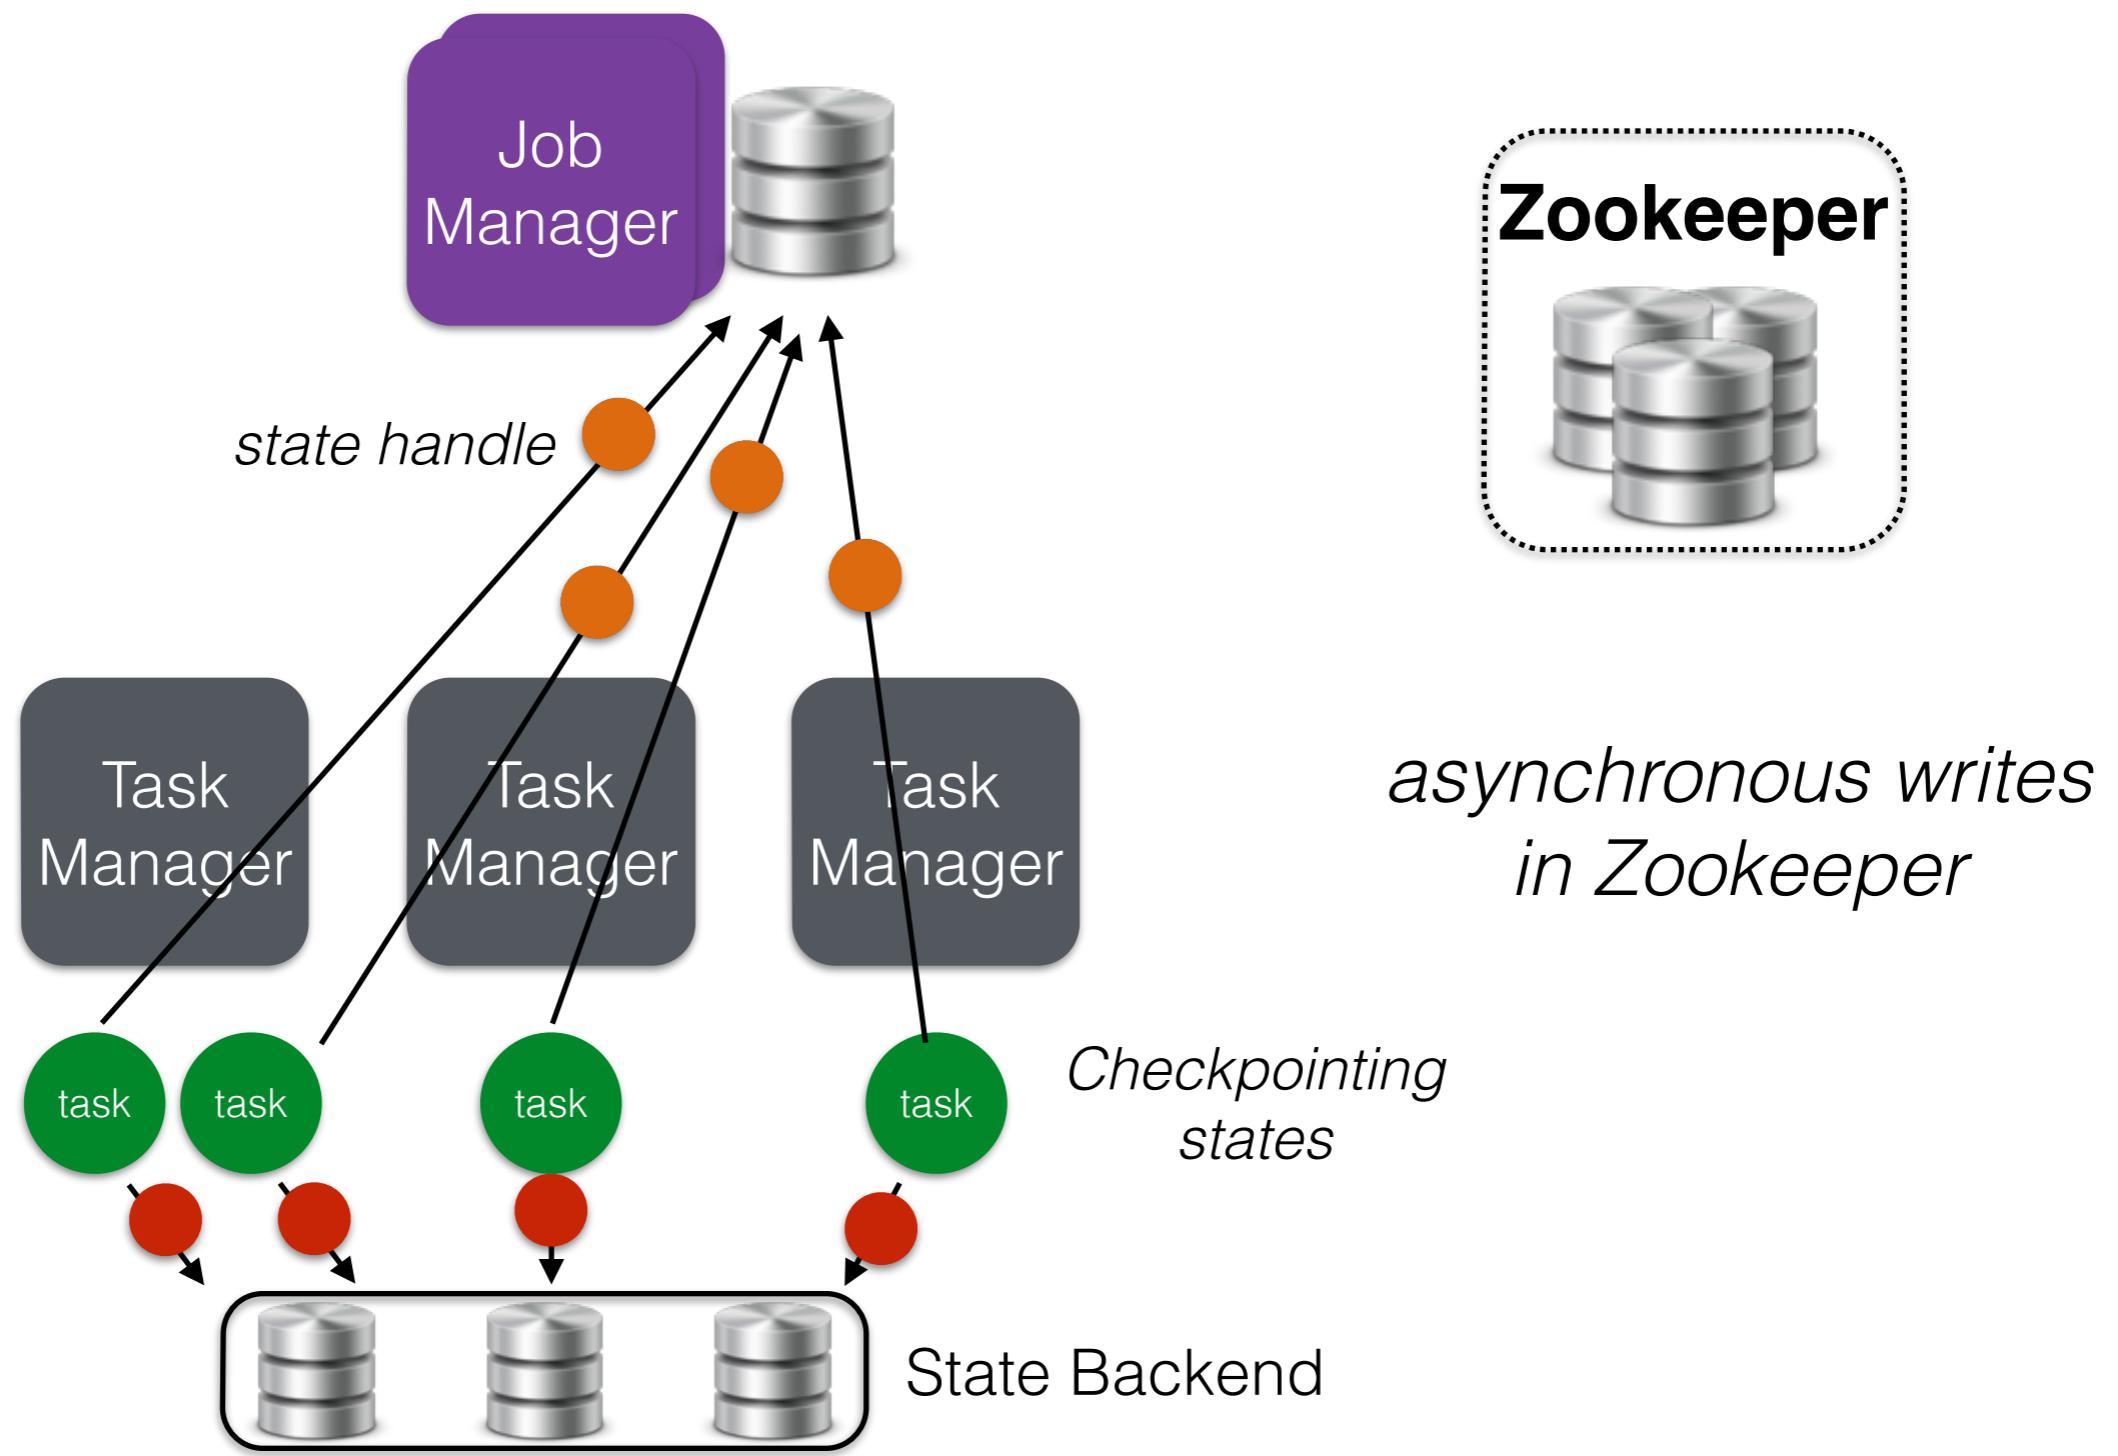
Logging Snapshots



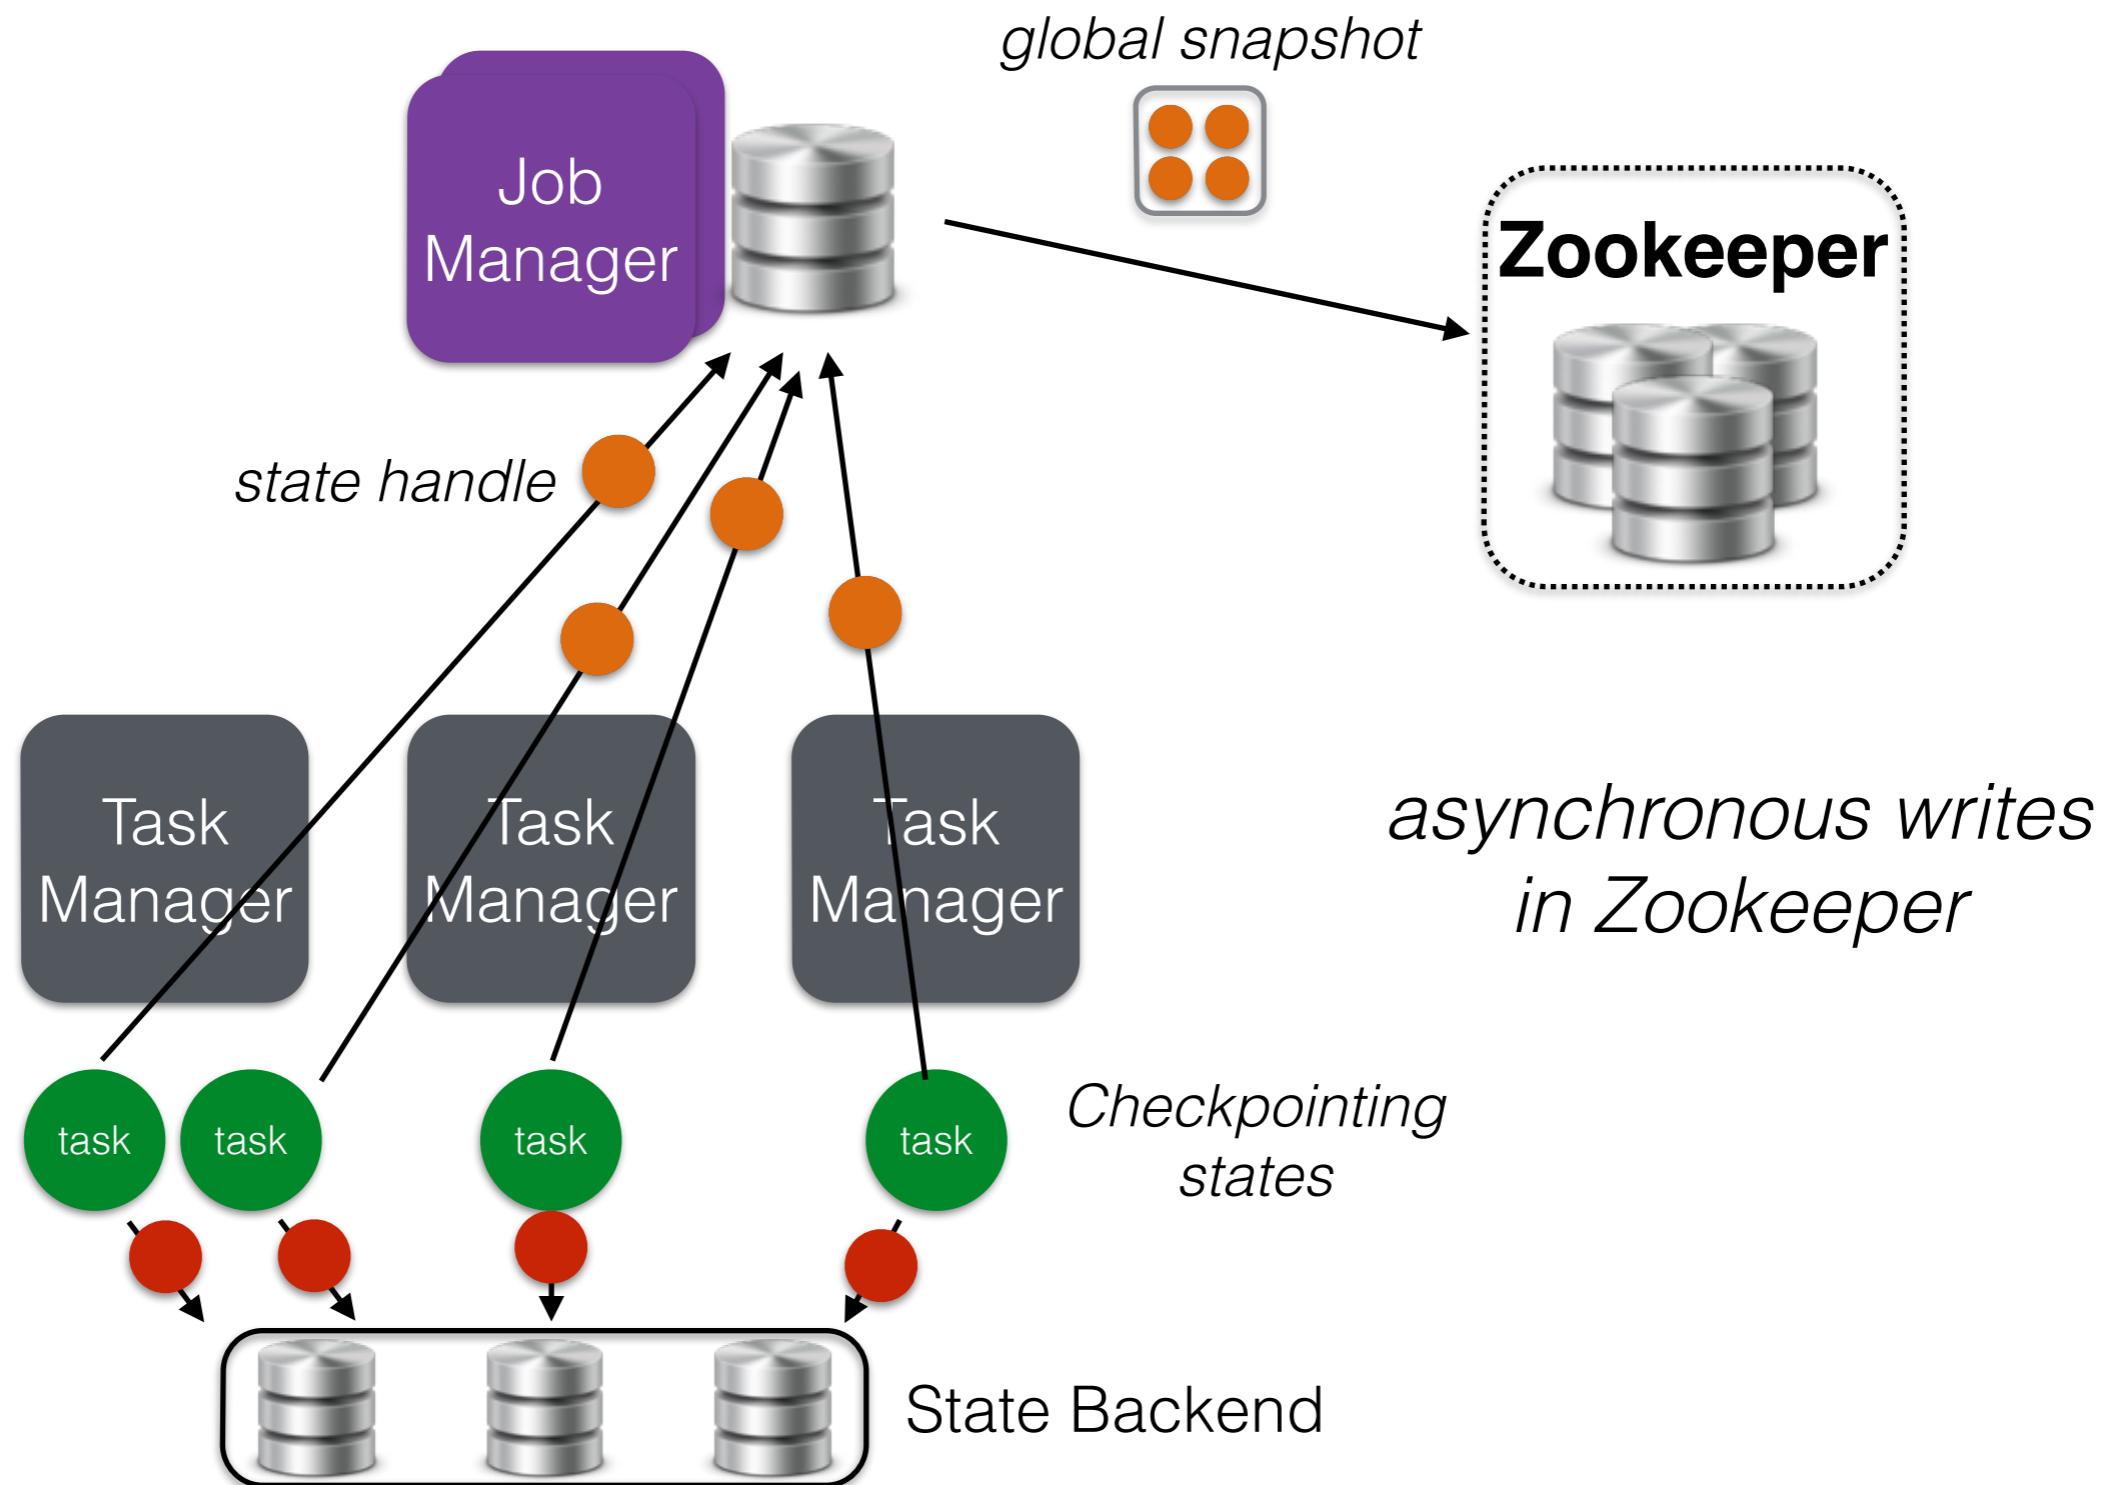
Logging Snapshots



Logging Snapshots



Logging Snapshots



HA in Zookeeper

- Task Managers query ZK for the current leader before connecting (with lease)
- Upon standby failover **restart** pending execution graphs
- **Inject** state handles from the last global snapshot

Summary

- The Flink execution monitors long running tasks
- Establishing exactly-once-processing guarantees with snapshots can be achieved without halting the execution
- The ABS algorithm persists the minimal state at a low cost
- Master HA is achieved through Zookeeper and passive failover